

Advancing Differentially Private Learning: Bridging Theory to Practice

Shuya Feng, Ph.D.

University of Connecticut, 2025

Privacy-preserving machine learning has become critical as organizations increasingly rely on sensitive personal data to train artificial intelligence systems. While Differential Privacy (DP) provides mathematically rigorous privacy guarantees, significant barriers prevent its widespread practical adoption. This dissertation addresses the fundamental theory-practice gap in differential privacy through a comprehensive research program spanning algorithmic innovation, systems development, and human factors analysis. Current differential privacy mechanisms face three interconnected challenges: algorithmic limitations in complex scenarios (infinite data streams, federated learning, high-dimensional data), framework fragmentation across specialized techniques with incompatible privacy accounting methods, and insufficient attention to human factors that determine real-world deployment success. This dissertation contributes four interconnected research thrusts that establish a pathway from theoretical innovation to practical deployment. We develop DPI (Differentially Private Infinite streams), the

first user-level DP solution with bounded privacy loss over infinite time horizons using sensitivity compression and multi-layer boosting. We establish UDP-FL (Universal Differentially Private Federated Learning), which harmonizes different DP mechanisms in distributed settings using Rényi DP as a unifying framework with convergence guarantees through mode connectivity analysis. We demonstrate DPED (Differentially Private Embedding Distillation), applying universal DP principles to high-dimensional text data through multi-layer noise injection and teacher-student distillation. Finally, we conduct comprehensive multi-stakeholder analysis examining data contributors, ML engineers, and model recipients to identify deployment barriers and provide practical guidelines. This research makes significant theoretical, practical, and methodological contributions to differential privacy. Theoretically, we establish the first universal DP frameworks with provable guarantees across streaming, federated, and high-dimensional scenarios. Practically, we deliver production-ready implementations with demonstrated performance improvements and comprehensive deployment guidelines. Methodologically, we demonstrate how rigorous theoretical advances can be systematically translated into real-world impact through integrated consideration of algorithmic design, systems implementation, and human factors. By providing both technical solutions and deployment methodologies that address stakeholder needs, this work establishes a comprehensive framework for advancing privacy-preserving technologies from theory to practice, enabling organizations to protect individual privacy while harnessing the collective value of sensitive data.

Advancing Differentially Private Learning: Bridging Theory to Practice

Shuya Feng

B.S., Harbin Institute of Technology, 2015

M.S., Illinois Institute of Technology, 2017

A Dissertation

Submitted in Partial Fulfillment of the

Requirements for the Degree of

Doctor of Philosophy

at the

University of Connecticut

2025

Copyright by

Shuya Feng

2025

Doctor of Philosophy Dissertation

**Advancing Differentially Private Learning: Bridging
Theory to Practice**

Presented by

Shuya Feng, B.S., M.S.

Approved by

Major Advisor : Yuan Hong

Associate Advisor: Derek Aguiar

Associate Advisor: Meisam Mohammady

Associate Advisor:

Associate Advisor:

University of Connecticut

2025

ACKNOWLEDGEMENTS

This dissertation was made possible through the support of many individuals and institutions. I offer my deepest gratitude to my advisor, Prof. Yuan Hong, for his unwavering guidance and commitment to bridging theory with practice. I am also immensely thankful to my dissertation committee, Prof. Derek Aguiar, Prof. Meisam Mohammady, Prof. Chuxu Zhang, and Prof. Suining He, for the valuable perspectives and challenging questions that significantly strengthened this research. My sincere thanks are extended to my collaborators, Prof. Binghui Wang, Dr. Ashish Kundu, and Prof. Doowon Kim, for their invaluable technical contributions. I am also grateful for the camaraderie and intellectual exchange with my lab colleagues: Han Wang, Hanbin Hong, Shenao Yan, Ali Atraste, Nima Naderloui, and Qin Yang, as well as the many students who worked alongside me. I gratefully acknowledge the University of Connecticut and the Department of Computer Science & Engineering for providing an inspiring academic home. Finally, I am profoundly grateful to my family for their unconditional love and encouragement. Their steadfast belief in me during moments of uncertainty has been my greatest source of strength and inspiration.

TABLE OF CONTENTS

1. Introduction	1
1.1 Background	1
1.2 Motivation	3
1.3 Contributions	4
2. DPI: Ensuring Strict <u>D</u>ifferential <u>P</u>rivacy for <u>I</u>nfinite Data Streaming	7
2.1 Abstract	7
2.2 Introduction	8
2.3 Preliminaries	16
2.4 Overview of DPI Framework	22
2.5 Detailed Design in DPI	24
2.6 Privacy and Utility Analysis	32
2.7 Experimental Evaluations	37
2.8 Discussion	49
2.9 Related Works	52
2.10 Proof and Further Discussions	55
2.11 Conclusion	64
3. Universally Harmonizing Differential Privacy Mechanisms for Federated Learning: Boosting Accuracy and Convergence	66
3.1 Introduction	66

3.2	Preliminaries	70
3.3	UDP-FL Framework	76
3.4	Theoretical Analyses	82
3.5	Experiments	92
3.6	Discussion	105
3.7	Related Works	107
3.8	Proofs and Further Discussion	111
3.9	Conclusion	125
4.	DPED: Multi-Layer Noise Distillation for Privacy-Preserving Text Embed-	
	dings	127
4.1	Abstract	127
4.2	Introduction	128
4.3	Related Work	130
4.4	Methodology	131
4.5	Privacy Analysis	133
4.6	Experiments	134
4.7	Detailed Privacy Analysis	137
4.8	Algorithm Pseudocode	139
4.9	Conclusion	139
5.	Usable Machine Learning Model Training with Differential Privacy . . .	142

5.1	Abstract	142
5.2	Introduction	143
5.3	Background and Related Works	147
5.4	Overview of Experiment Design	153
5.5	Experiment 1: Data Contributors	158
5.6	Experiments 2A & 2B: ML Engineers	163
5.7	Experiment 3: Model Recipients	168
5.8	Discussion	173
6.	Conclusions	177
	Bibliography	181

LIST OF FIGURES

2.1	Optional caption for list of figures	11
2.2	Optional caption for list of figures	21
2.3	Average MSE and KL divergence vs. ϵ for different values of the λ hyperparameter. The left column shows results for accumulative data disclosure, and the right column shows results for instantaneous data disclosure. . .	38
2.4	Average MSE and KL divergence vs. ϵ for different values of the μ hyperparameter. The left column shows results for accumulative data disclosure, and the right column shows results for instantaneous data disclosure.	39
2.5	Average MSE and KL divergence vs. ϵ for different values of the ζ hyperparameter. The left column shows results for accumulative data disclosure, and the right column shows results for instantaneous data disclosure. . .	40
2.6	Average MSE and KL divergence of statistical queries for ϵ and t across 1,000 time slots on COVID-19 dataset. (a), (b): Average MSE and KL divergence with varying ϵ . (c), (d): Average MSE and KL divergence over t when ϵ is 2.	44

2.7	Average Precision and Recall for anomaly detection over 2,700 time slots on Network Traffic dataset. (a), (b): anomaly detection with different ϵ . (c), (d): anomaly detection in time slot t ($\epsilon = 2$). Due to unsupervised learning on unlabeled data, non-private results are treated as the ground truth.	45
2.8	Average Precision and Recall for recommender system over 2,700 time slots on USDA dataset. (a), (b): Precision and Recall vs ϵ . (c), (d): Precision and Recall vs time slot t ($\epsilon = 2$). Due to unsupervised learning on unlabeled data, non-private results are treated as the ground truth. . .	46
2.9	Average MSE and KL divergence over 1,000 time slots. (a) (b): MSE vs t (varying ϵ). (c) (d): KL divergence vs t (varying ϵ).	48
2.10	Distribution spread of DPI output on the synthetic dataset with changing distribution. (a) (b): DPI output vs ϵ . (c) (d): DPI output vs t . The boxplots show the lower quantile (LQ), upper quantile (UQ), and median (M) of the DPI output densities. The spread of the distributions remains stable.	49
2.11	Real-time DPI outputs and true data (median).	50
2.12	Remaining budget for DPI with the random privacy budget allocation: (a) $\epsilon = 2$, and (b) $\epsilon = 10$	52
2.13	Evaluation of DPI on 200,000 time slots on synthetic datasets (a) KL divergence and (b) MSE: domain size 10,000.	52

2.14	Optional caption for list of figures	55
2.15	Information theory of 0-DP synopsis generation.	62
2.16	Total privacy loss of DPI and existing methods.	64
2.17	Evaluation of DPI on 10 synthetic datasets (each box includes 10 results for a specific privacy bound ϵ : the average MSE/KL of 2,700 time slots in each of the 10 datasets). (a) (b) MSE and KL divergence: domain size 1,000. (c) (d) MSE and KL divergence: domain size 10,000. Each experiment is repeated for 10 times, and the average results are plotted in the figures.	65
3.1	Accuracy and convergence results of UDP-FL and the baselines. 1) among the three mechanisms, the Staircase always performs the best with the same privacy budget; 2) UDP-FL obtains significantly better privacy- utility tradeoff and faster convergence than the baseline; and 3) UDP-FL (Staircase) even has a comparable accuracy with FedAvg (No DP). . . .	94
3.2	Optional caption for list of figures	96
3.3	Impact of the number of clients and sampling rate on UDP-FL. We observe that: 1) when the number of clients increases, shown in Figures (a) and (c), UDP-FL needs more epochs to converge; 2) with the increase of data sampling rate shown in Figures (b) and (d), UDP-FL converges faster.	98

3.4	Performance evaluation of UDP-FL on MNIST and Medical datasets. (a) and (c) present UDP-FL's accuracy under various differential privacy noise mechanisms, compared to a non-private baseline, with varying ϵ values. (b) and (d) illustrate the learning curves over training epochs for the MNIST and Medical datasets without privacy and with DP guarantees.	99
3.5	Optional caption for list of figures	106
3.6	Impact of Shuffler on UDP-FL with the privacy guarantee as Figure 3.4 (MNIST). We observe that: 1) the utility of UDP-FL can be marginally improved with the shuffler; 2) UDP-FL can converge slightly faster with the shuffler; and 3) UDP-FL using the Staircase mechanism performs the best.	121
3.7	Impact of Shuffler on UDP-FL with the privacy guarantee as Figure 3.4 (Medical dataset, with similar observations to the MNIST dataset). . .	122
4.1	DPED Overview. Disjoint teachers add noise to their outputs, which are then aggregated; the student learns from these privatized signals. . . .	129

4.2	Main experimental results on WikiText-2 (embedding training) and IMDB (downstream sentiment) under different privacy budgets. $\downarrow$ means lower is better, $\uparrow$ higher is better. Non-private is the upper bound with no noise. DP-SGD is training the embedding model with differentially private SGD. PATE is a baseline with teacher ensemble voting (single-layer). DPED is our full method with multi-layer distillation and rare-word-aware aggregator. Results are mean of 3 runs; standard deviations are shown for DP methods. Our approach consistently outperforms DP-SGD and PATE baselines in perplexity and downstream accuracy at equivalent (ϵ, δ)	135
4.3	Cosine similarity of DPED and DP-SGD regarding different word frequency.	141
5.1	Ecosystem of multi-stakeholder with DP-SGD.	144
5.2	Instructions of survey questions in Experiment 1, 2 and 3.	153
5.3	Web-based educational tools for DP-SGD concept delivery	158
6.1	Research roadmap: building secure and privacy-centric data solutions. . . .	178

LIST OF TABLES

2.1	Representative DP data streaming methods (not scalable to user-level protection over infinite data streams). Other existing methods share similar properties.	10
2.2	Characteristics of datasets (after pre-processing)	37
2.3	Expanded sensitivity (Δ) and privacy loss (actual DP) for the median query of packets in each time slot (0.5 seconds) on a network traffic dataset. $\epsilon = 0.5$ is used to generate noise for existing methods (ϵ is always bounded by 0.5 in DPI).	62
3.1	Comparison of representative DP-FL methods on accuracy, privacy, and convergence (observed based on theoretical analyses and empirical results in all the works).	70
3.2	Frequently used notations for convergence analysis.	82
3.3	Rényi DP of three main noise mechanisms utilized in UDP-FL along with their corresponding parameters.	84
3.4	UDP-FL accuracy vs. privacy guarantees.	100
3.5	Runtime of UDP-FL (sec) vs # training iterations.	100

3.6	Evaluation of the Shokri et al. [1], SOTA LiRA [2] and Canary [3] MIAs on CIFAR-10. TPR* denotes the TPR when FPR=0.01. The TPR*, ACC and AUC for Shokri et al. are 0.053, 0.710, and 0.757. The TPR*, ACC and AUC LiRA are 0.126, 0.651, and 0.716. The TPR*, ACC and AUC LiRA are 0.137, 0.649, and 0.719.	102
3.7	Evaluation on the SOTA InvGrad DRA [4] on CIFAR-10. The MSE, PSNR and SSIM for the Non-private method are 1.7104, 9.79, and 0.0751, respectively.	103
3.8	Evaluation on the SOTA AIA [5] on UTKFace. The Accuracy and F1-score for the Non-private method are 0.86 and 0.85, respectively.	104
3.9	Evaluation for data reconstruction attacks.	125
4.1	Ablation on use of hidden representation in distillation ($\epsilon = 2$). Providing the student with the noisy hidden representation in addition to the noisy label significantly improves performance.	136
5.1	Summary of Experimental Design	157

Chapter 1

Introduction

1.1 Background

In an era where data-driven insights fuel technological advancement, the tension between utility and privacy [6] has never been more pronounced. Organizations increasingly rely on sensitive personal data—from social media posts and location traces to financial records and health information—to train machine learning models [7] that power everything from personalized recommendations to automated decision-making systems. Yet this data-hungry paradigm raises fundamental questions: How can we harness the collective value of personal data while protecting individual privacy? How do we ensure that the models we train don't inadvertently reveal sensitive information about the people whose data they were trained on? Differential Privacy (DP) [8, 9], emerged as a mathematically rigorous answer to these questions. By providing a formal guarantee that the presence or absence of any individual's data has minimal impact on a system's output, DP offers a principled framework for privacy-preserving computation. The elegance of DP lies in its worst-case guarantee: even an adversary with

arbitrary auxiliary information and unlimited computational power cannot determine whether a specific individual's data was used in the computation. This theoretical foundation has sparked significant research interest, leading to sophisticated algorithms like Differentially Private Stochastic Gradient Descent (DP-SGD) [10], advanced privacy accounting methods using Rényi Differential Privacy [11], and specialized techniques for various data types and learning paradigms. Major technology companies have begun deploying DP methods at scale, with Google's RAPPOR system [12], Apple's privacy-preserving analytics [13], and Microsoft's WhiteNoise platform [14] serving as prominent examples. However, real-world applications present far more complex challenges than traditional DP mechanisms were designed to handle. Modern systems generate continuous data streams where traditional approaches suffer from unbounded privacy leakage. Federated learning scenarios involve multiple parties with heterogeneous privacy requirements. High-dimensional data types like text and images present unique challenges that often lead to excessive noise addition, destroying utility. Moreover, the DP research community has developed numerous specialized techniques that exist in isolation, each with different privacy accounting methods and incompatible noise mechanisms, creating a fragmented landscape that prevents unified deployment strategies.

1.2 Motivation

Despite significant theoretical advances and some industrial deployments, a critical gap persists between differential privacy theory and widespread practical adoption. This theory-practice gap manifests across three interconnected dimensions that collectively limit the real-world impact of DP research.

Algorithmic Limitations in Complex Scenarios. Traditional DP mechanisms were designed for relatively simple settings involving aggregate statistics over static datasets. However, real-world applications demand solutions for infinite data streams (social media feeds, IoT sensors), distributed learning environments with heterogeneous privacy preferences (federated learning), and high-dimensional complex data types (text, images) that suffer from the curse of dimensionality under traditional DP approaches.

Framework Fragmentation. The DP research community has developed numerous specialized techniques for different scenarios, but these approaches often exist in isolation with incompatible privacy accounting methods, different noise mechanisms, and domain-specific optimizations that don't generalize. This fragmentation forces organizations to choose between specialized solutions that may not integrate well or general-purpose approaches that may be suboptimal for their specific needs.

The Human Factors Blind Spot. Perhaps most critically, the DP research community has primarily focused on technical optimization while paying insufficient attention to the human factors that determine real-world adoption. Any practical DP system involves multiple stakeholders—data contributors who must trust privacy mechanisms,

ML engineers who must implement and maintain DP systems, and model recipients who must accept utility trade-offs. Current research typically considers these stakeholders in isolation, optimizing for mathematical privacy-utility trade-offs without considering implementation feasibility or user acceptance.

These challenges are interconnected: algorithmic limitations create the need for specialized solutions, contributing to fragmentation, while the lack of attention to human factors means that even technically sound solutions may fail in practice. What is needed is a comprehensive research approach that advances differential privacy mechanisms from theoretical foundations to practical deployment, addressing technical challenges while explicitly considering the human factors that determine real-world success.

1.3 Contributions

This thesis addresses the theory-practice gap through a comprehensive research program that advances differential privacy mechanisms across four interconnected dimensions. Our approach is guided by three core principles: universal mechanism design that adapts to diverse scenarios, rigorous theoretical foundations, and human-centered design that validates technical innovations through stakeholder analysis.

Strict DP for Infinite Data Streams (DPI). We address the fundamental problem of unbounded privacy leakage in streaming scenarios by developing the first user-level DP solution that provides bounded privacy loss over infinite time horizons [15]. Our key innovations include sensitivity compression using probability density function (PDF) data

structures, bi-directional boosting mechanisms that integrate multiple DP approaches, and random budget allocation strategies that prevent privacy budget depletion. This work establishes the theoretical foundation for user-level DP approaches while enabling privacy-preserving analytics on continuous data streams.

Universal DP Framework for Federated Learning (UDP-FL). We create a unified framework [16] that harmonizes different DP mechanisms in federated learning settings, enabling clients with diverse privacy preferences to collaborate effectively. Our contributions include a Rényi DP-based harmonizer that enables universal mechanism support, mode connectivity analysis for convergence guarantees across different mechanisms, and adaptive aggregation that optimizes for heterogeneous client requirements. This work extends universal DP principles to distributed learning scenarios with theoretical convergence guarantees.

Domain-Specific Applications: High-Dimensional Text Processing (DPED). We further demonstrate the applicability of DP principles to challenging high-dimensional text data through a novel teacher-student distillation framework [17]. Our innovations include multi-layer noise injection strategies that preserve semantic information, rare-word-aware aggregation that handles the unique challenges of sparse text features, and adaptive privacy mechanisms that optimize for text-specific characteristics. This work shows how universal DP frameworks can be specialized for complex data types while maintaining strong privacy guarantees.

Multi-Stakeholder Analysis of DP Deployment. We try to bridge the theory-practice

gap through the first comprehensive study of all stakeholders in the DP deployment pipeline, providing practical guidelines for real-world adoption [18]. This work includes data contributor perspective analysis on privacy mechanism acceptance, ML engineer usability studies across text and image modalities, and model recipient evaluation of privacy-utility trade-offs in deployed systems. By understanding the needs, constraints, and perspectives of all stakeholders, we will provide the missing link between technical innovation and practical deployment.

These contributions collectively represent a complete journey from theoretical innovation to practical deployment. The technical contributions (DPI, UDP-FL, DPED) establish universal DP frameworks with strong theoretical guarantees and empirical performance improvements, while the privacy attacks and human factors study validate that these innovations can be successfully deployed in real-world contexts. My broader research has also addressed other critical areas of data science, security, and machine learning. These complementary research activities inform the holistic, security-minded approach taken in this dissertation. Together, they provide not just technical solutions, but a methodology for advancing privacy-preserving technologies in ways that achieve both theoretical rigor and practical adoption, positioning differential privacy to fulfill its promise as a foundational technology for privacy-preserving machine learning.

Chapter 2

DPI: Ensuring Strict Differential Privacy for Infinite Data

Streaming

Shuya Feng, Meisam Mohammady, Han Wang, Xiaochen Li, Zhan Qin, Yuan Hong

2.1 Abstract

Streaming data, crucial for applications like crowdsourcing analytics, behavior studies, and real-time monitoring, faces significant privacy risks due to the large and diverse data linked to individuals. In particular, recent efforts to release data streams, using the rigorous privacy notion of differential privacy (DP), have encountered issues with unbounded privacy leakage. This challenge limits their applicability to only a finite number of time slots (“finite data stream”) or relaxation to protecting the events (“event or w -event DP”) rather than all the records of users. A persistent challenge is managing the sensitivity of outputs to inputs in situations where users contribute many activities and data distributions evolve over time. In this paper, we present a novel technique for *Differentially Private data streaming over Infinite disclosure* (DPI) that effectively

bounds the total privacy leakage of each user in infinite data streams while enabling accurate data collection and analysis. Furthermore, we also maximize the accuracy of DPI via a novel boosting mechanism. Finally, extensive experiments across various streaming applications and real datasets (e.g., COVID-19, Network Traffic, and USDA Production), show that DPI maintains high utility for infinite data streams in diverse settings. Code for DPI is available at <https://github.com/ShuyaFeng/DPI>.

2.2 Introduction

Most computer systems and applications continuously generate data streams for analyses. For instance, online media such as YouTube [19] or Instagram recommends videos to users based on their vast visiting records. Network monitoring [20] detects abnormal behavior and identifies anomalies while tracking the network traffic. Water management [21] makes timely decisions on purifying water per the real-time water quality data collected by sensors.

Existing solutions have not effectively addressed the privacy threats posed by streaming data, which is vulnerable due to its volume and diversity. Recent attempts to release data streams using the widely adopted privacy model, differential privacy (DP), have revealed persistent privacy leakage [22, 23, 24, 25, 26, 27]. In particular, Dwork et al. [22] (STOC’10) found that in streaming data, as the number of rounds increases, an unavoidable logarithmic increase in error for an advancing counter occurs. Similarly, a logarithmic growth in privacy leakage is observed when the error is bounded. This issue

arises from the infinite accumulation of two key factors in DP mechanisms: sensitivity (the maximum impact of each user on the result) and privacy budget composition (additional leakage from sequential result releases).

In the past decade, several solutions have been proposed to address “event-level privacy on finite or infinite streams” [23, 24, 25, 26, 27, 28, 29, 30], or “user-level privacy on finite streams under certain settings” (e.g., partially disclosing some randomized results and applying interpolation [31, 32, 33, 34, 35]). However, they have limitations in practice. The first category struggles when users’ contiguous events collectively reveal sensitive information. Event-level privacy safeguards individual location visits but fails to protect a user’s path across successive, potentially indefinite timestamps, like a sequence of locations. The second category is less practical for continuously and indefinitely collected data streams, since it is unrealistic to predict when services like traffic reporting will cease.

Thus, designing user-level DP mechanisms for infinite streaming involves addressing five fundamental challenges.

- **Unboundedness.** The data size in the stream is continuously growing and unbounded, with each user’s data potentially generated indefinitely, making it also unbounded in infinite data streams.
- **Dynamic Data Distribution.** The data distribution changes dynamically over time, requiring DP mechanisms to consistently offer accurate analysis and privacy

guarantees upon these changes.

- **Accumulative Privacy Loss over Infinite Streams.** This has been widely recognized but not addressed yet.
- **Utility-Privacy Tradeoff.** It is challenging to minimize the utility loss continually (under a bounded leakage) while leveraging dynamic changes in data distribution.
- **Real-Time Processing.** DP mechanisms for infinite data streams must balance real-time processing needs with privacy protection and computational efficiency.

Table 2.1: Representative DP data streaming methods (not scalable to user-level protection over infinite data streams). Other existing methods share similar properties.

Methods	Assumption	Properties in Existing Works		
	(Sensitivity, Setting)	Privacy	Error Bound	Applications
[24]	(1, event)	$O(\epsilon \cdot t)$	$O(\frac{\log t^{1.5} \cdot \log \frac{1}{\delta}}{\epsilon})$	count, top-k, range queries
[25]	(1, event)	$O(\epsilon \cdot t)$	-	count, event monitoring
[27]	(1, event)	$O(\epsilon \cdot t)$	$O(\tau(\log t)^{1.5})$	count, range queries, etc.
[28]	(1, event)	$O(\epsilon \cdot t)$	w.r.t. query range, threshold, ϵ, t	multiple analyses

To our best knowledge, existing methods have not adequately addressed the significant challenges in infinite data streams. Challenges in current methods include not supporting infinite data disclosure, assuming limited sensitivity instead of indefinite full event sequence protection, and limiting applications to fixed time slot queries. Table 2.1 provides an overview of the most crucial properties of representative existing works on data streaming methods with differential privacy. There are many limitations in existing methods. For instance, they all consider the special case of sensitivity (e.g., 1 for counting) and their privacy bound linearly increases with the number of rounds t .

In this paper, we present a novel DP framework for infinite data streams, known as DPI (*Differential Privacy for Infinite Data Streams*). DPI is designed with three key components to address the above challenges: (1) *Sensitivity Compression*, (2) *DPI Boosting Mechanism*, and (3) *Privacy Budget Allocation Mechanism*, discussed as below.




Fig. 2.1: DPI guarantees limited privacy bound and sensitivity by managing the pace of convergence and error margins, striving to attain the highest level of accuracy (*the output streaming PDFs can support most real-time applications as the input stream*).

2.2.1 DPI in a Glimpse

Sensitivity Compression. Bounding sensitivity is crucial for reliable DP models. DPI adopts a *probability distribution data structure* to effectively limit the unbounded sensitivity by projecting/compressing the maximum output difference to the ℓ_1 -norm of probability distribution difference, offering advantages over unbounded structures like counting.

DPI Boosting Mechanism. The design of DPI has two inherent crucial challenges: (1) preserving a (lower-bound) meaningful utility for the data stream, and (2) developing methods to capture dynamic patterns within the new coming data. A sequence of randomization mechanisms (e.g., additive noise [36], sampling [37] or randomized

response [38]) for all the time slots, inevitably destroys both cumulative and instantaneous information. For instance, applying a series of ϵ -DP Laplace mechanisms over a dataset D will generate results with an error of $O(2^n/(|D| \cdot \epsilon))$ [39].

Unlike atomic DP mechanisms, additive or accumulative learning algorithms are especially effective for streaming data, where they continuously update models in real-time to adapt to evolving patterns. In particular, Dwork et al. [40] and Moritz et al. [41] demonstrated that the differential privacy version of boosting mechanisms can be effective in achieving cumulative learning outcomes for a finite number of queries. However, direct application of these approaches to continual streaming data results in *unbounded privacy leakage* and *significant computational overhead* [40, 41]. We have identified that both deficiencies result from the use of an expensive *base synopsis generator* in [40, 41].¹

To handle infinite streams, we have devised a novel DPI boosting mechanism, which guarantees both bounded and efficient computation while integrating an infinite number of tiny-budget DP mechanisms into an additive learning game. While the outcomes of various queries on streaming data may vary over time, it is still possible to precisely associate each result with an adequately accurate state, such as a numerical value or a PDF. The DPI boosting mechanism introduces a groundbreaking approach to synopsis generation by creating a significantly smaller set of synopses from PDFs, resulting in faster processing. Such non-trivial synopsis generation is designed as a

¹ Such generator acts as a compact representation of query answers over a dataset and can be implemented using various forms, such as a synthetic database or any data structure such as kernel or probability density function (PDF) that approximates query results (“synopsis”) [40, 41].

privacy-free (0-DP) mechanism, generating an ample number of synopses that precisely address a set of queries across an infinite number of rounds. It necessitates a rigorous analytical analysis to determine the number of synopses required to achieve a meaningful level of accuracy while ensuring privacy. Also, the existing uni-directional boosting mechanism [40, 41] (updating weights solely over queries based on data-dependent synopses) should be revised into a *bi-directional boosting* approach that considers both data-independent synopses and queries for infinite data streaming (*otherwise, the privacy budgets will not be sufficient*). With these considerations in DPI, we aim to achieve provable utility while maintaining the bounded privacy guarantees.

Privacy Budget Allocation and Boosting. By saving budget on the synopsis generation process in DPI, we develop a bounded budget boosting mechanism that optimally utilizes an infinite number of tiny boosting mechanisms for making refinements, focusing the privacy budgets solely on *reweighting* and not synopses. To achieve this, we establish an optimal allocation strategy with a converging sum series to maximize the utility.

To prevent budget depletion after a certain number of releases, typically expected with a decaying series, we employ a sophisticated deployment of this infinite decaying series of privacy budgets. During each iteration (time slot) of DPI, we randomly generate the privacy budget from the series, allowing for efficient and controlled management of the privacy budget throughout the continuous process.

Figure 2.1 shows an overview of DPI, including the input data stream, the three key components, and the output results (as streaming PDFs). The DPI framework takes an

incoming data stream as input and processes it through sensitivity compression, privacy budget allocation, and a boosting mechanism to output an infinite series of differentially private results that maximize accuracy and utility over time (see the detailed design in Section 2.5). This end-to-end process allows DPI to ensure strong privacy for users while handling the challenges of infinite streaming data. Note that DPI may require a restart to refresh the privacy budget in possible uncommon cases of extremely long-term, highly dynamic data streams, as discussed in Section 3.6.

2.2.2 Contributions

To sum up, DPI makes the following major contributions:

Unique and Significant Benefits. First, to our best knowledge, DPI provides the first user-level DP solution for protecting users in applications relying on continuously collecting data over infinite streams, e.g., YouTube [19] or Instagram recommends videos to users based on their vast visiting records indefinitely. Second, unlike most existing works which usually take various limited assumptions (e.g., bounded sensitivity, protecting the presence/absence statuses of events rather than users, a limited number of disclosures), DPI strictly satisfies generic ϵ -DP for protecting all the users (end-to-end) involved in the unlimited data streams. Third, DPI provably maximizes the accuracy of the disclosure in the long run with substantial theoretical analyses.

Novel Methods. We have designed a novel dynamic boosting mechanism (DPI boosting) that optimally balances both privacy and utility over time. DPI significantly differs from

the “AdaBoosting” [40] for DP: (1) DPI replaces the base synopsis generator with a pre-determined set of synopses *independent of the dataset*. We demonstrate that these synopses can be created by the performance of the output and strong synopses can be integrated into the boosting mechanism with no impact on privacy, and (2) DPI boosting has been designed to generate a PDF data structure that tracks the overall PDF of the stream instead of accurately answering pre-defined queries. We have also established lower bounds on the utility of DPI based on this PDF data structure.

Comprehensive Evaluations. We comprehensively evaluate the performance of DPI on different applications and datasets. Our primary objectives in these evaluations are: (1) investigating the impact of DPI’s hyperparameters on its performance over time (in Section 2.7.2), (2) evaluating DPI on different datasets with different domain sizes/distributions and a wide variety of real applications, such as statistical queries, anomaly detection, recommender systems (in Sections 2.7.3 and 2.7.4), and (3) compare DPI with existing methods and show that they violate the privacy requirement ϵ after a few rounds (in Appendix 2.10.9) while DPI strictly bounds the privacy leakage with ϵ indefinitely and still has considerable share of unspent privacy budget after a large number of rounds.

2.3 Preliminaries

2.3.1 Streaming Model

We consider data collection and analysis in the context of interactive and continuous observation over infinite data streams. In this setting, a trusted curator maintains a dynamic dataset, represented as a collection of streaming input $S = \{S_1, S_2, \dots\}$, where S_t refers to the indexed input collected at time slot t . The universe of possible individual records is denoted as X .

Interactive Queries. The data analyst interacts with the curator by submitting a sequence of queries over time. Let $\mathcal{Q}$ be the space of all possible queries, and at each time slot t , the data analyst requests a subset of queries $\mathcal{Q}_t \subseteq \mathcal{Q}$. Each query q_t is a function that takes the data until the current stream S_t and outputs a response $y_t = q_t(S_1, S_2, \dots, S_t)$, which is released by the curator. Note that q can remain the same as previous time slots or dynamically change.

2.3.2 Privacy Models

To protect the privacy of individual records in the context of continual observation, the trusted curator employs Differential Privacy (DP) mechanisms represented by $\mathcal{M}(\epsilon)$, where $\epsilon > 0$ is the privacy budget allocated for each query in the infinite data stream. Each query q_t outputs a response $y_t = q_t(S_1, S_2, \dots, S_t)$, and the curator releases a private response denoted as $\tilde{y}_t$. The DP mechanism $\mathcal{M}(\epsilon_t)$ ensures that $\tilde{y}_t$ does not disclose any information about the presence or absence of sensitive information from

any particular individual in all the streamed data. By limiting the privacy budget ϵ_t , the curator controls the amount of privacy provided for each query q_t .

We adopt the standard definition of ϵ -DP [36]. Let Γ be the domain of datasets, and let D, D' be two datasets in Γ such that D' can be obtained from D by adding or removing the data of a single individual (user). A randomization mechanism $\mathcal{M} : \Gamma \times \Omega \rightarrow \mathbb{R}$ is ϵ -differentially private if, for all $\Theta \subseteq \mathbb{R}$ and for all $D, D' \in \Gamma$ such that D' is adjacent to D , i.e., D and D' differ in the data of a single individual, the following inequality holds:

$$\Pr(\mathcal{M}(D) \in \Theta) \leq e^\epsilon \Pr(\mathcal{M}(D') \in \Theta) \quad (2.1)$$

where Ω is a sample space by randomization $\mathcal{M}$ to generate the output. To achieve differential privacy, we need to consider the concept of sensitivity. The sensitivity Δ [9] of a query q is defined as the maximum ℓ_1 -distance between the exact query answers on any two neighboring databases D and D' , i.e., $\Delta q = \max_{D, D'} \|q(D) - q(D')\|_1$. Bounding sensitivity is crucial for DP models. Even with a constrained privacy budget, unbounded sensitivity can result in overly sensitive outputs, risking privacy breaches.

2.3.3 Boosting Mechanisms

We next introduce some background for the AdaBoosting mechanism and DP with AdaBoosting, which will be significantly revised on both *privacy* and *efficiency* for DPI.

AdaBoosting Mechanism

AdaBoosting (Adaptive Boosting) [40] is a powerful machine learning technique used to enhance the accuracy of weak classifiers by combining them into a strong ensemble classifier. Formally, given a training dataset $D = (x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$, where x_i and y_i represent the data points and their respective labels, AdaBoosting constructs an ensemble classifier $H(x) = \text{sign}\left(\sum_{t=1}^T \alpha_t h_t(x)\right)$. Here, T is the number of weak classifiers, $h_t(x)$ is the t -th weak classifier, and α_t is the corresponding weight assigned to $h_t(x)$. The objective of boosting is to convert a weak learner, which produces a hypothesis only slightly better than random guessing, into a strong learner that achieves high accuracy. The AdaBoosting proceeds in rounds. In each round, two steps are performed:

1. The base learner is run on the current distribution $\mathcal{D}_t$, generating a classification hypothesis h_t .
2. The hypotheses $h_1, h_2, \dots, h_t$ are used to reweight the samples and define the updated distribution $\mathcal{D}_{t+1}$.

It continues for a predefined number of rounds or until the combination of hypotheses is sufficiently accurate.

DP with AdaBoosting

DP with AdaBoosting, introduced by Dwork and Rothblum [40], combines AdaBoosting with differential privacy. As detailed in Figure 2.14 (in Appendix 2.10.1), the process iteratively updates the weights assigned to each query, and samples queries based on their weights as follows. Let $Q = \{q_1, q_2, \dots, q_m\}$ represent a set of queries. The method assigns uniform weights w_i to each query q in Q in its initialization phase ($P_Q(t = 1) \rightarrow$ uniform). Next, it will operate on the “base synopsis generator” to produce “weak” DP synopses $\mathcal{A}_{Q_s}^t$, which are moderately accurate for r sampled queries Q_s .

Definition 1 (Sampled Queries Q_s). *Let Q be the set of requested queries, and $P_Q(t)$ be the PDF for the query sampling at iteration t in DP AdaBoosting [40]. At each iteration t , the set of sampled queries Q_s^t is defined as a k -size subset of Q , sampled from Q per $P_Q(t)$. The sampling process reflects AdaBoosting’s practice of assigning larger weights to weaker queries to improve their performance.*

Definition 2 (Synopsis). *A synopsis [40, 42] generated from a dataset D is defined as a compact representation of answers to a set of queries $Q = \{q_1, q_2, \dots, q_k\}$ over D . Synopsis can be in the form of a synthetic database, or any arbitrary data structure, e.g., a kernel or PDF, that can be queried to return an approximation of the query’s result.*

This data structure is referred to as a *synopsis* of the database. General privacy-preserving synopses are of interest as they may be easier to construct than privacy-

preserving synthetic databases. Also, stronger limitations are known for constructing synthetic databases than general privacy-preserving synopses.

Definition 3 (Base Synopsis Generator [40]). *For any data universe X and set of queries $Q : \{X^n \rightarrow \mathbb{R}\}$ with sensitivity Δ_Q , a $(r, \lambda, \varepsilon_1, \kappa)$ -base synopsis generator for Q is a method that:*

- *Outputs a synthetic database y of size n based on $r = O(n \cdot \log |X| \cdot \kappa)$ queries from Q , where κ is the probability that the chosen synopsis is λ -accurate (the synopsis closely approximates the true data within a λ margin of error).*
- *Adds Laplace noise $O(\kappa\sqrt{r} \cdot \frac{\Delta_Q}{\varepsilon_1})$ to each query answer.*
- *Guarantees $(\varepsilon_1, \exp(-\kappa))$ -differential privacy.*
- *It has a complexity of $|X|^n \cdot \text{poly}(n, \kappa, \log(1/\varepsilon_1))$.*

The base synopsis generator outputs the lexicographically first database y satisfying $|q(y) - \hat{q}_s^t(x)| \leq \lambda/2$ for every query q in the sampled query set Q_s , where $\hat{q}_s^t$ is the noisy query answers. If no such database exists, it outputs $\perp$ (i.e., fails). The $(r, \lambda, \varepsilon_1, \kappa)$ -base synopsis generator privately generates a synthetic database approximating the responses of the original database to a set of queries Q .

The goal is to iteratively improve the base generator using the boosting algorithm to create a synopsis that performs well for all requested queries Q while preserving privacy.




Fig. 2.2: The DPI framework in time slot t . (1) Q_s^t can be a single query or a subset of queries, and can remain the same or dynamically change over time, (2) all the privacy budgets are pre-computed with converging series and ϵ_t is sampled in real-time, and halved for query-sampling and synopsis-sampling, (3) generator returns the PDF output (sampled synopsis) for the current time slot t to the interface/analyst, and also sends it for boosting, and (4) DPI boosting (bi-directional) reweights the PDFs (both query-sampling probability and synopsis-sampling probability) for the next time slot $(t+1)$.

The next steps aim to strengthen the queries with bad results by adaptive sampling and refining them. At each iteration, the query sampling probability $P_Q(t)$ is reweighted based on the performance of the base synopsis generator over all queries such that the probability modification is to iteratively assign higher weights to the (new) queries with bad results, allowing the subsequent weak model to focus on these challenging examples and improve overall accuracy. Finally, it integrates a learning rate η_t^Q in the reweighting process, which is proportional to the privacy budget allocated for boosting (ϵ_2). Then, two separate privacy budgets are required: one for generating synopses (ϵ_1) and the other for the reweighting step (ϵ_2). Consequently, AdaBoosting with DP [40] becomes a privacy-expensive mechanism.

DP with AdaBoosting's runtime depends on the number of queries and the base synopsis generator's runtime. The accuracy of the mechanism scales with $\sqrt{n} \log m$ (n : size of the data, m : number of queries). For details of the synopses and base synopsis

generator, see [40, 42].

2.4 Overview of DPI Framework

In this section, we present the DPI framework comprising the DPI Interface and the DPI Brain. The DPI Interface manages temporal input streams, stores sequential DP results for pre-defined queries (“query synopses”), and delivers query results to analysts over time. On the other hand, the DPI Brain focuses on online learning. The details of the DPI framework are shown in Figure 2.2.

2.4.1 DPI Components

DPI Interface plays a central role in managing the interaction between analysts and the underlying data, including:

- *Input Module* collects temporal data streams from users over successive timestamps. It also receives temporal queries $Q_1, \dots, Q_t$ from the analyst.
- *Sensitivity Compression* processes the input data stream with a PDF data structure to bound sensitivity.
- *Output Module* returns temporal output query synopses as a sequence of PDFs with DP guarantees for a pool of queries $Q_1, \dots, Q_t$, which can be converted to support most data analyses on the input stream.

DPI Brain comprises the following modules.

- *Privacy Budget Allocation* generates the privacy budget ϵ_t for each time slot t .
- *Generator Module* generates the sampled queries (Q_s^t) and a set of sampled synopses ($\mathcal{A}_{Q_s^t}^t$) for each query $q \in Q_s^t$. It interacts with the DPI Boosting module, receiving the updated configurations in terms of PDFs to effectively adjust the sampling process.
- *DPI Boosting Module (bi-directional)* effectively captures the dynamic nature of data streams (new coming data and queries) when answering queries.

2.4.2 DPI Workflow

- **Step 1:** *Input Streaming and Sensitivity Compression.* See details in Section 2.5.1.
- **Step 2:** *Privacy Budget Allocation.* In time slot t , DPI randomly samples a pre-computed privacy budget ϵ_t from the series, halved for query-sampling and synopsis-sampling (see details in Section 2.5.4).
- **Step 3:** *0-DP Synopsis Generation.* The Generator module generates a set of parameters, including a sampled subset of queries Q_s^t over the pre-defined pool of queries Q and a set of sampled synopses $\mathcal{A}_{Q_s^t}^t$ for queries $q \in Q_s^t$. To derive the sampled synopses $\mathcal{A}_{Q_s^t}^t$, we need to generate a pool of synopses $\mathcal{A}$ that can represent all possible distributions in a dataset. This generation can avoid

consuming privacy budget (0-DP, see details in Section 2.5.2).

- **Step 4: *DPI Boosting*.** The DPI Boosting module improves accuracy via bi-directional reweighting of queries and synopses (see details in Section 2.5.3).
 - *Reweighting Synopsis-Sampling Probability.* In this phase, the DPI Boosting module uses a portion of the privacy budget allocated for DPI to perform reweighting of the Synopsis-Sampling Probability. This step aims to improve the accuracy of synopses and update the sampling PDF accordingly.
 - *Reweighting Query-Sampling Probability.* In this phase, another portion of the privacy budget is used to reweight the Query-Sampling Probability. This step ensures that the most informative queries are given higher probabilities in the sampling process.
- **Step 5: *Output Private Streams (PDFs)*.** DP query results (temporal) are provided to analysts over time.

2.5 Detailed Design in DPI

2.5.1 Sensitivity Compression

To enable effective privacy for infinite data streams, the sensitivity of the analysis must be bounded. Sensitivity represents the maximum change in the output resulting from the presence or absence of any single user’s data. This motivates the need for sensitivity

compression in the infinite stream setting. Our key insight is adopting a PDF data structure to represent the data stream. This provides an inherent sensitivity bound, as changing one user’s data can only alter the PDF by a maximum total variation distance of 2 (considering the two PDFs of two adjacent inputs are two normalized vectors with all the entries summing up to 1 in each). Denote D and D' as the output distributions of two neighboring inputs for any query at time slot t . Thus, for all D and D' , their sensitivity is derived as $\Delta_Q = ||D - D'||_1 \leq 2$. By bounding the PDF difference, we effectively compress sensitivity to a fixed constant for all the queries. This allows formal privacy guarantees to be achieved over infinite data. Furthermore, the PDF representation maintains high utility for the analyses in most real-time applications.

2.5.2 0-DP Synopsis Generator

The synopsis generation in DPI is responsible for generating a sampled synopsis for each query $q \in Q_s$. Our key observation is that establishing a universal synopses pool can be done in a privacy-free (0-DP) and efficient way, enabling the system to privately answer a wide range of queries.

Establishing a Pool of PDFs. DPI utilizes a pioneering 0-DP synopsis generation process to navigate this challenge (AdaBoosting [40] cannot solve). This method constructs a set of PDFs that can universally emulate any PDF with minimal error (see detailed analysis in Appendix 2.10.6). These PDFs are then employed to generate answers for a set of queries, delivering approximated results without requiring access

to raw data. Various types of queries can be addressed using these PDFs, including but not limited to Point Queries, Range Queries, Aggregation Queries, Top-K Queries, Join Queries, and Correlation Queries [43, 44, 45, 46].

Quantization. The first step in the 0-DP synopsis generation process is quantization, which discretizes the values that PDFs can take. We define the precision level of quantization as p , representing one unit of probability (e.g., $p = 0.001$). By dividing the range of possible values into intervals of size p , we can express the quantized value of a continuous variable x as $\text{quantize}(x) = \text{round}(x/p) \cdot p$. This ensures that the values are discretized into intervals of width p .

Sampling Synopses with Multinomial Distribution. After quantization, we can generate all possible PDFs by distributing the available units of probability (denoted as $n_p = \frac{1}{p} + 1$) among the domain of interest, denoted as $\mathcal{K}$. Since n_p should be the same as the dataset size n , we can derive the value of p . Then, we can model this distribution using a multinomial PDF, denoted as $P(x)$, where x represents the random variable that takes on values from the domain $\mathcal{K}$. The multinomial PDF captures the probabilities of outcomes for this discrete random variable with k categories.

During each multinomial sampling, we treat each category equally and sample them uniformly for n_p times, leading to a generation of a synopsis (the distribution of n data points with k categories). To create the 0-DP synopses pool, we leverage the multinomial sampling. Let n be the number of samples drawn from the multinomial distribution and denote the sampled values as $(x_1, x_2, \dots, x_n)$, where each x_i

represents a possible outcome from the multinomial distribution defined by $P(x)$. We should repeat the multinomial sampling process for N times to get the pool of synopses $(P_1, P_2, \dots, P_N)$. The algorithm of 0-DP synopsis generator is shown in Appendix 2.10.6.

Next, for queries in Q_s^t , a set of synopses $\mathcal{A}_{Q_s^t}^t$ should be sampled according to the most recently updated sampling probability $P_{\mathcal{A}}(t)$. We will discuss how to update the sampling weights for different synopses in the DPI Boosting. The DPI Boosting algorithm iteratively improves the accuracy of DP queries on streaming data through a bi-directional reweighting process that favors challenging queries and high-quality synopses. By continually reweighting query and synopsis sampling probabilities based on estimated errors, DPI Boosting adapts the new coming data to optimize accuracy under DP constraints.

0-DP Guarantee for the Synopsis Generator. The synopsis generation process in DPI is uniquely designed to ensure that it does not consume any privacy budget since it operates without accessing to any raw data. Initially, the Synosis Generator invokes quantization to discretize potential values for any PDFs without using any raw data. Subsequently, through the procedure of multinomial synopses sampling, all potential PDFs are generated by distributing available units of probabilities across a defined domain of interest, bypassing the need for raw data access. The detailed steps and theoretical analysis are given in Appendix 2.10.6.

2.5.3 DPI Boosting (Bi-directional)

The DPI Boosting continuously evaluates and adjusts the importance of a subset of queries (“sampled queries”). It allows less accurate sampled queries to be reweighted higher, while the reweighting process is applied to focus more on accurate synopses for each sampled query. This bi-directional boosting process effectively captures the dynamic nature of data streams when answering queries. In this section, we will provide a comprehensive understanding of the DPI Boosting module and its role in enhancing the utility and privacy guarantees of the DPI.

Bi-directional Boosting. DPI Boosting significantly revises the AdaBoosting under DP constraints [40] to reweight the sampling weights. Specifically, we focus on the optimization techniques used to reweight the query and synopsis distribution during each round of DPI Boosting. The goal is to assign higher weights to poorly handled queries, and lower weights to well-handled queries, ultimately leading to enhanced accuracy and utility of the final synopsis.

Recall that we sample a subset of synopses $\mathcal{A}_{Q_s^t}^t$ for the sampled subset of queries Q_s^t . In the DPI Boosting, we should evaluate the difference between the true query result and the query result based on the synopsis. Our weight boosting differs from the AdaBoosting [40] from two aspects: (1) there is a subset of synopses instead of one synopsis to reweight the sampling probability, and (2) the reweighted sampling probability should be used both for queries in a pool and synopses in a pool. The details of DPI boosting are shown in Algorithm 1. Notice that the size of query pool Q may

be different from the size of synopsis pool $\mathcal{A}$ and it should not affect the reweighting process since normalization is applied for queries and synopses.

Algorithm 1: DPI Algorithm

Input: data stream S , privacy budget ϵ , synopsis error bound λ , error bound for overfitting μ , decay rate of decaying series ζ , query pool Q , synopsis pool $\mathcal{A}$ (Algorithm 3)

Output: private output streams $\mathcal{A}_{Q_s}^t$ (PDFs)

- 1 Initialize $P_Q(1), P_{\mathcal{A}}$ to uniform distribution
- 2 **foreach** S_t in Stream S **do**
 - /* Random Budget Allocation */
 - 3 $\epsilon_t = \text{RBA}(t)$ /* Algorithm 2 */
 - 4 $\epsilon_{Q,t} = \epsilon_{\mathcal{A},t} = \frac{\epsilon_t}{2}$ /* Q_s^t sampling */
 - 5 $Q_s^t \sim P_Q(t)$ /* $\mathcal{A}_{Q_s^t}^t$ sampling */
 - 6 $\mathcal{A}_{Q_s^t}^t \sim P_{\mathcal{A}}(t)$ /* DPI Boosting */
 - 7 **if** $|q(S_t) - \mathcal{A}_{Q_s^t}^t| < \lambda$ **then**
 - 8 $P_{\mathcal{A}}(t)[q] \leftarrow 1, P_Q(t) \leftarrow -1$
 - 9 **if** $|q(S_t) - \mathcal{A}_{Q_s^t}^t| \geq \lambda + \mu$ **then**
 - 10 $P_{\mathcal{A}}(t)[q] \leftarrow -1, P_Q(t) \leftarrow 1$
 - 11 **else**
 - 12 $P_{\mathcal{A}}(t)[q], P_Q(t) \leftarrow 1 - 2(|q(D) - \mathcal{A}_{Q_s^t}^t| - \lambda)/\mu$
 - 13 $P_{\mathcal{A}}(t+1) = \text{Normalize } P_{\mathcal{A}}(t)$
- 14 $u_{q,a} \leftarrow \exp(-\alpha_t \cdot \sum_{j=1}^t P_{\mathcal{A}}(j)[q])$, where $\alpha_t = \frac{1}{2} \ln[(1 + 2\eta_t)/(1 - 2\eta_t)]$ and η_t is the learning rate in time slot t
- 15 normalization factor $Z_t \leftarrow \sum_{q \in Q} u_{q,a}$
- 16 update $P_Q(t+1) = u_{q,a}/Z_t, P_{\mathcal{A}}(t)[q] = u_{q,a}/Z_t$
- 17 **return** $\mathcal{A}_{Q_s^t}^t$ as private output streams

Choice of Parameters in DPI Boosting The reweighting depends on the parameters λ , μ , and η , which all evaluate the accuracy of synopses. Note that λ denotes the error bound of the chosen synopses, while μ refers to an additional error bound that is

tolerated to avoid overfitting during DPI boosting. The parameter η is determined by the specified privacy budget ϵ . We find the optimal value of η_t for each time slot with the given ϵ to minimize the error bound of DPI framework under the privacy constraints. The details are presented in Section 2.6.2 (Theorem 4).

2.5.4 Random Budget Allocation (RBA)

As outlined in Algorithm 2 to mitigate the risk of *systematic depletion* over time, DPI employs a random selection process when defining ϵ_t (“Random Budget Allocation (RBA)”) rather than sequentially picking elements from the infinite series of potential budget values. It employs an exponential PDF with $p(\epsilon_t = x) = \exp(-\Lambda \cdot x)$, where Λ is an astronomically large number. This could generate very tiny samples, with the probability of obtaining such small values approaching 1.

Algorithm 2: Random Budget Allocation (RBA)

```

Data: available privacy budget space  $\mathcal{E}_t = \{\epsilon_1, \epsilon_2, \dots\}$ 
Data: tuning parameter  $\Lambda = 10^8$ 
/*  $\Lambda$  is a large number tunable to scale of remaining
   epsilon values in  $\mathcal{E}_t$  */
/* sample from exponential distribution */
1  $S \sim \text{Exp}(\Lambda)$ ;
/* find the closest remaining privacy budget to  $S$  */
2  $\epsilon_t \leftarrow \arg \min_{\epsilon \in \mathcal{E}_t} |\epsilon - S|$ ;
/* update available privacy budget space */
3  $\mathcal{E}_{t+1} \leftarrow \mathcal{E}_t \setminus \{\epsilon_t\}$ ;
/* return the selected privacy budget */
4 return  $\epsilon_t$ ;

```

After generating such very tiny values, they are mapped to the elements in the series $\mathcal{E}$ sampled from the exponential PDF. The closest one to the sampled value is

then returned, denoted as $\epsilon_t \leftarrow \arg \min_{\epsilon \in \mathcal{E}_t} |\epsilon - S|$. We note that since RBA generates samples without replacement in the series, they tend to get smaller and smaller over time. However, the sum of t independent samples from an exponential PDF with the rate parameter Λ follows a Gamma distribution with a shape parameter t and a rate parameter Λ . The mean (ν) of a Gamma distribution with shape parameter ω and rate parameter θ is given by $\frac{\omega}{\theta}$. In our case, for the sum of t exponential samples, given the shape parameter t and the rate parameter Λ , a mean (ν) of $\frac{t}{\Lambda}$ can be generated. Therefore, for any $t < \Lambda$, we can effectively preserve $\epsilon - \frac{t}{\Lambda}$. In practice, this bound must be much better as samples of the series are tending to be smaller than exponential PDF's samples due to sampling without replacement. An alternative RBA scheme based on *ranges* is given in Appendix 2.10.7.

Finally, as shown in Figure 2.2, each ϵ_t is further equally divided into $\epsilon_{Q,t}$ and $\epsilon_{\mathcal{A},t}$ for reweighting the query and synopses distribution, respectively. Specifically, it determines the learning rate that controls how aggressively the query sampling distribution is updated based on the empirical errors. Also, $\epsilon_{\mathcal{A},t}$ plays a key role in governing the privacy budget utilized when reweighting the synopses distribution $P_{\mathcal{A}}(t)$. By separating the budgets for querying and synopses sampling, DPI allows more flexible control to boost each of these two key components in the overall mechanism. Figure 2.12 plots the remaining budgets after t time slots and it clearly demonstrates that RBA is a thoughtful budget consumption strategy over the infinite disclosure.

DPI Framework. Algorithm 1 presents the details for DPI, which outputs PDF(s) in

each time slot to dynamic queries with ϵ -DP over infinite streams. W.l.o.g., we take the single-query case for Q_s^t as an example. A set of queries for Q_s^t only need to split the budget for multiple output PDFs.

2.6 Privacy and Utility Analysis

2.6.1 Privacy Loss after t Iterations

To understand the privacy guarantees of our DPI, we analyze the privacy budget consumed over successive disclosures. Theorem 1 shows how to calculate the total privacy budget ϵ by DPI after t iterations (e.g., time slots).

Theorem 1. *Consider a series of DPI mechanisms, each triggered with η_i^Q and $\eta_i^{\mathcal{A}}$. The overall privacy budget of DPI after t iterations is given by:*

$$\epsilon = \frac{4}{\mu} \sum_{i=1}^t \log \left[\left(\frac{1 + 2\eta_i^{\mathcal{A}}}{1 - 2\eta_i^{\mathcal{A}}} \right) \left(\frac{1 + 2\eta_i^Q}{1 - 2\eta_i^Q} \right) \right] \quad (2.2)$$

Proof. Each round of disclosure t results in the use of an additional privacy budget $D_\infty(\mathcal{D}_t || \mathcal{D}'_t)$. Dwork et al. [40] showed that, under the assumption that synopses are known, this quantity is $\frac{2}{\mu} \log \left(\frac{1+2\eta_t}{1-2\eta_t} \right)$ privacy budget. Briefly, for each query $q \in Q$, let $d_q^t = |q(x) - A_t(q)|$ and $d_q^{t'} = |q(x') - A_t(q)|$. It follows that $|d_q^t - d_q^{t'}| \leq \Delta_Q$. This implies $|P_{\mathcal{A}}(t)[q] - P'_{\mathcal{A}}(t)[q]| \leq \frac{2\Delta_Q}{\mu}$, and thus:

$$e^{-2\alpha \cdot t \cdot \frac{\Delta_Q}{\mu}} \leq \frac{u_t, q}{u'_t, q} \leq e^{2\alpha \cdot t \cdot \frac{\Delta_Q}{\mu}}$$

Define normalization factor for two adjacent database as $Z_t = \sum_{q \in Q} u_t, q$ and $Z'_t = \sum_{q \in Q} u'_t, q$, we have:

$$e^{-2\alpha \cdot t \cdot \frac{\Delta_Q}{\mu}} \leq \frac{Z_t}{Z'_t} \leq e^{2\alpha \cdot t \cdot \frac{\Delta_Q}{\mu}}$$

The claim follows from the above because $\mathcal{D}_{t+1}[q] = \frac{u_t, q}{Z_t}$, $\mathcal{D}'_{t+1}[q] = \frac{u'_t, q}{Z'_t}$ and replacing $\alpha = \frac{1}{2} \log \left(\frac{1+2\eta}{1-2\eta} \right)$.

Therefore, in DPI, since we update two PDFs $P_Q(t)$ and $P_{\mathcal{A}}(t, q)$, each of these two PDFs will consume a privacy budget of $\frac{4}{\mu} \sum_{i=1}^t \log \left(\frac{1+2\eta_i^{\mathcal{A}}}{1-2\eta_i^{\mathcal{A}}} \right)$ and $\frac{4}{\mu} \sum_{i=1}^t \log \left(\frac{1+2\eta_i^Q}{1-2\eta_i^Q} \right)$, respectively. This completes the proof. $\square$

The convergence of the privacy bound in the DPI can be analyzed through the total privacy budget ϵ utilized after t iterations. Next, assuming the privacy budet is divided equally between the two reweighting mechanisms, i.e., $\epsilon_Q = \epsilon_{\mathcal{A}} = \frac{\epsilon}{2}$, the optimal series can be derived.

2.6.2 Total Privacy and Utility Loss

According to Dwork et al. [40] (Claim 4.3), after t rounds of DP with AdaBoosting, while the mechanism produces λ -inaccurate results, the loss probability $\mathcal{L}(t)$, derived via the cardinality of a subset of queries ($Q_{\text{bad}} \subset Q$), is upper-bounded by $\left(\sqrt{1 - 4\eta^2} \cdot t \right) \cdot |Q|$. A very similar observation to the proof but utilizing different η_i in each iteration i (from a converging series) will lead to $\sum_{i=1}^t \mathcal{L}_i = \sum_{i=1}^t \frac{1}{2} \log [(1 + 2\eta_i)(1 - 2\eta_i)]$.

Theorem 2 (Proof in Appendix 2.10.2). *Consider a series of DPI mechanisms, and each triggered with η_i^Q and η_i^A . The overall error of DPI, represented as $\sum_{i=1}^t \mathcal{L}_i = [\mathbb{P}(|\mathcal{A}(q, t) - q(S_1, S_2, \dots, S_t)| > \lambda + \mu)]$, is upper-bounded by:*

$$\frac{1}{4} \sum_{i=1}^{\infty} \log \left[(1 - 4(\eta_i^Q)^2) \cdot (1 - 4(\eta_i^A)^2) \right] \quad (2.3)$$

Theorems 2.6.1 and 2 allow us to formulate an optimization problem for calculating the parameter η to maximize utility under a total privacy constraint.

$$\begin{aligned} \min_{\eta_i, i \in [\mathbb{N}]} \quad & \frac{1}{4} \sum_{i=1}^{N \rightarrow \infty} \log \left[(1 - 4(\eta_i^Q)^2) (1 - 4(\eta_i^A)^2) \right] \\ \text{w.r.t.} \quad & \frac{4}{\mu} \sum_{i=1}^{N \rightarrow \infty} \log \left[\left(\frac{1+2\eta_i^A}{1-2\eta_i^A} \right) \left(\frac{1+2\eta_i^Q}{1-2\eta_i^Q} \right) \right] = \epsilon \end{aligned} \quad (2.4)$$

To derive the optimal series, we model the problem using the continuous representations $\eta(x)$ and $\mathcal{L}(\eta(x))$. This simplifies the analysis using the integral calculus.

Definition 4. *Continuous functions $\eta : \mathbb{R}^+ \rightarrow (0, 0.5)$ and $\mathcal{L} : \eta(x) \rightarrow \frac{1}{2} \log(1 - 2\eta(x)) (1 + 2\eta(x))$ are defined to be, the continuous representation of η_i , and its corresponding continuous loss function obtained by interpolating a continuous function over the sequences of $\langle \eta_i^A, \eta_i^Q \rangle$.*

Moreover, we make two reasonable assumptions:

1. The interpolation error to be negligible.

2. The function $\eta(x)$ is twice continuously differentiable, indicating that it belongs to the class C^2 .

Total Privacy over Infinite Iterations. Under this continuous model, we derive key relationships between the total privacy budget ϵ and the overall utility loss L in Remark 2.6.1 and Theorem 3.

Remark 2.6.1 (Proof in Appendix 2.10.3). *For the overall probability loss $L = \int_1^\infty \mathcal{L}(x) dx$, there exist constants ζ, m, M such that*

$$L \geq \frac{1}{4\zeta} \left(Li_2(M^2) - Li_2(m^2) + 4Li_2(m) - 4Li_2(M) \right) \quad (2.5)$$

where $Li_2(\cdot)$ is the poly-logarithm Spence's function.

Theorem 3 (Proof in Appendix 2.10.4). *The total privacy bound of DPI over infinite time slots is*

$$\epsilon \geq \frac{2\Delta_Q}{\mu\zeta} \left(Li_2(M^2) - Li_2(m^2) \right) \quad (2.6)$$

Since the $Li_2(M^2)$ is the poly-logarithms of the Spence's function, the overall privacy loss is growing logarithmically. Thus, Theorem 3 signifies that the privacy loss is converging and effectively bounded in DPI.

Utility Loss We now present our final result: a series that maximizes the accuracy of PDF estimation in DPI.

Theorem 4 (Proof in Appendix 2.10.5). *Given total privacy budget ϵ , DPI mechanism*

achieves an optimal utility (in a probability loss sense) with the series $\eta_t = \frac{1}{2} \cdot \frac{e^X - 1}{e^X + 1}$:

$$|L| \geq \frac{1}{4\zeta} \left| \frac{\epsilon}{C} - \frac{2\pi^2}{3} + 4Li_2 \left(\sqrt{Li_2^{-1} \left(\frac{\pi^2}{6} - \frac{\epsilon}{C} \right)} \right) \right|. \quad (2.7)$$

where $Li_2^{-1}(\cdot)$ represents the inverse of poly-logarithm of the Spence's function,

$$X = \frac{1 - (Li_2^{-1}(\frac{\pi^2}{6} - \frac{\epsilon}{C}))^t}{t^2|\zeta|} \quad |\zeta| = \sup_{\{t=2,3,\dots\}} \{e^{|\eta_{k-1} - \eta_k|}\} \text{ and } C = \frac{\Delta_Q}{\mu|\zeta|}.$$

DPI leverages the inverse of the poly-logarithm of the Spence's function which can be represented as a converging infinite series to guarantee the loss bound. As t grows, the utility loss will converge. The converging series provides a tight utility guarantee in Theorem 4.

Summary. Our theoretical analyses have shown that DPI can provide strong cumulative privacy and high utility for infinite data streams. Specifically, we prove DPI's overall privacy budget grows logarithmically with iterations (Theorem 3). This allows continuous operation under a fixed total privacy budget ϵ . We also derive an upper bound on DPI's total error over time (Theorem 2), quantifying its utility. Furthermore, we obtain the optimal strategy for η_t across iterations to maximize utility under ϵ (Theorem 4). These results show that DPI provably converges to $O(\epsilon)$ error for answering pre-defined queries under privacy budget ϵ .

2.7 Experimental Evaluations

2.7.1 Experimental Setting

Experimental Datasets. We conduct all the experiments on three real-world and some synthetic streaming datasets to evaluate the performance of our DPI. Table 2.2 shows the characteristics of the datasets (the time period is partitioned into specific numbers of equal-length time slots).

Table 2.2: Characteristics of datasets (after pre-processing)

Dataset	User #	Domain Size	Slot #	Period
COVID-19	48,925	110	1,000	6 month
Network Traffic	5,074,413	65,534	2,700	1 hour
USDA Production	1,695,038	111,989	2,700	1960-2022

COVID-19 Dataset [47] provides a chronological sequence of COVID-19 confirmed cases, deaths, and recoveries, allowing researchers to analyze the trend and pattern of the pandemic over time. Moreover, the attributes are disaggregated by 110 countries and subregions aiding in regional and comparative studies, enhancing the understanding of the geographical variation and spread of the virus.

Network Traffic Dataset [48] is used to study a variety of real-world DDoS attacks. This dataset contains network traffic such as source and destination IP addresses, etc. 65,534 source IP addresses and their network traffic records were extracted for our experiments.

USDA Production Supply and Distribution Dataset [49] collects the agricultural production data for 300 commodities with over 100 attributes, including 111,989 cate-

gories of 1,695,038 production records from 1960 to 2022.




Fig. 2.3: Average MSE and KL divergence vs. ϵ for different values of the λ hyperparameter. The left column shows results for accumulative data disclosure, and the right column shows results for instantaneous data disclosure.

Applications for Evaluation. We next introduce three real-world applications for evaluating the performance of DPI.

Statistical Queries. Since our framework DPI could release the data distribution in each time slot, we first evaluate the utility for querying the item distribution: the difference between the original data’s item distribution and the DPI output using the




Fig. 2.4: Average MSE and KL divergence vs. ϵ for different values of the μ hyperparameter. The left column shows results for accumulative data disclosure, and the right column shows results for instantaneous data disclosure.





(a) MSE vs ϵ (b) MSE vs ϵ (c) KL divergence vs ϵ (d) KL divergence vs ϵ

Fig. 2.5: Average MSE and KL divergence vs. ϵ for different values of the ζ hyperparameter. The left column shows results for accumulative data disclosure, and the right column shows results for instantaneous data disclosure.

MSE and KL divergence metrics.

Anomaly Detection. In addition to statistical queries, the data distribution released by DPI facilitates various downstream analyses [50, 51, 52, 53, 54, 55, 56, 57]. In our experiments, we conduct experiments on anomaly detection (*most existing methods cannot support such analysis*) by adopting a distribution-based anomaly detection method [58] to detect the anomalies with the DPI output and evaluate the accuracy.

Recommender System. DPI also enhances recommender systems by providing privacy guarantees without compromising accuracy. Singular Value Decomposition (SVD) algorithms [59, 60] used in recommender systems can investigate the relationships between items by using the distribution of users and items. In our experiments, DPI enables precise suggestions by providing differentially private data distributions to the SVD algorithm.

Benchmarks. We have shown that existing methods violate the privacy bound ϵ after a few time slots (see Appendix 2.10.9). They are incomparable with DPI in our infinite settings due to their limitations (e.g., unbounded privacy, sensitivity assumption, event-level differential privacy).

2.7.2 Performance on Hyperparameters

The utility of DPI is influenced by parameters such as ϵ , λ (error bound for synopsis), μ (error bound for avoiding overfitting), and ζ (decay rate of the decaying series), which are all related to the reweighting phase. To explore how they affect the utility in real-time

disclosure, we choose the COVID-19 dataset to observe the utility trends with varying hyperparameters. The privacy requirement ϵ ranges from 0.1 to 10, and ζ varies from 0.1 to 0.5. λ is set in the range 0.05 to 0.5 and μ varies from 0.05 to 0.3.

We show the MSE and KL divergence regarding ϵ for 1,000 time slots in Figure 2.3, 2.4, 2.5. Figures 2.3(a), 2.3(c), 2.4(a), 2.4(c), 2.5(a), and 2.5(c) are the MSE and KL divergence with different λ , μ , and ζ values for accumulative data distribution disclosure (from time slot 1 to 1,000). Figures 2.3(b), 2.3(d), 2.4(b), 2.4(d), 2.5(b), and 2.5(d) are the MSE and KL divergence with different λ , μ , and ζ values for instantaneous data distribution disclosure (time slot 1,000).

First, Figure 2.3, 2.4, 2.5 demonstrate the MSE and KL divergence by varying ϵ for accumulative data distribution disclosure and instantaneous data distribution disclosure. With an increasing ϵ , the MSE and KL divergence both decrease, demonstrating that the DPI output is closer to the true distribution (privacy/utility tradeoff). Second, Figures 2.3(c) and 2.3(d) show the KL divergence decreases faster when λ is very small (purple line) since the parameter λ evaluates the error of query based on the sampled synopsis and a small λ means that the synopsis is more accurate. Thus, the output results of DPI tend to be accurate. Furthermore, the smaller additional error bound μ (for avoiding overfitting) indicating the confidence for the universal synopsis output also affects the performance of DPI. Figures 2.4(c) and 2.4(d) demonstrate that we can have good utility with small KL divergence given small μ (e.g., 0.05 or 0.1). In our settings, μ is set to be less than half of λ . Finally, Figures 2.5(c) and 2.5(d) validate that small ζ (e.g., 0.05,

0.1, or 0.2) returns higher accuracy (since smaller ζ produces larger η , leading to more accurate synopses).

Similarly, we can draw similar observations from Figures 2.3(a) and 2.3(b), 2.4(a) and 2.4(b), 2.5(a) and 2.5(b) (with the MSE metric). This confirms that the effect of different parameters on DPI aligns with our theoretical studies.

2.7.3 Utility Evaluation

Statistical Queries. We consider two queries: sum and mean. We evaluate the KL divergence of total 1,000 time slots with a varying ϵ and present the results in Figures 2.6(a) and 2.6(b). The results show that the utility of DPI improves as ϵ increases. Even with a small ϵ , DPI ensures low MSE values, e.g., ~ 0.12 when $\epsilon = 0.1$. In addition, we also evaluate the MSE and KL divergence in each of the 1,000 time slots, as shown in Figures 2.6(c) and 2.6(d). The MSE values for all statistical queries are quite low, up to 0.16.

Anomaly Detection. Another application focuses on anomaly detection in the Network Traffic dataset. In this task, we will apply the Histogram-based Outlier Score (HBOS) [58] (equal to a distribution-based anomaly detection technique) and isolation forest [61] to identify anomalies. The HBOS value reflects the rarity of data in the learned distribution and the extent of deviation of certain features from this distribution. The network monitor can quickly identify and report any anomalies from incoming packets, along with their severity, to the respective tenants.




Fig. 2.6: Average MSE and KL divergence of statistical queries for ϵ and t across 1,000 time slots on COVID-19 dataset. (a), (b): Average MSE and KL divergence with varying ϵ . (c), (d): Average MSE and KL divergence over t when ϵ is 2.

We conduct experiments on the Network Traffic dataset and define the items with a score over a threshold as anomalies. DPI is expected to detect the anomalies from its real-time disclosures. We adopt Precision and Recall as the evaluation metrics compared to the non-private results, as shown in Figure 2.7. The Precision and Recall of both HBOS and Isolation Forest increase linearly with a growing ϵ . Moreover, anomaly detection can be accurately performed by dynamically tracking the new data





(a) Precision vs ϵ (b) Recall vs ϵ (c) Precision (moving average) vs t (d) Recall (moving average) vs t

Fig. 2.7: Average Precision and Recall for anomaly detection over 2,700 time slots on Network Traffic dataset. (a), (b): anomaly detection with different ϵ . (c), (d): anomaly detection in time slot t ($\epsilon = 2$). Due to unsupervised learning on unlabeled data, non-private results are treated as the ground truth.

over a large number of time slots (as observed from the fluctuated Precision and Recall scores). Finally, anomaly detection over all the time slots (accumulative) tends to show relatively better results than that in specific time slots (instantaneous).

Recommender System. For recommender systems, we utilize the SVD algorithm [59, 60], which capitalizes on data distribution for delivering recommendations by efficiently managing dimensionality reduction and highlighting essential features. It




Fig. 2.8: Average Precision and Recall for recommender system over 2,700 time slots on USDA dataset. (a), (b): Precision and Recall vs ϵ . (c), (d): Precision and Recall vs time slot t ($\epsilon = 2$). Due to unsupervised learning on unlabeled data, non-private results are treated as the ground truth.

decomposes a user-item matrix into three constituent matrices, capturing the interaction between users and items. By reducing the dimensionality, SVD effectively uncovers latent features indicative of user preferences. Then, SVD leverages the data distribution to recommend items that align with users' interests based on their past behavior. We also test the utility with the K-means [62, 63] algorithm. Figure 2.8 shows that both Precision and Recall increase as ϵ increases. They can be very high by considering the

non-private results as the ground truth (e.g., 90%+).

2.7.4 Highly Dynamic Data Distributions

To evaluate DPI on highly dynamic data streams, we generated a synthetic dataset over 3,000 time slots. For each time slot, we generate a synthetic dataset by initializing 100 items in each dataset and sampling their counts with the Gaussian distributions (mean 100 and variance randomly chosen from 1, 4, 9, 16, 25, $\dots$, 100). Figure 2.11 shows the true values in the synthetic datasets and the DPI outputs under different ϵ (since there are numerous values in each time slot, we plot the median values in the figure). Then, we can observe that the DPI outputs are close to the true values by well preserving the original data distribution.

Furthermore, Figure 2.9 shows the MSE and KL divergence of DPI outputs over time under different ϵ . From the plots, we observe the values of MSE and KL divergence are very small even in case of small ϵ (e.g., 0.5 and 1). Moreover, although the distribution changes significantly over time, the MSE and KL divergence of DPI are very stable. The boxplots in Figure 2.10 show the lower quantile (LQ), upper quantile (UQ), and median (M) of the DPI output densities. The output density here refers to the distribution of DPI outputs in terms of their frequency and spread. The spread of the distributions remains stable over time.

Finally, we conduct experiments on 20 more synthetic datasets and different $\epsilon \in [0.5, 10]$ (see Figure 2.17 in Appendix 2.10.10). First, we generate 10 distinct




Fig. 2.9: Average MSE and KL divergence over 1,000 time slots. (a) (b): MSE vs t (varying ϵ). (c) (d): KL divergence vs t (varying ϵ).

synthetic datasets by initializing 1,000 items in each dataset and sample their counts with the Gaussian distributions (mean 1,000 and variances 1, 4, 9, 16, 25, $\dots$, 100 for 10 datasets, respectively). Second, we generate another 10 distinct synthetic datasets with the domain size 10,000 and similar settings. Figure 2.17 further proves the stability and effectiveness of DPI across varying data distributions and domain sizes.





(a) Density vs ϵ (b) Density vs ϵ (c) Density vs t (d) Density vs t

Fig. 2.10: Distribution spread of DPI output on the synthetic dataset with changing distribution. (a) (b): DPI output vs ϵ . (c) (d): DPI output vs t . The boxplots show the lower quantile (LQ), upper quantile (UQ), and median (M) of the DPI output densities. The spread of the distributions remains stable.

2.8 Discussion

Boosting and DPI. Boosting is a powerful technique in machine learning that combines multiple weak models to produce a stronger and more robust model. In the case of DPI, boosting is utilized to enhance privacy and accuracy through the continuously updated streaming data in each dynamic batch. In particular, we only run T boosting rounds in




Fig. 2.11: Real-time DPI outputs and true data (median).

the first batch to get a good sampler PDF. Subsequently, DPI will continuously update the synopses sampling distribution in each time slot to the underlying data distribution with privacy preservation. Upon that, the efficiency of DPI can be significantly improved. DPI can also be modified to utilize boosting technique for each round. In this way, the accuracy can be better but cost more privacy budget.

Complexity Analysis. The complexity analysis for DPI focuses on the primary computational components: the Boosting Algorithm, and the 0-DP Synopsis Generation algorithm. In DPI, the Boosting Algorithm updates and normalizes distribution synopses with a complexity of $O(n)$, and the 0-DP Synopsis Generation algorithm generates synopses through probabilistic sampling also with a complexity of $O(n)$. The overall time complexity $O(n)$ also aligns well with the standard streaming algorithms (over time series). Moreover, DPI consumes memory proportional to the domain size (which is tolerable even for large domains).

No Numerical Overflow for Extremely Tiny Budgets. The randomization in DPI

is based on “DP-Boosting”, which effectively translates extremely tiny budgets into negligible or even no reweighting. Thus, different from noise-additive DP mechanisms (e.g., Laplace and Gaussian), extremely tiny budgets in DPI do not entail any numerical overflow.

Limitations. First, DPI presents a new privacy-utility tradeoff compared to traditional noise-additive mechanisms. Utility loss manifests as ineffective tracking of new stream data, potentially in extreme cases of dynamic data over extremely long periods, where privacy budget depletes, possibly necessitating DPI system restart for optimal utility with a renewed budget. However, DPI may not necessarily need to restart in practice for two reasons: (1) our random budget allocation ensures that budgets are not depleted even after a large number of time slots (Figure 2.12 demonstrates that budgets are adequate after 500 time slots), and (2) given extremely tiny privacy budgets, DPI maintains low MSE and KL divergence for different data distributions over extremely long periods. Figure 2.13 shows that the results are still stable even after 200,000 time slots (no need to restart), evaluated on the synthetic data generated in Section 2.7.4. Indeed, we cannot find extreme cases of the low utility of DPI in our extensive empirical studies.

Second, DPI outputs data as a series of PDFs instead of count histograms, suitable for most downstream analyses. Nevertheless, counts can be estimated from these PDFs by discretizing the domain into bins, calculating each bin’s probability mass, and scaling these probabilities by the total data size. For instance, a 0.02 probability in a bin for ages 15-20 in a million-record dataset implies an estimated 20,000 individuals in that age




Fig. 2.12: Remaining budget for DPI with the random privacy budget allocation: (a) $\epsilon = 2$, and (b) $\epsilon = 10$.




Fig. 2.13: Evaluation of DPI on 200,000 time slots on synthetic datasets (a) KL divergence and (b) MSE: domain size 10,000.

group. To privately disclose counts in certain applications, DPI only needs to assume the release of the non-private total data size at time slot t .

2.9 Related Works

Data Streaming with DP. Since the early studies on the private data streaming [64, 65],

DP models were proposed to protect the streaming data. Li et al. [66] introduce a

framework that enables differentially private data release from data streams by allowing controlled delays, thereby enhancing utility while ensuring privacy. Ebadi et al. [67] proposed a personalized DP for dynamically adding records to the database. Meanwhile, Liu et al. [68] also proposed personalized DP with the weighted posterior sampling to reduce the extra Gaussian noise to the parameter space. To further apply DP to streaming data, many works [23, 24, 25, 26, 30, 69, 70] focus on continuously publishing statistics computed over events. Chan et al. [24] propose several methods that can handle binary streams of different users over potential infinite streams. Several works [25, 27, 28] follow a similar idea of applying partition algorithms to handle numeric values in the streaming data. However, event-level privacy is not strong enough to prevent sensitive data. Later, Kellaris et al. [29] proposed the w -event privacy to protect event sequence occurrence in w timestamps. This notion converges to user-level privacy when w is set to infinite, but the noise would be unbounded. Aiming to achieve user-level DP, FAST [31] was proposed to release real-time statistics without full DP infinite data stream.

Query Boosting with DP. Boosting has been widely used for improving the accuracy of learning algorithms [71]. There are also many algorithms that boost the DP results. In [72], they combined private learning with DP, using learning theory to comprehend database probabilities, potentially enhancing DP's distortion reduction. For general counting queries, Dwork et al. proposed a base synopsis generator [42], and also constructs an appropriate base synopsis generator for any set of low-sensitivity queries (not just counting queries) [40]. These well-designed synopsis generators can significantly

boost the utility of DP results. Recall that our synopsis generator significantly differs from them to support infinite data streaming.

Data Streaming with LDP. LDP [38] involves users perturbing their data before sending it to an untrusted server for aggregation, posing challenges in privacy accumulation for streaming data with limited utility (e.g., a sequence of locations [73]). Joseph et al. [74] proposed a framework for continuous data sharing under LDP, consuming privacy budget based on distribution changes rather than collection periods to reduce the overall privacy expenditure. Wang et al. [28] proposed an algorithm using the Exponential mechanism with a quality function to publish a stream of real-valued data under both centralized and LDP. Li et al. [75] proposed an LDP approach to heavy hitter detection on data streams with bounded memory. These works can provide event-level privacy protection. Bao et al. [76] proposed a novel correlated Gaussian mechanism for (ϵ, δ) -LDP on streaming data aggregation. However, CGM needs to update the privacy budget periodically. Compared to these, Ren et al. [77] then proposed a population division framework that not only avoids the high sensitivity of LDP noise to the budget division but also requires less communication. However, it can only provide the w -event privacy.

Applications. Differential privacy (DP) has been the de facto rigorous privacy solution for learning algorithms [78, 79, 80, 81, 82, 83, 84, 85]. McSherry et al. [86] applied DP to provide personal recommendations. Chen et al. [87] proposed to publish the rating matrix of the source domain with DP guarantee. Feng et al. [88] proposed a topic privacy-relevance parameter method for top- k recommendations. Okada et al.

[89] used three queries on statistical aggregation on outliers to detect the occurrence of anomalous situations with differential privacy guarantee.

2.10 Proof and Further Discussions

2.10.1 DP with AdaBoosting [40]




Fig. 2.14: The architecture of DP AdaBoosting [40]

2.10.2 Proofs

Proof of Theorem 2

As stated in Claim 4.3 in [40], the probability of inaccuracy produced by a DP boosting algorithm can be described in terms of (η, β) -base learners. In DPI, we can analogously

consider $(\eta_i^{\mathcal{A}(q,t)}, \eta_i^Q, \beta)$ -base learners at each iteration i .

After t rounds of DP boosting, the probability of inaccuracy in the DPI framework is upper-bounded by $\prod_{i=1}^t \sqrt{\left[(1 - 4(\eta_i^Q)^2) (1 - 4(\eta_i^{\mathcal{A}})^2) \right]}$. By taking the natural log of this bound, a version similar to the privacy bound (the sum of natural logs) is obtained. Specifically,

$$\begin{aligned} \log [\mathbb{P}(\text{bad outcome})] &\leq \\ \frac{1}{2} \sum_i &= 1^{N \rightarrow \infty} \log (1 - 2\eta_i) (1 + 2\eta_i) \end{aligned}$$

Thus, this completes the proof. $\square$

2.10.3 Proof of Remark 2.6.1

Our construction involves a decreasing sequence of privacy budgets, thus proving that η_i must also decrease. We assume $\eta(x) < 0.5$ is a converging function of at least an exponential order ($\exp(-\zeta \cdot x)$). Then, we have

$$\eta'(x) \geq -\zeta \eta(x), \quad \forall x \in \mathbb{R}^+ \quad (2.8)$$

where $\zeta = \left| \inf_{x \in [1, \infty)} \frac{d\eta(x)}{dx} \right|$. By replacing $2\eta(x)$ with u , and for constants $m = \min_{x \geq 1}(\eta(x))$ and

$M = \max_{x \geq 1}(\eta(x))$, we obtain

$$|L| \geq -\frac{1}{4} \int_m^M \frac{\log[(1+u)(1-u)]}{\zeta u} du$$

but given that the Spence's function [90] is defined as $Li_2(z) = -\int_0^z \frac{\log[(1-u)]}{u} du$, for $|z| \leq 1$, we have

$$|L| \geq \frac{1}{4\zeta} \left(Li_2(M^2) - Li_2(m^2) + 4Li_2(m) - 4Li_2(M) \right)$$

Here, we have utilized the following two results:

$$\int_0^z \frac{\log[(1+u)]}{u} du = -Li_2(-z) \quad (2.9)$$

$$Li_2(z) + Li_2(-z) = \frac{1}{2} Li_2(z^2)$$

Thus, this completes the proof. $\square$

2.10.4 Proof of Theorem 3

According to Equation 2.4, the continuous format of DPI utilizes the following privacy budget.

$$\epsilon = \frac{4}{\mu} \int_0^\infty \log \left[\frac{1 + 2\eta(x)}{1 - 2\eta(x)} \right] dx,$$

By following a similar process as that described in the Remark 2.6.1, we will have

$$\epsilon \geq \frac{4}{\zeta\mu} \int_m^M \frac{\log \left[\frac{1+u}{1-u} \right]}{u} dx.$$

Finally, by utilizing the two results in Equation 2.9, we have

$$\epsilon \geq \frac{2}{\mu\zeta} \left(Li_2(M^2) - Li_2(m^2) \right).$$

Thus, this completes the proof. $\square$

2.10.5 Proof of Theorem 4

By establishing lower bounds for both privacy and accuracy loss, we can determine the optimal series. Thus, our optimization problem can be succinctly stated as follows.

$$\begin{aligned} \min_{0 \leq m < M \leq 1} \frac{1}{8\zeta} (Li_2(M^2) - Li_2(m^2) + 4Li_2(m) - 4Li_2(M)) \\ \text{w.r.t. } \frac{2}{\mu\zeta} (Li_2(M^2) - Li_2(m^2)) = \epsilon, \end{aligned} \quad (2.10)$$

but Spence's function has the following infinite series form:

$$Li(z) = \sum_{t=1}^{\infty} \frac{z^t}{t^2}$$

Therefore, we have $Li_2(M) - Li_2(m) = \sum_{t=1}^{\infty} \frac{M^t - m^t}{t^2}$. Due to

$$Li_2(M^2) - Li_2(m^2) = \sum_{t=1}^{\infty} \frac{M^{2t} - m^{2t}}{t^2} = O(\epsilon), \quad (2.11)$$

However, we note that minimizing *probability loss* is equivalent to if we maximize $m^t + M^t$. This is true because our objective function has two opposite sign terms $A = [Li_2(M^2) - Li_2(m^2)] > 0$ and $B = 4Li_2(m) - 4Li_2(M) < 0$, where B appears in A , and minimization is translated into maximizing $(m^t + M^t)$ (using the sequence of *mean theroems*). Since Spence's function is strictly increasing, maximizing terms $m^t + M^t, \forall c$ means that $M = 1$.

Thus, when $M = 1 \Rightarrow C (Li(1) - Li(m^2)) = \epsilon$, we have

$$\begin{aligned}
 &\Rightarrow Li(m^2) = Li(1) - \frac{\epsilon}{C} = \frac{\pi^2}{6} - \frac{\epsilon}{C}, \quad C = \frac{2}{\mu\zeta} \\
 &\Rightarrow m^2 = Li_2^{-1} \left(\frac{\pi^2}{6} - \frac{\epsilon}{C} \right) \\
 &\Rightarrow m = \sqrt{Li_2^{-1} \left(\frac{\pi^2}{6} - \frac{\epsilon}{C} \right)} \tag{2.12}
 \end{aligned}$$

Since each term in the series representation of the overall ϵ (Equation 2.11) has to be equal to the budget spent over the same index's disclosure, we have

$$\begin{aligned}
 &\frac{2}{\mu\zeta} \cdot \sum_{t=1}^{\infty} \frac{M^{2t} - m^{2t}}{t^2} = \frac{4}{\mu} \log \left(\frac{1 + 2\eta_t}{1 - 2\eta_t} \right) \\
 &\Rightarrow e^{\frac{1 - m^{2t}}{2\zeta t^2}} = \frac{1 + 2\eta_t}{1 - 2\eta_t} \\
 &\Rightarrow \eta_t = \frac{\left[e^{\frac{1 - (Li_2^{-1}(\frac{\pi^2}{6} - \frac{\epsilon}{C}))^t}{t^2 \zeta}} - 1 \right]}{2 \cdot \left[e^{\frac{1 - (Li_2^{-1}(\frac{\pi^2}{6} - \frac{\epsilon}{C}))^t}{t^2 \zeta}} + 1 \right]}
 \end{aligned}$$

Thus, this completes the proof. □

2.10.6 Approximation of the Universe of PDFs

When n is sufficiently large, we can employ an asymptotic approximation to analyze the behavior.

This approximation corresponds to the normal distribution limit of the Binomial distribution.

For large n , we have: $H \approx k - \frac{1}{2} \ln(2\pi ne) + \frac{1}{2} \sum_{i=1}^k \ln p_i$, where the natural logarithm is

Algorithm 3: 0-DP Synopsis Generator

Input: streaming dataset D , size of dataset n , sample size N

Output: the universe of synopses $\mathcal{A}$

- 1 Initialize the quantization precision $p \in (0, 1)$ to make the corresponding precision level $n_p = \frac{1}{p} + 1$ equal to dataset size n
 - 2 $k \leftarrow$ distinct value number of data D
 - 3 Initialize the synopses pool $\mathcal{A}$ as empty
 - 4 **for** $i \in N$ **do**
 - 5 Set equal sampling probability for each distinct value in the domain
 $Prob = \frac{1}{k}$
 /* Sampling n times from k distinct values with
 $Prob = \frac{1}{k}$ and have a number of times each distinct
 value show among n times */
 - 6 $t \leftarrow \text{Multinomial}(n, k, Prob)$
 /* Sum of P_i is 1 */
 - 7 $P_i \leftarrow t * p$
 - 8 $\mathcal{A}.\text{append}(P_i)$
 - 9 **return** the universe of synopses $\mathcal{A}$
-

used for entropy calculation. Note that this approximation is not intended to provide an upper or lower bound, but rather an approximate form for large n . In the special case where the p_i are uniform, the expression simplifies to:

$$H \approx k - \frac{1}{2} \ln(2\pi n e) - \frac{k}{2} \ln k$$

While this is not an exact expression, it yields a highly accurate estimate of H when n is large.

Lemma 1. Let $H(\mathcal{A})$ denote the entropy of the constructed pool of synopses $\mathcal{A}$, and let H_{ideal} denote the entropy of the ideal multinomial PDF. As the number of samples drawn from the multinomial distribution, denoted as N , approaches infinity, the ratio $\frac{H(\mathcal{A})}{H_{ideal}}$ converges to 1.

Proof. This proof is based on the concept that as the number of samples approaches infinity, the constructed pool of synopses becomes more representative of the ideal multinomial PDF. By constructing $\mathcal{A}$ from the multinomial distribution, which is designed to distribute the available probability units uniformly, the synopses generated cover a comprehensive range of possible outcomes. As a result, the entropy of $\mathcal{A}$ approaches the entropy of the ideal multinomial PDF. Therefore, the ratio $\frac{H(\mathcal{A})}{H_{\text{ideal}}}$ converges to 1. $\square$

Theorem 5. *Let $H(\mathcal{A}; n_p, k)$ denote the entropy of the constructed pool of synopses $\mathcal{A}$, given a precision level n_p (equivalent to $1/p$, where p is the quantization precision) and a domain size k . For sufficiently large values of N (number of samples drawn), $H(\mathcal{A}; n_p, k)$ approximates the entropy expression for the ideal multinomial PDF.*

Proof. The proof of Theorem 5 involves comparing the entropy expression for the ideal multinomial PDF, denoted as H_{ideal} , with $H(\mathcal{A}; n, k)$. For a large value of n , the entropy of a heterogeneous multinomial PDF can be approximated as:

$$H_{\text{approx}}(N, k) \approx \frac{k-1}{2} \ln(2\pi N e) - \frac{k}{2} \ln(k), \quad (2.13)$$

where N represents the number of samples drawn from the multinomial distribution [91].

By examining the entropy expression for the ideal multinomial PDF, which depends on k , and comparing it with $H(\mathcal{A}; n_p, k)$ for different values of n and k in the constructed pool of synopses, we can establish a relationship between the two. Specifically, for sufficiently large values of N , $H(\mathcal{A}; n_p, k)$ approaches $H_{\text{approx}}(N, k)$, indicating that the constructed pool of synopses accurately represents the ideal multinomial PDF. Figure 2.15 shows the information theory of 0-DP Synopsis Generation. $\square$




Fig. 2.15: Information theory of 0-DP synopsis generation.

Table 2.3: Expanded sensitivity (Δ) and privacy loss (actual DP) for the median query of packets in each time slot (0.5 seconds) on a network traffic dataset. $\epsilon = 0.5$ is used to generate noise for existing methods (ϵ is always bounded by 0.5 in DPI).

	Δ after 5 slots	Δ after 50 slots	Actual DP after 5 slots	Actual DP after 50 slots
Chan et al. [24]	7	67	0.590	1.771
Chen et al. [25]	7	67	1.841	2.166
Perrier et al. [27]	7	67	1.734	2.248
Wang et al. [28]	7	67	1.674	1.972
DPI (ours)	2	2	0.001	0.005

2.10.7 Random Budget Allocation (Alternative)

Random budget allocation with a non-depleting total budget output can also be achieved by dividing the total privacy budget ϵ into several ranges like *high*, *medium*, and *low*. In each time slot t , ϵ_t is randomly selected from one of these ranges. Refer to Algorithm 4 for details of this alternative approach. In essence, the random allocation of the privacy budget allows us to cater to diverse and changing data streams, ensuring the high utility of DPI while preserving privacy.

Algorithm 4: RBA: Alternative Approach

```

1 initialize queues smallNumbers, mediumNumbers, largeNumbers
2 define the ranges: smallRange, mediumRange, largeRange
3 for each i in smallRange, mediumRange, and largeRange do
4   correspondingQueue.enqueue(i)
5 while true do
6   randomQueue  $\leftarrow$  randomly choose from (smallNumbers,
7     mediumNumbers, largeNumbers)
8   if randomQueue is not empty then
9     number  $\leftarrow$  randomQueue.dequeue()
10    return number

```

2.10.8 Additional Results

2.10.9 Privacy Loss of DPI vs Existing Methods

We additionally conduct experiments to compare the privacy loss of DPI with representative existing methods [24, 25, 27, 28]. Existing methods for differentially private data stream disclosures have several limitations. Most critically, they cannot preserve overall privacy loss in a bounded or converging manner. To prove this, we measure the expansion of sensitivity across the first 5 time slots and 50 time slots, respectively (in terms of the user-level DP protection) using the median query of packets in each time slot over a network traffic dataset in Table 2.3. In addition, we derive the privacy loss via the notion of Rényi differential privacy (e.g., deriving the Rényi differential privacy guarantees after injecting the noise and then converting the Rényi differential privacy to ϵ -DP [11]). Figure 2.16 further demonstrates that existing methods cannot bound the privacy over time. Thus, we do not benchmark DPI with them in the infinite stream settings.




Fig. 2.16: Total privacy loss of DPI and existing methods.

2.10.10 Additional Results on Synthetic Data

To further show that DPI can produce accurate output results on datasets with diverse characteristics, we conducted additional experiments on more synthetic data. Figure 2.17 shows the utility of DPI (w.l.o.g., using the accumulative query on the data distribution as an example) evaluated on additional 20 synthetic datasets and different privacy budgets ϵ varying from 0.5 to 10. Out of 594,000 query results (2,700 time slots, 11 different ϵ , 10 datasets, 2 different domain sizes), DPI can still ensure low and stable MSE and KL divergence in all these different settings. These results demonstrate high utility on datasets with diverse characteristics.

2.11 Conclusion

This paper addresses the significant challenges in an open problem: differential privacy for infinite data streams, which is crucial for various real-time monitoring and analytics applications. DP has been widely used to protect streaming data, but it has key limitations in terms of unbounded privacy leakage and sensitivity for ensuring user-level protection. We propose a novel solution, DPI, to effectively bound privacy leakage and enhance accuracy in real-time




Fig. 2.17: Evaluation of DPI on 10 synthetic datasets (each box includes 10 results for a specific privacy bound ϵ : the average MSE/KL of 2,700 time slots in each of the 10 datasets). (a) (b) MSE and KL divergence: domain size 1,000. (c) (d) MSE and KL divergence: domain size 10,000. Each experiment is repeated for 10 times, and the average results are plotted in the figures.

analysis on infinite data streams. We have conducted extensive theoretical studies to prove the convergence of privacy bound and the bounded errors (utility loss) for DPI over infinite data streams. We have also conducted comprehensive experiments to validate DPI on various real streaming applications and datasets.

Chapter 3

Universally Harmonizing Differential Privacy Mechanisms for Federated Learning: Boosting Accuracy and Convergence

Shuya Feng, Meisam Mohammady, Hanbin Hong, Shenao Yan, Ashish Kundu, Binghui Wang,

Yuan Hong

3.1 Introduction

As the volume of data generated by emerging applications continues to grow exponentially, data analysis tasks are often outsourced to cloud servers or distributed storage. However, the sharing of data is a critical issue due to its sensitivity, and any leakage could result in significant losses. Federated learning (FL) [92] has emerged as a solution that allows multiple parties to train a machine learning model jointly without exchanging local data. Despite the absence of local data exposure, sensitive information about the data can still be leaked through the exchanged model parameters via, e.g., membership inference attacks [1, 2, 93, 94, 95], data reconstruction attacks [4, 96, 97, 98], and attribute inference attacks [5, 99, 100, 101, 102].

Differential privacy (DP) has been proposed to provide rigorous privacy guarantees, ensuring that any data sample or user's data at any client cannot influence the output of a

function (e.g., the gradient or model parameters in FL) [10, 103, 104, 105]. However, directly applying existing DP mechanisms and accounting approaches to FL can result in excessive noise addition and loose privacy guarantees.

Recent techniques, such as the advanced accounting of privacy [10, 106, 107, 108] and Rényi Differential Privacy (RDP) [109], have shown remarkable results in optimizing accounting for privacy loss in machine learning. These techniques significantly improved the tradeoff between data privacy and model utility. Nevertheless, applying the Moments Accountant (MA) as a budget economic solution to other DP mechanisms can be challenging. One reason is that deriving the moments accountant for each DP mechanism often requires a heavy analysis of the tails in their probability density function (PDF), which can be difficult and may differ from one mechanism to another. Accountants rely heavily on the Gaussian mechanism to ensure DP, but this often results in excessive perturbation of gradients. This can make it hard to achieve a satisfactory balance between privacy and utility in existing DP-FL methods [110, 111, 112, 113, 114, 115, 116, 117], which are dominantly based on the Gaussian mechanism and the DP-SGD variants.

Building upon the limitations of existing DP-FL methods [110, 111, 112, 113, 114, 115, 116], we propose a novel universal solution for DP-FL called UDP-FL, which offers a comprehensive approach for achieving DP in FL that extends beyond the popular DP-SGD algorithm (Gaussian). It universally adopts different DP mechanisms (e.g., Staircase [118] *which greatly outperforms Gaussian in terms of noise magnitude*¹) and harmonizes privacy guarantees under a unified framework. It allows for tighter budget accounting and comparison of privacy

¹ The Staircase mechanism [118] has been proven to be optimal for ℓ_1 and ℓ_2 metrics for a wide range of privacy budget ϵ [82, 118].

guarantees between different DP-FL techniques, e.g., the Gaussian, Laplace and Staircase noise additive mechanisms, using the Rényi DP notion as a mediator variable [11, 108]. This approach provides greater flexibility and generalizability for real-world scenarios. Specifically,

DP-noise Harmonizer. To address the challenges of distributed learning systems, we rely on the Harmonizer component to enhance the budget accountant and management and improve the applicability and convergence of DP-FL across diverse scenarios. It harmonizes different DP mechanisms with the Rényi divergence, providing a generalized, flexible, and universal approach to DP-FL that adapts to various requirements, applications, and scenarios while ensuring the best privacy-utility trade-off for each case. Moreover, the Harmonizer maps the Gaussian Moments Accountant [10] to the corresponding Rényi DP of other DP mechanisms for FL, allowing it to measure privacy loss using Rényi divergence [10, 11, 108]. This algorithm calculates privacy leakage for each training round and guarantees that the leakage of the adopted DP mechanism (e.g., Staircase mechanism [118]) does not exceed the Gaussian version (viz. DP-SGD and other variants).

Universal Convergence Analysis. In this work, we also take the first cut to use the concept of mode connectivity² for universally analyzing the *convergence of different DP-FL frameworks* (including UDP-FL). This analysis is operationalized through a series of transformative steps: transitioning from DP mechanisms to mode connectivity [119, 120, 121], conducting a mode connectivity-based convergence analysis, and then relating the findings back to a broader range of DP mechanisms. Notably, UDP-FL demonstrates a markedly accelerated convergence rate compared to baselines (e.g., NbAFL [114], DP-SGD [10]) and sometimes even non-private

² Loss surfaces of deep neural networks can have many connected regions or “modes”. These modes can be connected by paths of low loss, which may be possible to move between during training.

aggregation algorithm FedAvg (No DP) [122], underscoring the effectiveness of our methodological innovations in convergence analysis.

Robustness against Privacy Attacks. UDP-FL demonstrates significant resilience against a spectrum of privacy attacks, aligning with advanced theories presented in recent work [123]. Through extensive empirical studies, our results reveal a substantial reduction in the success rate of Membership Inference Attacks (i.e., LiRA [2]) compared with non-private FL, highlighting the framework’s effectiveness in preserving privacy. While DP is not inherently designed to combat Attribute Inference Attacks, we observed a reduced correlation in feature learning. Against Data Reconstruction Attacks, UDP-FL has proven to be adept at preventing the reconstruction of original training data, thereby reinforcing its robustness in protecting data privacy within DP-FL environments.

Thus, the key contributions of this work are summarized below:

1. To our best knowledge, we propose the first DP-FL framework (UDP-FL), which universally harmonizes different differential privacy mechanisms with tighter privacy bounds and faster convergence compared to SOTA methods. The tighter privacy bounds and faster convergence of UDP-FL significantly improve the model accuracy (given the same privacy guarantees) and efficiency (i.e., reducing the computation and communication overheads) of different DP solutions for federated learning. Table 3.1 shows the superior performance of our UDP-FL framework compared to existing DP-FL methods.
2. To our best knowledge, we also take the first step to introduce a mode connectivity-based method for analyzing the convergence of DP-FL models. This approach provides novel insights into the complex dynamics of privacy-utility trade-offs, and our mode

connectivity-based method bridges the gap between theoretical analyses and practical applications for DP-FL (including UDP-FL) on convergence.

3. We comprehensively evaluate the robustness of our UDP-FL on defending against advanced privacy attacks, including the membership inference [1, 93], data reconstruction [124, 125], and attribute inference attacks [5, 99]. To our best knowledge, these are not fully explored for SOTA methods.

Table 3.1: Comparison of representative DP-FL methods on accuracy, privacy, and convergence (observed based on theoretical analyses and empirical results in all the works).

Method	Noise	Accuracy	Privacy	Conv.
DP-FedAvg [92, 115]	Gaussian	High	DP	Slow
NbAFL [114]	Gaussian	Low	DP	Slow
RDP-PFL [126]	Gaussian	High	DP	Slow
ALS-DPFL [127]	Gaussian	Low	DP	Slow
LDP-Fed [128]	Rand. Response	Low	LDP-Shuffle	N/A
CLDP-SGD [129]	Rand. Response	Low	LDP-Shuffle	N/A
LDP-FL [130]	Rand. Response	Low	LDP-Shuffle	N/A
COFEL-AVG [131]	Laplace	Low	DP	Slow
UDP-FL (Ours)	Universal	High	DP	Fast

3.2 Preliminaries

3.2.1 System and Adversaries

In this work, we follow the standard semi-honest adversarial setting for differentially private federated learning (DP-FL) where the adversary can possess arbitrary background knowledge. The server is honest-but-curious by following the protocol but attempting to derive private information about the client’s data from the exchanged messages during the training process. Clients are also categorized as “honest-but-curious”, by strictly adhering to the protocol without deviating

from established procedures. Key responsibilities include refraining from manipulating local model updates and avoiding the use of poisoned or false data in the training. Upholding these guidelines is essential for maintaining the integrity and security of the global model, ensuring its reliability and robustness.

In terms of privacy, both UDP-FL (across all mechanisms) and the DP-SGD are adding noise into local gradients during the training process. Despite this, the disclosure of trained model parameters is proven to preserve (ϵ, δ) -DP [10]. The model parameters, viewed as post-processed results of the DP guaranteed noisy gradients, do not affect the privacy leakage.

Despite the distributed nature of federated learning offering collaborative model training, the challenge of preserving data privacy persists. The privacy concerns stem from the potential leakage of sensitive information through clients' local model updates. We also empirically evaluate the performance of UDP-FL against privacy attacks, including membership inference attacks (MIAs) [1, 93], data reconstruction attacks (DRAs) [124, 125], and attribute inference attacks (AIAs) [5, 99]. Their settings (which are different from DP-FL) will be discussed in Section 3.5.

3.2.2 Federated Learning

FL is an emerging distributed learning approach that enables a central server to coordinate multiple clients to jointly train a model without accessing to the raw data. Assuming the FL system has N clients $C = \{C_1, C_2, \dots, C_N\}$ and each client C_k owns a private training dataset $\mathcal{D}_k = \{(\mathbf{x}_j^k, y_j^k)\}$ with $|\mathcal{D}_k|$ samples and each sample $\mathbf{x}_j^k$ has a label y_j^k . Then, FL considers the

following distributed optimization problem:

$$\min_w F(w) = \sum_{k=1}^N p_k F_k(w), \quad (3.1)$$

where $p_k \geq 0$ is the client C_k 's weight and $\sum_{k=1}^N p_k = 1$; Each client C_k 's local objective is defined by $F_k(w) = \frac{1}{|\mathcal{D}_k|} \sum_{j=1}^{|\mathcal{D}_k|} \ell(w; (\mathbf{x}_j^k, y_j^k))$, with $\ell(\cdot; \cdot)$ a user-specified loss function, e.g., cross-entropy loss.

FedAvg [122] is the *de facto* FL algorithm to solve Equation (3.1) in an iterative way. It has the following steps:

1. **Global Model Initialization.** The server initializes a global model w^0 , selects a random subset S_n of n clients from C , and broadcasts w^0 to all clients in S_n .
2. **Local Model Update.** In each global epoch t , each client C_k receives the global model w^t , initializes its local model w_k^t as w^t , and updates the local model by minimizing $F_k(w^t)$ on the local dataset D_k . E.g., when running SGD, we have: $w_k^t \leftarrow w_k^t - \eta_t \nabla_{w_k^t} F_k(w^t)$, where η_t is the learning rate in the t -th epoch.
3. **Global Model Update.** The server collects the updated client models $\{w_k^t\}$ and updates the global model w^{t+1} for the next round via an aggregation algorithm. For instance, when using FedAvg [122], the updated global model is: $w^{t+1} \leftarrow \frac{N}{n} \sum_{C_k \in S_n} p_k w_k^t$, which is then broadcasted to clients for the next round.
4. Repeat Steps 2 and 3 until the global model converges.

3.2.3 Differential Privacy and Rényi Accountant

The use of DP in FL enhances the benefits of collaborative model training with the need for protecting data privacy. It ensures that each data sample or user's contribution to the model training process is indistinguishable from others, and it can be implemented by adding noise to the gradients or parameters of the model or by using secure aggregation techniques. The notion of DP can be defined as below.

Definition 5 ((ϵ, δ) -Differential Privacy [9, 132]). *A randomization algorithm $\mathcal{A}$ is (ϵ, δ) -differentially private if for any adjacent databases d, d' that differ on a single element, and for any output set $\Omega \subseteq \text{range}(\mathcal{A})$, we have $\Pr[\mathcal{A}(d) \in \Omega] \leq e^\epsilon \Pr[\mathcal{A}(d') \in \Omega] + \delta$, and vice versa.*

FL with (ϵ, δ) -DP generally requires hundreds of training rounds to obtain a satisfactory model. Rényi accountant [107, 108] via Rényi Differential Privacy (RDP) [11] has been proposed to provide tighter privacy bounds on the privacy loss than the standard DP. RDP is defined over the Rényi divergence [109]. Recall that for two probability distributions P and Q , their Rényi divergence is defined as $\mathcal{D}_\alpha(P||Q) = \frac{1}{\alpha-1} \log \mathbf{E}_{x \sim Q}(\frac{P(x)}{Q(x)})^\alpha$ where x denotes a random variable and $\alpha > 1$ is the Rényi divergence order. Thus, the RDP can be defined as below.

Definition 6 ((α, γ) -Rényi Differential Privacy [11]). *A randomized mechanism $\mathcal{A}$ is said to have γ -Rényi differential privacy of order α , if for any adjacent datasets d, d' that differ on a single element, and for any output set $\Omega \subseteq \text{range}(\mathcal{A})$, the Rényi divergence $\mathcal{D}_\alpha[\mathcal{A}(d) = \Omega || \mathcal{A}(d') = \Omega] \leq \gamma$ holds.*

Rényi accountant [10, 108, 133] is a method for managing and assessing the cumulative privacy loss in a sequence of DP operations. It operates by tracking the cumulant generating

function (CGF) of the privacy loss random variable over the sequence of operations. Specifically, the Rényi accountant evaluates the CGF at a series of fixed points corresponding to different orders of Rényi divergence, thus enabling the calculation of an overall privacy guarantee for the sequence. This overall guarantee is expressed in terms of Rényi Differential Privacy, providing a more nuanced and tighter estimation of privacy loss compared to traditional methods. The Rényi accountant is particularly effective in complex scenarios, such as those encountered in machine learning algorithms, where multiple DP operations are composed over time.

Although the efficacy of the Rényi accountant in providing a refined estimation of privacy loss becomes increasingly significant when addressing privacy loss in FL, several challenges still exist. These include non-optimal noise mechanisms like Gaussian or Laplace that degrade accuracy, loose DP guarantees in complex systems, difficulty in tracking privacy loss across diverse clients, and reduced convergence speed. These limitations underscore the need for developing an enhanced DP-FL framework that *universally ensures tighter privacy out of diverse DP mechanisms* while maintaining fast convergence and accuracy. The Rényi accountant [133] adopted in UDP-FL helps to address these challenges by providing tighter accounting of privacy loss across numerous training rounds. Moreover, other recent accountants [10, 106, 107, 108] can also act as viable alternatives, offering flexibility in the choice of DP composition. Our framework’s design is orthogonal to the specific choice of accountant, meaning it is adaptable and could incorporate even tighter accounting methods as they become available in the future.

3.2.4 Mode Connectivity

Machine learning involves searching for a model parameter that minimizes a trainer-specified loss function [134]. We can visualize this process geometrically as moving on a loss surface

to search the global minima. However, recovering the global minima becomes harder as the dimension of the model parameters increases, making it more likely that training terminates at a local minimum for deep models [135, 136]. Recent research has shown that local minima on the loss surface (aka. *modes* of the loss function) are connected by simple curves, a phenomenon known as mode connectivity [119]. Mode connectivity is a strong ensembling technique (aka Fast Geometric Ensembling) developed to facilitate transporting from one mode (local minima) to another mode with possibly lower loss.

To find a curve between two modes, we perform a *curve finding procedure* similar to [119]. Let w_1 and w_2 be the model parameters in $\mathbb{R}^d$ for two local minima, and $\ell(w)$ be the cross-entropy loss function. Furthermore, let $\phi_\theta : [0, 1] \rightarrow \mathbb{R}^d$ be a continuous piece-wise smooth parametric curve, with parameters θ , such that the curve endpoints are $\phi_\theta(0) = w_1$ and $\phi_\theta(1) = w_2$. Our goal is to find the parameters θ that minimize the expectation of the loss $\ell(\phi_\theta(p))$ ($p \in [0, 1]$) w.r.t. a uniform distribution on the *curve parameter*:

$$L(\theta) = \mathbb{E}_{p \sim U(0,1)}[\ell(\phi_\theta(p))] \quad (3.2)$$

The generic curve function $\phi_\theta(p)$ has been characterized as a polygonal chain or Bezier curve in prior works [137, 138]. An example of the polygonal chain characterization can be seen in Equation 3.3. The bend in the curve is parameterized by θ , and the curve finding procedure minimizes Equation (3.2).

$$\phi_\theta(p) = \begin{cases} 2(p\theta + (0.5 - p)w_1), & 0 \leq p \leq 0.5 \\ 2((p - 0.5)w_2 + (1 - p)\theta), & 0.5 \leq p \leq 1 \end{cases} \quad (3.3)$$

The curve finding procedure begins by initializing θ as the midpoint of the line segment between w_1 and w_2 . Then, in each training round, we sample $\hat{p}$ from a uniform distribution, and a gradient step is applied to θ with respect to $\ell(\phi_\theta(\hat{p}))$. This effectively fixes the curve endpoints w_1 and w_2 while allowing the midpoint (i.e., parameters θ) to move in the weight space. To recover the model parameters from each point of a trained curve, we first train a polygonal curve using the described procedure. Once we obtain the optimal midpoint parameters θ , we can recover the model parameters by setting a value for the curve parameter p and evaluating the curve at that point. For instance, if we set $p = 0.75$, then the point $\phi_\theta(0.75)$ corresponds to the midpoint of the line segment between θ and the original curve endpoint w_2 , which can be expressed as $0.5w_2 + 0.5\theta$. Consequently, the value at each dimension of the recovered model is equal to the average value of that dimension between w_2 and the optimal parameters θ^* .

3.3 UDP-FL Framework

In this section, we propose a comprehensive framework, called universal DP-FL (UDP-FL), for achieving superior privacy-utility tradeoff and faster convergence in FL.

3.3.1 Building Blocks of UDP-FL

DP Mechanisms. UDP-FL’s ability to integrate diverse DP mechanisms makes it a universal framework for FL applications. It harmonizes different DP mechanisms (e.g., Laplace, Staircase [118], and also Gaussian) to provide strong privacy guarantees while preserving the quality of training outcomes. This diversity in DP mechanisms caters to the specific requirements of different FL applications, offering a more adaptable and effective approach than Gaussian-only methods. The Staircase mechanism, for instance, proves its optimality w.r.t. ℓ_1 and ℓ_2 metrics

(except the very small ϵ) [82, 118], showcasing its superiority in scenarios demanding stronger privacy guarantees without sacrificing model accuracy. This multi-mechanism strategy enhances UDP-FL’s flexibility and robustness, making it a versatile and powerful tool in DP-FL. In this paper, due to the optimality under specific ϵ settings, we adopt the Staircase mechanism [118] as our instantiated DP mechanism for FL, along with the Gaussian and Laplace mechanisms. UDP-FL with Gaussian are equivalent to DP-SGD in FedAvg which is further used as one of the baselines in the experiment section. Other alternative optimal or complex DP mechanisms [82, 139] would also work in a similar manner. The choice of mechanism depends on the specific requirements of the FL task, such as the nature of the data, the desired privacy-utility trade-off, and the computational constraints.

Harmonizer. The Harmonizer, pivotal in UDP-FL, stands as a fundamental tool enabling accurate, customized privacy accounting for a diverse range of DP mechanisms in FL. Adapting from the popular Moments Accountant technique, originally formulated for Gaussian noise, the Harmonizer extends its applicability to mechanisms like the Staircase, Laplace, and more, thereby addressing both ℓ_1 and ℓ_2 data sensitivities. The Harmonizer in UDP-FL serves as a centralized budget management system akin to Cohere’s budget control system. It allows clients to specify their own privacy budgets and noise mechanisms and then unifies these diverse preferences using Rényi divergence. The Harmonizer tracks the privacy budget consumption at a fine-grained level, considering both individual clients and the shuffler. This enables more efficient budget utilization and tighter privacy analysis through Rényi divergence.

In FL, where numerous training epochs are necessary, the Harmonizer plays a crucial role in balancing the privacy-utility tradeoff. It leverages the RDP order of α to track the privacy

loss across multiple iterations. This flexibility is essential in determining when to halt training based on accumulated privacy loss, thus preventing excessive privacy leakage.

While the Gaussian mechanism is a common choice in moments accountant for noise generation, UDP-FL primarily focuses on the Staircase mechanism for its superior adaptability and flexibility, especially in scenarios with limited knowledge about clients' private training data. Staircase does not undermine the utility of other mechanisms like Laplace, which can be optimal in specific cases, but positions Staircase as the more comprehensive choice in UDP-FL.

The noise multiplier is a scalar value that determines the magnitude of noise added to the model updates in differentially private federated learning. It is a crucial parameter that balances the trade-off between privacy and utility. The calculation of the noise multiplier prior to training, a key function of the Harmonizer, involves several steps: taking inputs like privacy parameters, training rounds, and client data sampling rate; initially setting and then adjusting the noise multiplier through trial runs; and computing Rényi divergence for each iteration to ensure privacy loss is within the set bounds (see detailed procedures in Algorithm 7 in Appendix 3.8.2). In addition, Table 3.3 presents the Rényi divergence for several example DP mechanisms (including the Laplace and Gaussian mechanisms as baselines). It exemplifies the Harmonizer's capability to harmonize various DP mechanisms, making UDP-FL a robust and adaptable framework in the realm of differentially private FL.

Overall, the Harmonizer not only selects the optimal noise multipliers but also enables clients to customize their privacy settings based on their specific needs and risk tolerance. By using Rényi divergence, it harmonizes various privacy preferences and noise mechanisms, ensuring efficient management of the overall privacy budget across all participants in the federated

learning process.

3.3.2 UDP-FL Framework

In this section, we present the main steps in UDP-FL.

1. **Local Data Preparation.** The clients collect and store their data locally. Once their data is ready, clients specify privacy parameters (ϵ, δ) and send them to the server.
2. **Global Model Initialization.** Server initializes a global model, selects a random subset of clients, and sends the current global model parameters to the selected clients.
3. **Local Model Update.** Selected clients update their model locally. First, initialize local model parameters with current global model parameters. Then, sample local data uniformly at random and split them into batches. Perform local model updates using the **Harmonizer**: (1) compute and clip gradients; (2) calculate the noise multiplier based on Algorithm 7 in Appendix 3.8.2; (3) add noise to the clipped gradients; (4) update local model parameters using noisy gradients and learning rate; (5) send the updated local model parameters to the server.
4. **Global Model Update.** The server receives the selected clients' model parameters and aggregates them via an aggregator, e.g., FedAvg, to update the global model.

The detailed procedures of UDP-FL are illustrated in Algorithm 5.

Algorithm 5: UDP-FL Framework

Input: Global epochs T , local epochs I , privacy bound ϵ , clients

 $C = \{C_1, \dots, C_N\}$, clip parameter c , noise parameters Δ, μ
Output: Model parameters w^*

```

1  $t \leftarrow 0$  // Initialize global epoch counter
2  $w^t \leftarrow$  initialize global model
3 while  $t < T$  do
    // Server polling
4   Select  $S_n \subset C$ 
5   Broadcast  $w^t$  to  $S_n$ 
    // Client training
6   forall  $C_k \in S_n$  do
7     Set local model  $w_k^t = w^t$ 
8     for  $i = 0$  to  $I - 1$  do
9       Sample from  $\mathcal{D}_k$  with rate  $q$ 
10      forall sampled data do
11        Compute gradient  $g_k = \nabla_{w_k^t} F_k(w^t)$ 
12        Clip gradient  $g_k^c = g_k \cdot \min\left(1, \frac{c}{\|g_k\|}\right)$ 
13        Add noise  $g_k^c + \text{Noise}(\Delta, \mu)$ 
14        Update  $w_k^t \leftarrow w_k^t - \eta g_k^c$ 
15        Send  $w_k^t$  to Server
16      end
    // Server model update with Mode Connectivity
17    while  $|S_n| > 1$  do
18      Initialize  $W_{\text{int}}$  // List of intermediate models
19      forall pairs  $(w_i, w_j)$  in  $S_n$  do
20        Initialize  $\theta_{ij} = 0.5(w_i + w_j)$  // Midpoint initialization
21        Define  $\phi_{\theta_{ij}}(p)$  as a piecewise function
22        do
23          Sample  $\hat{p} \sim U(0, 1)$ 
24          Compute  $L(\theta_{ij}) = \ell(\phi_{\theta_{ij}}(\hat{p}))$ 
25          Compute gradient  $\nabla_{\theta_{ij}} L(\theta_{ij})$ 
26          Update  $\theta_{ij} \leftarrow \theta_{ij} - \eta \nabla_{\theta_{ij}} L(\theta_{ij})$ 
27        while not converged
28        Add  $\theta_{ij}$  to  $W_{\text{int}}$ 
29       $S_n \leftarrow W_{\text{int}}$ 
30    end
31    Update  $w^t \leftarrow S_n[0]$ 
32     $t \leftarrow t + 1$ 
33 end
34 return  $w^*$ 

```

3.3.3 UDP-FL Harmonizer

We can enhance Algorithm 5 to address heterogeneity in individual privacy guarantees of clients and harmonize them through mode connectivity. Let $\epsilon_{\max} = \max_k \epsilon_k$ represent the largest (weakest) privacy parameter among all clients. The UDP-FL Harmonizer algorithm penalizes each local model update based on its deviation from the strongest DP guarantee. Specifically, a penalty term is incorporated into the update process, where the penalty is proportional to the difference between the client's privacy parameter and $\epsilon_{\max}$. The update rule for client C_k will be modified as follows:

$$w_k^t \leftarrow w_k^t - \eta (g_k^c + \lambda_k (w_k^t - w_{\max}^t))$$

Where:

- w_k^t denotes the local model of client C_k at global epoch t .
- g_k^c represents the clipped gradient computed by client C_k .
- λ_k is the penalty coefficient for client C_k , calculated as $\lambda_k = \frac{\epsilon_{\max} - \epsilon_k}{\epsilon_{\max}}$.
- $w_{\max}^t$ corresponds to the model with the largest (weakest) DP guarantee among all clients at global epoch t .

It enables UDP-FL to mitigate privacy heterogeneity among clients while promoting convergence towards a unified global model.

3.4 Theoretical Analyses

In this section, we will theoretically analyze the privacy and utility of UDP-FL and discuss the impact of its parameters on the privacy analysis. Moreover, we present a comprehensive convergence analysis for general DP-FL and UDP-FL. These analysis utilizes the notation listed in Table 3.2.

Table 3.2: Frequently used notations for convergence analysis.

Symbol	Description
N	the number of clients
T	number of global epochs
p_k	the weight of the k -th device
n_k	the number of training data samples in the k -th device
w_t^k	the latest model on the k -th device
η_t	the decaying learning rate at round t
v_t^k	the immediate result of one step SGDUpdate from w_t^k
w^*	the model which minimizes F
ξ_t^k	a sample uniformly chosen from the local data
$F(\cdot)$	the weighted average loss
$F_k(\cdot)$	a user-specified loss function
$F_k^*(\cdot)$	the minimum of $F_k(\cdot)$
F^*	the minimum of F
σ_k^2	the upper bound variance of stochastic gradients in the k -th device
G	the upper bound expected squared norm of stochastic gradients in k -th device
μ	the strongly convex constant of $F_k(\cdot)$
L	the smoothness constant of $F_k(\cdot)$
θ	mode connectivity parameter
r	mode connectivity percentage

3.4.1 Rényi Differential Privacy

Rényi Differential Privacy (RDP) [11] is a refinement of standard Differential Privacy (DP) that provides tighter privacy guarantees. It is defined using the Rényi divergence, which measures the

difference between probability distributions. Consider a randomized mechanism $\mathcal{A}$ operating on datasets d, d' that differ by at most one element. The mechanism $\mathcal{A}$ achieves $(\alpha, \epsilon_\alpha)$ -RDP if the Rényi divergence between the outputs of $\mathcal{A}$ on d and d' is bounded by ϵ_α , denoted as:

$$\mathcal{D}_\alpha[\mathcal{A}(d) = \Omega || \mathcal{A}(d') = \Omega] \leq \epsilon_\alpha$$

Here, α is the order of the Rényi divergence, and ϵ_α quantifies the privacy loss. RDP exhibits key properties akin to differential privacy, including adaptive composition and conversion to traditional DP.

- **Adaptive Composition of RDP:** If M_1 obeys (α, ϵ_1) -RDP and M_2 obeys (α, ϵ_2) -RDP, their composition obeys $(\alpha, \epsilon_1 + \epsilon_2)$ -RDP (proof in [11]).
- **RDP to DP Conversion:** If mechanism M obeys $(\alpha, \epsilon_\alpha)$ -RDP, then M obeys $(\epsilon_\alpha + \log(1/\delta)/(\alpha - 1), \delta)$ -DP for all $0 < \delta < 1$ (proof in [11]).

RDP accounting after T rounds of training will budget DP mechanisms as

$$\epsilon(\delta, T) = \min_{\alpha > 1} \left\{ \epsilon_\alpha + \frac{\log(1/\delta)}{\alpha - 1} \right\}.$$

Table 3.3 first summarizes three representative DP mechanisms in UDP-FL (including Gaussian) that can realize the Rényi moments accountant for federated learning, as well as the Rényi divergence used to keep track of privacy loss during the training process.

Table 3.3: Rényi DP of three main noise mechanisms utilized in UDP-FL along with their corresponding parameters.

DP Mechanism	Rényi divergence ($\alpha > 1$)	Parameters for ϵ -RDP
Gaussian	$\frac{\alpha\Delta}{2\sigma^2}$	$\sigma = \frac{\Delta}{\alpha\sqrt{2\ln(1.25/\delta)}}$
Laplace	$\frac{1}{\alpha-1} \log \left(\frac{\alpha}{2\alpha-1} \exp \left(\frac{\Delta(\alpha-1)}{\lambda} \right) \right) + \frac{\alpha-1}{2\alpha-1} \exp \left(\frac{-\Delta\alpha}{\lambda} \right)$	$\lambda = \frac{\Delta(\alpha-1)}{\epsilon}$
Staircase	$\frac{1}{2}e^{(\alpha-1)\epsilon} + \frac{1}{2}e^{-\alpha\epsilon} + (e^{(\alpha-1)\epsilon} + e^{-\alpha\epsilon})(1-\nu) + 2\nu-1 e^{-\text{sgn}(\frac{1}{2}-\nu)\epsilon} \frac{1-e^{-1}}{2(\nu+e^{-\epsilon}(1-\nu))}$	$\nu = \frac{1}{2} \left(1 + e^{-\epsilon} - \sqrt{1 + (e^{-\epsilon} - 1)(e^{2(\alpha-1)\epsilon} - 1)} \right)$

3.4.2 Error Bounds Analysis of UDP-FL

The Staircase mechanism can be considered as a geometric mixture of uniform probability distributions. It ensures an optimal privacy-utility tradeoff, particularly for medium and relatively large ϵ values. This mechanism generates noise in a controlled manner by mixing uniform probability distributions, adjusting for the privacy budget and other requirements, and adding this noise to the query responses to preserve privacy without significantly degrading accuracy.

The Staircase mechanism for a function f is defined as:

$$f = \begin{cases} e^{-\rho\lambda y} & \|x\|_1 \in [\rho\Delta, (\rho+\nu)\Delta] \\ e^{-(\rho+1)\lambda y} & \|x\|_1 \in [(\rho+\nu)\Delta, (\rho+1)\Delta] \end{cases}$$

for $\rho \in \mathbb{N}$, where:

$$y \triangleq \frac{1 - e^{-1}}{2\Delta(\nu + e^{-\lambda}(1 - \nu))}$$

Theorem 6 enables the Harmonizer in UDP-FL to support the MA that tracks the privacy

loss of the Staircase mechanism during federated learning.

Theorem 6 (Proof in Appendix 3.8.1). *For any $\alpha > 1$, $\epsilon_\alpha > 0$, the Staircase mechanism satisfies $(\alpha, \epsilon_\alpha)$ -Rényi differential privacy (RDP), where ϵ_α is*

$$\frac{1}{2}e^{(\alpha-1)\lambda} + \frac{1}{2}e^{-\alpha\lambda} + \left((e^{(\alpha-1)\lambda} + e^{-\alpha\lambda})(1-\nu) + |2\nu-1|e^{-\text{sgn}(\frac{1}{2}-\nu)\lambda} \right) \times \frac{1-e^{-1}}{2(\nu+e^{-\lambda}(1-\nu))}$$

We now provide the error bounds for the DP mechanism (i.e., noise applied to the model parameters).

Lemma 2 (Proof in Geng et al. [118]). *Given the Staircase mechanism $f(\lambda, \Delta, \nu)$, when $\nu = \frac{1}{1+e^{\lambda/2}}$, the minimum expectation of noise amplitude is $\Delta \frac{e^{\lambda/2}}{e^\lambda - 1}$.*

As demonstrated by Geng et al. [118], when the privacy budget ϵ is sufficiently small, the Staircase mechanism saves at least $\Delta^2 \left(\frac{1}{12} - \frac{\epsilon^2}{720} + O(\epsilon^4) \right)$ perturbation in variance compared to Gaussian and Laplacian mechanisms. Since our privacy budget spent in each round is negligible ($\epsilon \rightarrow 0$), we can capitalize on this improved accuracy payoff per round, leading to enhanced model performance while maintaining robust privacy guarantees.

Theorem 7 (Proof in Appendix 3.8.1). *For any $\alpha > 1$, $\gamma > 0$, Staircase mechanism satisfies (α, γ) -Rényi differential privacy, where*

$$\gamma = \frac{1}{2}e^{(\alpha-1)\lambda} + \frac{1}{2}e^{-\alpha\lambda} + \left\{ (e^{(\alpha-1)\lambda} + e^{-\alpha\lambda})(1-\nu) + |2\nu-1|e^{-\text{sgn}(\frac{1}{2}-\nu)\lambda} \right\} \frac{1-e^{-1}}{2(\nu+e^{-\lambda}(1-\nu))} \quad (3.4)$$

Theorem 8 (Proof in Appendix 3.8.1). *The expectation of the ℓ_1 distance for the output model parameters preserved by UDP-FL with the Staircase mechanism after T training rounds is:*

$$\frac{mT}{1 - e^{-\lambda}} \left(v^2 \Delta^2 + e^{-\lambda} \Delta^2 - e^{-\lambda} v^2 \Delta^2 + \Delta e^{-\lambda} \right)$$

where m is the length of the loss function, and v, ρ, λ are the noise multipliers computed by UDP-FL.

Faster Convergence of UDP-FL (e.g., Staircase). Furthermore, the utilization of Staircase noise has been demonstrated to significantly accelerate convergence compared to the baseline, as empirically validated in Figure 3.3, Figure 3.5, and Table 3.1. The enhanced convergence speed is a byproduct of applying the optimal noise for ℓ_1 and ℓ_2 distance metrics (for a wide range of ϵ). In essence, when DP is fixed to guarantee ϵ , this noise has been proven to minimize both ℓ_1 and ℓ_2 distances. This implies that all noise-additive operations, including gradient perturbation and client model averaging, yield more accurate results, closer to the non-private scenario.

This observation aligns with findings by Subramani et al. [140] in the context of the JIT framework, where a solid relationship between noisy and less noisy schemes is explored (along with convergence analysis). Specifically, they report a faster epoch for less noisy schemes.

3.4.3 Mode Connectivity-based Convergence

We can utilize mode connectivity to study the convergence of deep neural network (DNN) models, as it provides a flexible procedure for selecting subsequent minima without imposing

stringent conditions on the model parameters.

First, we introduce the worst-case estimate for the additional number of rounds required by UDP-FL. Then, we present an optimal strategy to accelerate convergence. Let w_t^k be the model parameter maintained in device k in training round t .

Theorem 9 (Proof in Appendix 3.8.1). *Let w^* represent the model parameters at a local minimum. Suppose θ_t^k and $\tilde{\theta}_t^k$ are the parameters trained over the curve function $\phi_\theta(p)$ for the non-DP and DP models, respectively, where the DP model is perturbed by the Staircase mechanism with parameters Δ and ϵ . The system UDP-FL will transport $\tilde{\theta}_t^k$ to w^* with less than $O\left(\Delta^2 \cdot \frac{e^\epsilon}{(e^\epsilon - 1)^2}\right)$ additional rounds compared to the baseline federated learning (FL) for θ_t^k to reach w^* .*

Theorem 10. *The average of the optimal choice of parameter $\bar{\theta}_{t+1}$ which minimizes Equation 3.2 over a quadratic Bezier curve is*

$$\bar{\theta}_{t+1}^* = \frac{1.2}{L} + 1.1\bar{w}_t - 0.1\bar{v}_{t+1} \quad (3.5)$$

For a FedAvg algorithm trained on a quadratic Bezier curve with curve parameter r (equivalent to p in Section 3.2.4) randomly sampled from $U(0, 1)$, the update of FedAvg with full clients active can be described as

$$w_{t+1}^k = (1 - p)^2 v_{t+1}^k + 2(r - r^2)\theta_{t+1}^k + r^2 w_t^k, \quad (3.6)$$

where $v_{t+1}^k = w_t^k - \eta_t \nabla F_k(w_t^k, \xi_t^k)$ and θ_{t+1}^k is the Bezier function parameter which minimizes Equation 3.2.

Similar to [141, 142], we define two virtual sequences $\bar{v}_{t+1} = \sum_{k=1}^N p_k v_{t+1}^k$ and $\bar{w}_t = \sum_{k=1}^N p_k w_t^k$ where the former originates from a single step of SGD from the latter. Finally, it is clear that we always have

$$\bar{w}_{t+1} = (1 - \bar{r})^2 \bar{v}_{t+1} + 2(\bar{r} - \bar{r}^2) \bar{\theta}_{t+1} + \bar{r}^2 \bar{w}_t$$

Proof. Suppose the step-size between the two sequential average model weights $\bar{w}_{t+1}$ and $\bar{w}_t$, i.e., $|w_{t+1}^k - w_t^k|$, is α . From the L-smoothness of F_k , it follows that

$$\begin{aligned} E_k \left\{ E_{\xi_{t+1}^k} (F_k(w_{t+1}^k, \xi_{t+1}^k)) \right\} &\leq E_k \left\{ E_{\xi_t^k} (F_k(w_t^k, \xi_t^k)) \right\} \\ &\quad - \left| \frac{\alpha^2 L}{2} - \alpha \right| \cdot E_k \left\{ \|\nabla F_k(w_t^k, \xi_t^k)\|_2^2 \right\} \end{aligned} \quad (3.7)$$

Let $\bar{F}(\bar{w}_t) = E_k \left\{ E_{\xi_t^k} (F_k(w_t^k, \xi_t^k)) \right\}$ and $\bar{\nabla} F(\bar{w}_t) = E_k \left\{ \nabla F_k(w_t^k, \xi_t^k) \right\}$, we can rewrite Equation (3.7) as follows.

$$\bar{F}(\bar{w}_{t+1}) \leq \bar{F}(\bar{w}_t) - \left| \frac{\alpha^2 L}{2} - \alpha \right| \cdot \bar{\nabla} F(\bar{w}_t) \quad (3.8)$$

A high-accuracy path between v_{t+1}^k and w_t^k leverages the parameters θ that minimize the expectation over a uniform distribution on the curve [119], i.e.,

$$\theta_{t+1}^{k*} = \arg \min_{\theta} \left\{ E_{r \sim U(0,1)} [F_k(w_{t+1}^k, \xi_{t+1}^k)] \right\} \quad (3.9)$$

From Equation (3.7), it follows that

$$\min_{\theta} \left\{ E_{r \sim U(0,1)} [\bar{F}(\bar{w}_{t+1})] \right\} \Leftrightarrow \max_{\theta} \left| \frac{\alpha^2 L}{2} - \alpha \right| \quad (3.10)$$

But the expression on the right hand side is strictly increasing w.r.t. α and its maxima occurs at $\alpha = \frac{1}{L}$. Therefore, we will solve the following problem.

$$\bar{\theta}_{t+1}^* = \arg \min_{\theta} |\alpha(\theta) - \frac{1}{L}| \quad (3.11)$$

From the expression for $\bar{w}_{t+1}$ it follows that

$$\alpha(\theta) - \frac{1}{L} = (1 - \bar{r})^2 \bar{v}_{t+1} + 2(\bar{r} - \bar{r}^2) \bar{\theta}_{t+1} + (\bar{r}^2 - 1) \bar{w}_t - \frac{1}{L}$$

Therefore,

$$\begin{aligned} \min_{\theta} |\alpha(\theta) - \frac{1}{L}| &\Leftrightarrow \\ \min_{\theta} \{ |(1 - \bar{r})^2 \bar{v}_{t+1} + 2(\bar{r} - \bar{r}^2) \bar{\theta}_{t+1} + (\bar{r}^2 - 1) \bar{w}_t - \frac{1}{L}| \} \end{aligned}$$

By rearranging the term inside the minimum function and substituting $\bar{r}^n$ by $\frac{1}{n}$ (refer to moments of uniform random variable defined over $[0, 1]$), we have

$$\begin{aligned} \bar{\theta}_{t+1}^* &= \arg \min_{\bar{\theta}_{t+1}} |\bar{v}_{t+1} - \bar{w}_t - \frac{1}{L} + \bar{\theta}_{t+1} - \bar{v}_{t+1} + \frac{1}{12}(\bar{v}_{t+1} + \bar{w}_t - 2\bar{\theta}_{t+1})| \\ &= \arg \min_{\bar{\theta}_{t+1}} | -\frac{1}{L} + \frac{5}{6}\bar{\theta}_{t+1} + \frac{1}{12}\bar{v}_{t+1} - \frac{11}{12}\bar{w}_t | \end{aligned}$$

Finally, by forcing the term inside the absolute value to be zero we must have

$$\bar{\theta}_{t+1}^* = \frac{1.2}{L} + 1.1\bar{w}_t - 0.1\bar{v}_{t+1}.$$

Thus, this completes the proof. $\square$

As a result, the average loss is minimized if all clients set their mode connectivity parameter to $\theta_{t+1}^k = \frac{1.2}{L} + 1.1w_t^k - 0.1v_{t+1}^k$. Alternatively, if the updates are subjected to clipping, FedAvg can maintain a fixed step size of $\frac{1}{L}$. These constraints are integrated into our UDP-FL, as shown in Algorithm 6 in Appendix 3.8.1.

3.4.4 UDP-FL Framework with Shuffler

The Shuffler [143] is popular for preserving user privacy while enabling data usage from multiple sources. It anonymizes the network by shuffling client identities, enhancing differential privacy guarantees [129, 130, 143, 144]. UDP-FL supports this model for tighter privacy bounds and faster convergence. When integrated with UDP-FL, the Shuffler concatenates model updates and applies a random permutation to clients' IDs and updates before server aggregation:

1. **Permutation of Client IDs and Model Parameters:** Obscures the correlation between clients' IDs and their model parameters, ensuring anonymity.
2. **Updating Moments Accountant:** Updates the Moments Accountant with current privacy parameters (ϵ, δ) to manage cumulative privacy loss.
3. **Calculation of Current Privacy Loss:** Computes the privacy loss for the current round, ensuring adherence to privacy constraints.
4. **Computation of Final Privacy Loss:** Utilizes the Moments Accountant and Rényi divergence to compute overall privacy loss across all rounds.

By shuffling the order of local model updates, the Shuffler obscures individual contributions, making it harder for adversaries to infer sensitive information. This results in (α, γ_s) -RDP, with γ_s being less than γ in the previous step, thus requiring less noise to meet the same privacy requirements.

Privacy Amplification by Shuffler

Integrating a Shuffler in UDP-FL amplifies privacy protection [145, 146], proportional to the number of clients. The Shuffler can be used with any DP mechanism (e.g., Gaussian, Laplace, and Staircase), extending the moments accountant setting under Rényi differential privacy for federated learning.

Lemma 3. *Given N clients, any $\gamma \geq 0$, and any integer $\alpha \geq 2$, the Rényi divergence of the shuffle model is upper-bounded by γ_u :*

$$\gamma_u = \frac{1}{\alpha - 1} \log \left(1 + \binom{\alpha}{2} \frac{4(e^\gamma - 1)^2}{N} \right)$$

Lemma 4. *Given N clients, any $\gamma \geq 0$, and any integer $\alpha \geq 2$, the RDP of the shuffle model is lower-bounded by γ_l :*

$$\gamma_l = \frac{1}{\alpha - 1} \log \left(1 + \binom{\alpha}{2} \frac{(e^\gamma - 1)^2}{Ne^\gamma} \right)$$

These lemmas, proven in Girgis et al. [130], indicate that the Shuffler amplifies privacy in proportion to the number of clients.

Evaluation of UDP-FL with Shuffler

Experiments, shown in Figure 3.4, reveal that integrating a Shuffler marginally improves utility, accelerates convergence, and yields the best performance when combined with the Staircase mechanism. These benefits increase with the number of clients.

3.5 Experiments

In this section, we will evaluate the performance of UDP-FL on privacy, accuracy and efficiency. The key objectives of our evaluations are: (1) Assessing the accuracy of UDP-FL in comparison with SOTA mechanisms. (2) Investigating the convergence behavior of UDP-FL to understand how its hyperparameters influence the training performance. (3) Demonstrating UDP-FL’s computational and communication efficiency against baseline methods. (4) Rigorously testing UDP-FL’s resilience against common privacy attacks, including membership inference, data reconstruction, and attribute inference attacks, to validate its defense performance.

3.5.1 Experiment Setup

Datasets. We conduct experiments to evaluate UDP-FL primarily on the following datasets:

- MNIST: 70,000 images of handwritten digits 0 – 9 [147].
- Medical MNIST: 60,000 grayscale images across 10 medical conditions [148].
- CIFAR-10: 60,000 32×32 color images in 10 classes (50,000 training, 10,000 testing images) [149].

ML Models. For the MNIST dataset, we utilize a two-layer CNN with ReLU activation, max pooling, and two fully connected layers. It processes single-channel images to predict class labels. In contrast, the Medical MNIST dataset employs a four-layer CNN with additional fully connected layers designed for 3-channel image classification. Both networks use ReLU activation and log softmax in the final output layer with cross-entropy loss. For CIFAR-10, we use the ResNet-18 architecture [150] other than pre-trained models. This choice, focusing on initial training, may lead to lower accuracy but is crucial for assessing UDP-FL’s influence on early-stage learning and privacy in federated learning.

Parameters Setting. We have used a learning rate of 0.01 in all experiments. For the MNIST and Medical datasets, we have set the clipped gradient value of the ℓ_2 norm to 1 and 0.1, respectively. For the CIFAR-10, and UTKFace datasets (to be used in the defense evaluation against privacy attacks in Section 3.5.5), the clipped gradient value is set to 0.01. The default number of clients is set at 10, and the sampling rate is set at 0.05. We set the local communication round to be 150 and the local training epoch to be 2.

Experiments are performed on two servers: one is equipped with two NVIDIA H100 Hopper GPUs (80GB each) while the other one is installed with multiple NVIDIA RTX A6000 GPUs (48GB each).

3.5.2 Accuracy Comparison vs Baseline Methods

Due to the diverse settings for DP guarantee, trust model, noise injection, and model architecture in DP-FL, there is no universally accepted benchmark for evaluating DP-FL methods. Therefore, in this work, we will compare UDP-FL with NbAFL [114] and DP-SGD applied to FedAvg (equivalent to UDP-FL with Gaussian noise), as they share similar settings (sample-level




Fig. 3.1: Accuracy and convergence results of UDP-FL and the baselines. 1) among the three mechanisms, the Staircase always performs the best with the same privacy budget; 2) UDP-FL obtains significantly better privacy-utility tradeoff and faster convergence than the baseline; and 3) UDP-FL (Staircase) even has a comparable accuracy with FedAvg (No DP).

DP within each client and local noise injection before aggregation). We also apply the classic FedAvg [92] without DP guarantees as the baseline in the experiments. Figure 3.1 shows the comparison results. We can observe from Figure 3.1(a) and 3.1(b) that UDP-FL, when using the Staircase mechanism, obtains higher accuracy and faster convergence rates compared to other methods. For instance, on the MNIST dataset, UDP-FL achieves 90% accuracy in about 25 epochs, while other methods struggle to reach this level even after 100 epochs. This faster

convergence is particularly evident in the Medical MNIST dataset, where UDP-FL converges in approximately 30 epochs, while other methods fail to converge even after 100 epochs. Furthermore, UDP-FL achieves nearly the same accuracy as FedAvg without DP guarantee. The primary reason is that we evenly distribute the datasets to each client so that their local datasets are unique.

Moreover, each time we randomly select some clients to update their models, this may cause the global model to take more time to converge because the data distribution from clients is more heterogeneous. Thus, the noise generated from UDP-FL helps to balance the unbiased distribution and contributes to faster convergence. UDP-FL requires fewer training epochs to reach the optimal performance (as shown in Figure 3.1(c) and 3.1(d)) whereas other methods take more training epochs but still result in lower accuracy.

3.5.3 Convergence Results

Figure 3.1(c) and 3.1(d) further prove that even when the data distributions from different clients are heterogeneous, UDP-FL can still preserve good performance and converge faster than other baselines. For instance, in Figure 3.1(c), it takes about 25 epochs for UDP-FL to achieve an accuracy of 0.9 and the accuracy tends to be stable since then. However, the accuracy of other baselines is low and fluctuated. Similar results can be seen in Figure 3.1(d), where it only takes about 30 epochs for UDP-FL to converge on the medical dataset. In contrast, other baselines did not converge even with 100 epochs. Thus, UDP-FL converges faster and has better performance.

Figure 3.2 for UDP-FL on CIFAR-10 shows that with a small privacy budget, UDP-FL with Staircase and Laplace mechanism converge more smoothly than DP-SGD. At a larger privacy budget, UDP-FL achieves rapid convergence comparable to the non-private baseline,




Fig. 3.2: UDP-FL on CIFAR-10 when (a) $\epsilon = 2$ and (b) $\epsilon = 8$. For a small privacy budget ($\epsilon = 2$), DP-SGD yields better performance, while at a larger privacy budget ($\epsilon = 8$), UDP-FL with Staircase mechanism outperforms DP-SGD and Laplace.

matching the theoretical results on faster convergence.

3.5.4 UDP-FL Evaluation

The Number of Clients. As discussed in Section 4.5, the enhanced privacy is related to a number of clients N and sampling rate q . Analyzing the impact of these hyperparameters allows us to gain insights into the scalability and adaptability of UDP-FL, ensuring optimal performance across different settings. We use UDP-FL with the Staircase mechanism to represent the optimal randomization (w.r.t. a range of ϵ) and Laplace mechanism as another baseline besides DP-SGD. We will present the performance of UDP-FL with other mechanisms in Section 3.5.2.

The number of clients significantly impacts the performance, communication overhead, model convergence, and privacy-utility tradeoff in FL frameworks. We experimented with 50, 100, 150, 200 clients (as shown in Figure 3.3), selecting 10% of all the clients randomly per training round. Each chosen client trains locally on 5% of their data for 2 epochs. As shown

in Figures 3.3(a) and 3.3(c), the accuracy slightly decreases with more clients, likely due to the increased complexity in aggregating diverse model updates. This variance affects convergence and generalization but, notably, the performance drop is not substantial, showing UDP-FL’s effectiveness in handling scalable, heterogeneous client scenarios in FL.

Sampling Rate. The sampling rate can significantly impact the convergence speed, model performance, and privacy-utility tradeoff in UDP-FL. A higher sampling rate means more data samples are used in each training epoch, which can lead to faster convergence and better model performance. However, this increased usage may cause higher privacy loss. Thus, more noise will be used to preserve privacy. Figure 3.3(b) and 3.3(d) confirm the results in the theoretical analysis (as discussed in Section 4.5). Our experiments showed that as the sampling rate increases, the model accuracy will be improved, and the convergence becomes faster. At a low sampling rate (e.g., 0.01), the model accuracy will be lower, and it will take more rounds to achieve convergence. Conversely, at a high sampling rate (e.g., 0.5), the model can achieve higher accuracy and converge more quickly. These results indicate that including more data samples in each training round can lead to better model performance.

Performance across Privacy Settings. In our comprehensive evaluation, we first delve into the privacy-utility trade-off on the MNIST and Medical datasets, as depicted in Figure 3.4. UDP-FL’s performance with the Staircase mechanism exhibits a steady increase in accuracy as the privacy budget increases and outperforms other baseline methods. The accuracy versus epochs on both datasets reveals UDP-FL’s capability for consistent learning over time, even outperforming the non-private baseline (FedAvg) during early epochs on the Medical dataset. These results have validated the practicality of UDP-FL in scenarios where stringent privacy is




Fig. 3.3: Impact of the number of clients and sampling rate on UDP-FL. We observe that: 1) when the number of clients increases, shown in Figures (a) and (c), UDP-FL needs more epochs to converge; 2) with the increase of data sampling rate shown in Figures (b) and (d), UDP-FL converges faster.

required without substantially compromising model performance.

Subsequently, we extend our evaluation to the CIFAR-10 dataset, which offers a more complex challenge due to its higher dimensionality and diverse image representations. This further examination on CIFAR-10 aims to validate UDP-FL's robustness and scalability in more intricate visual data scenarios. From the experimental results on CIFAR-10 dataset in Table 3.4, we observed that DP-SGD has shown consistent moderate accuracy across varying privacy




Fig. 3.4: Performance evaluation of UDP-FL on MNIST and Medical datasets. (a) and (c) present UDP-FL’s accuracy under various differential privacy noise mechanisms, compared to a non-private baseline, with varying ϵ values. (b) and (d) illustrate the learning curves over training epochs for the MNIST and Medical datasets without privacy and with DP guarantees.

levels, peaking when no privacy constraint was applied ($\epsilon = \infty$). UDP-FL with the Laplace mechanism, while less accurate at stricter privacy settings, improved as privacy constraints were relaxed. Notably, UDP-FL with the Staircase mechanism initially underperformed at lower ϵ values but significantly improved with relaxed privacy, surpassing Gaussian. This observation demonstrates that when the privacy protection is medium and relatively weaker, UDP-FL with the Staircase mechanism can effectively improve the trade-off between privacy and utility, even

in the FL setting.

Table 3.4: UDP-FL accuracy vs. privacy guarantees.

Datasets	Mechanisms	$\epsilon = 2$	$\epsilon = 4$	$\epsilon = 6$	$\epsilon = 8$	$\epsilon = \infty$
CIFAR-10	Gaussian	0.704	0.714	0.719	0.742	0.871
	Laplace	0.475	0.461	0.506	0.638	0.871
	Staircase	0.633	0.691	0.733	0.780	0.871

Table 3.5: Runtime of UDP-FL (sec) vs # training iterations.

Iterations	50	100	500	1000
MNIST (50 clients)	1049.71	2056.34	10500.80	20803.10
Medical (50 clients)	644.97	1328.23	6482.50	13177.98
CIFAR-10 (50 clients)	1808.67	3515.08	16443.23	31956.76
MNIST (100 clients)	1108.91	2196.13	10934.62	21579.99
Medical (100 clients)	714.02	1445.75	7064.78	14353.64
CIFAR-10 (100 clients)	1873.93	3653.36	17215.17	33562.24
Computing Parameters	50.50	50.60	51.60	54.30

Computation and Communication Overheads. The efficiency of a DP-FL framework is a critical aspect to be considered, where the computation and communication overheads can accurately reflect the overall system performance. In this section, we evaluate the computation and communication overheads of UDP-FL. Specifically, we will present the total local training time and noise multiplier computation time in UDP-FL, which are executed on the Flower platform [151]. All the results are shown in Table 3.5 (can greatly reduce the training time due to faster convergence).

Moreover, since each client sends a local model to the server with a size of $\sim 2\text{MB}$ in each communication round, with a faster convergence by UDP-FL, the total bandwidth consumption can be reduced by more than 50%, e.g., 40GB bandwidth reduction for 10 clients involved in the FL.

3.5.5 Defense against Privacy Attacks

In the following, we will evaluate the performance of our UDP-FL against several common privacy attacks in the domain of federated learning, specifically Membership Inference Attack (MIA) [1], Data Reconstruction Attack (DRA) [124, 125], and Attribute Inference Attack (AIA) [5, 152]. Notably, we underscore the significance of DP’s core advantage lies in its ability to offer plausible deniability [153], maintaining the privacy defenses of UDP-FL against these attacks with strong indistinguishability. This fundamental attribute ensures that, regardless of an attack’s sophistication, the indistinguishability introduced by DP mechanisms significantly complicates the accurate reconstruction or direct association of any data with individual participants. Since these privacy attacks are primarily developed based on deterministic results, the evaluation against these attacks can only be based on several sampled random results instead of the entire output space. Thus, given the nature of randomized output by DP mechanisms, *the clients and data owners can still deny their inferred results* [153], even if the privacy attacks achieve high accuracy on some sampled random results.

Membership Inference Attacks. In our investigation, we meticulously assess the resilience of UDP-FL against three advanced Membership Inference Attacks (MIAs) using the CIFAR-10 dataset. The first MIA method we used is Shokri et al.[1] (We use Shokri to name this attack), which assesses how machine learning models reveal information about their training datasets. The second, the Likelihood Ratio Attack (LiRA) [2], utilizes shadow models to statistically ascertain if a data point was used in training. The third, the Canary attack, employs synthetic images, refined through iteration, to probe the model’s disclosure of training data characteristics. Our goal is to identify and mitigate potential privacy risks in the model’s outputs. We adopt the

evaluation setting from Canary [3] that maintains the True Positive Rate (TPR) at a stringent 0.01 to evaluate the False Positive Rate (FPR), Area Under the Curve (AUC), and Accuracy (ACC). This conservative approach aims to minimize false positives, reflecting the grave legal and ethical implications of erroneous membership inferences, underscoring the importance of a defense mechanism that accurately prevents unauthorized inferences while minimizing errors.

Table 3.6: Evaluation of the Shokri et al. [1], SOTA LiRA [2] and Canary [3] MIAs on CIFAR-10. TPR* denotes the TPR when FPR=0.01. The TPR*, ACC and AUC for Shokri et al. are 0.053, 0.710, and 0.757. The TPR*, ACC and AUC LiRA are 0.126, 0.651, and 0.716. The TPR*, ACC and AUC LiRA are 0.137, 0.649, and 0.719.

Attack	ϵ	DP-SGD			UDP-FL (Laplace baseline)			UDP-FL (Staircase)		
		TPR*	ACC	AUC	TPR*	ACC	AUC	TPR*	ACC	AUC
Shokri	2	0.009	0.502	0.493	0.011	0.509	0.508	0.009	0.506	0.501
	4	0.009	0.505	0.499	0.013	0.512	0.509	0.009	0.505	0.502
	6	0.013	0.510	0.509	0.014	0.505	0.501	0.009	0.504	0.499
	8	0.011	0.503	0.499	0.007	0.505	0.500	0.008	0.504	0.496
LiRA	2	0.013	0.506	0.497	0.009	0.506	0.501	0.008	0.504	0.491
	4	0.014	0.513	0.505	0.013	0.505	0.495	0.008	0.510	0.501
	6	0.017	0.514	0.511	0.011	0.505	0.493	0.009	0.510	0.504
	8	0.011	0.504	0.497	0.009	0.513	0.505	0.009	0.507	0.492
Canary	2	0.016	0.519	0.513	0.011	0.504	0.497	0.009	0.506	0.495
	4	0.016	0.510	0.504	0.020	0.512	0.498	0.013	0.512	0.509
	6	0.015	0.523	0.528	0.011	0.510	0.507	0.009	0.507	0.500
	8	0.010	0.521	0.518	0.015	0.516	0.507	0.012	0.511	0.509

Table 3.6 provides valuable insights into the effectiveness of various noise mechanisms in mitigating membership inference attacks. Both attacks show that the Staircase noise addition under UDP-FL consistently yields the lowest False Positive Rate (FPR), indicating superior privacy preservation against MIAs. The Laplace baseline also effectively reduces FPR, although not as consistently low as the Staircase, suggesting good but variable privacy protection. In terms of Accuracy and AUC, the Staircase and Laplace mechanisms demonstrate moderate success in maintaining model utility while ensuring privacy.

Data Reconstruction Attacks. We evaluate the data reconstruction attacks on UDP-FL on CIFAR-10. These attacks aim to reconstruct training data points from a target model. Instead of training a separate reconstruction model, we directly optimize synthetic inputs to match the gradient information obtained from the target model, following [4]. To evaluate multi-image attacks, we average the gradients from batches of up to 100 images before running the reconstruction. We assess the attack’s efficacy by measuring the mean squared error (MSE) and Structural Similarity (SSIM) between the reconstructed and original images.

The evaluation of data reconstruction attacks on CIFAR-10 demonstrates the efficacy of UDP-FL against DRA threats. Notably, across varying privacy budgets (ϵ values), Staircase noise consistently results in higher MSE and lower SSIM, substantially reducing the attackers’ ability to reconstruct original images accurately. While both DP-SGD and UDP-FL with Laplace baseline offer comparable levels of protection, evidenced by similar MSE and SSIM values. More evaluations on DRA are presented in the Appendix 3.8.5.

Table 3.7: Evaluation on the SOTA InvGrad DRA [4] on CIFAR-10. The MSE, PSNR and SSIM for the Non-private method are 1.7104, 9.79, and 0.0751, respectively.

ϵ	DP-SGD			UDP-FL (Laplace baseline)			UDP-FL (Staircase)		
	MSE	PSNR	SSIM	MSE	F1	SSIM	MSE	PSNR	SSIM
2	2.2646	8.51	0.0195	2.2686	8.63	0.0573	2.3399	8.37	0.0096
4	2.2058	8.69	0.0414	2.1840	8.68	0.0629	2.2405	8.39	0.0204
6	2.1532	8.76	0.0417	2.1532	8.83	0.0692	2.1910	8.54	0.0207
8	2.1463	8.78	0.0519	2.1290	8.95	0.0746	2.1832	8.73	0.0225

Attribute Inference Attacks. While DP is primarily designed to protect against privacy leaks like MIAs, its efficacy against AIAs is not empirically explored. We adopted the methodology outlined in [5] on UTKFace dataset which collects over 20,000 facial images with annotations for age, gender, and ethnicity [154]. In AIAs, the adversary possesses complete knowledge about the

Table 3.8: Evaluation on the SOTA AIA [5] on UTKFace. The Accuracy and F1-score for the Non-private method are 0.86 and 0.85, respectively.

ϵ	DP-SGD		UDP-FL (Laplace baseline)		UDP-FL (Staircase)	
	Acc	F1	Acc	F1	Acc	F1
2	0.8392	0.8239	0.8297	0.7987	0.8267	0.8124
4	0.8497	0.8408	0.8285	0.8248	0.8383	0.8260
6	0.8519	0.8439	0.8390	0.8096	0.8322	0.8131
8	0.8615	0.8544	0.8320	0.8169	0.8317	0.8174

target model, including its internal architecture and parameters. This allows for a more in-depth exploration of the model’s vulnerabilities, particularly in inferring sensitive attributes from the model’s outputs in a federated learning context. We extracted embeddings from the penultimate layer of our target model, utilizing them as inputs for a 2-layer MLP attack classifier. This classifier was trained with embeddings and known gender attributes from an auxiliary dataset. Based on the standard F1-score and accuracy, the evaluation focused on the attack model’s ability to predict gender accurately. From the results in Table 3.8, we can tell that UDP-FL with Staircase slightly reduces the attack inference accuracy compared to the other two baseline methods and non-private methods. Our findings validate the interplay between MIAs and AIAs in [123] that even if a system is resistant to direct MIA, it may still be vulnerable to AI attacks if the adversary has additional knowledge or capabilities, such as adversary’s ability to access the entire model or the ability to reconstruct data. However, such comprehensive adversary knowledge is typically unrealistic, as it assumes access levels that are difficult to achieve. As shown in Table 3.7, successful data reconstruction by adversaries is hard, reinforcing the strength of our approach against common attack vectors.

3.6 Discussion

Diverse DP mechanisms. Our experiments mainly evaluate the Staircase, Laplace and Gaussian mechanisms with UDP-FL. As a universal DP-FL work, UDP-FL is flexible and can be extended to incorporate more advanced DP mechanisms, such as the Matrix Variate Gaussian (MVG) mechanism [139], R^2 DP mechanism [82], DP Boosting [15, 40], and more. Moreover, UDP-FL is designed to be flexible and adaptable to various FL settings. One of the key aspects of this flexibility is the compatibility of our framework with other aggregation functions commonly used in FL, beyond the widely-used FedAvg and FedSGD algorithm [92, 122].

Support Diverse Aggregation Functions. Several aggregation functions have been proposed to address the limitations of FedAvg, such as dealing with non-IID data, mitigating the effects of stragglers, or improving convergence rates. Some of these alternative aggregation functions include Scaffold [155], FedMed [156], FedProx [157]. To show the possibility of integrating alternative aggregation functions into UDP-FL, we have conducted another experiments to evaluate it. We use a wind forecasting dataset [158], and train a simple CNN to predict hourly power generation up to 48 hours ahead at 7 wind farms. The baselines use NON-DP Scaffold aggregation functions in the FL frameworks. The sampling rate is 0.05, and the client number is 10, and clipped gradient value is 10.

From Figure 3.5, we can see UDP-FL with the Staircase mechanism still outperforms the STOA DP-FL-based NbAFL [114]. Moreover, when ϵ is smaller than 5, the performance of UDP-FL is better than FedAvg without DP. One possible reason is that each class of noise is optimal for one specific metric and Staircase is designed for ℓ_1 and ℓ_2 metrics, and the noise may serve as a regularization to mitigate the unbiased distribution of clients' data.




Fig. 3.5: Accuracy results on FL with the Scaffold aggregation [155]. On small training epochs, UDP-FL with the Staircase mechanism can achieve better accuracy more quickly.

DP Accounting Extensions. UDP-FL can also be extended for other accounting of differential privacy. Recent work has proposed using characteristic functions of the privacy loss random variable as an alternative approach for optimal privacy accounting [106, 107, 108]. This technique provides a natural composition similar to RDP while avoiding RDP’s limitations. UDP-FL’s architecture is flexible enough that it could potentially be extended to incorporate characteristic functions instead of just Rényi divergence. Specifically, the Harmonizer could be adapted to compute and track characteristic functions for each mechanism. The analytic Fourier accountant method could also replace the MA for conversion to (ϵ, δ) -DP. With these modifications, UDP-FL could achieve tighter accounting and flexibility by leveraging characteristic functions. The modularity of UDP-FL makes these extensions possible without changing the overall framework.

Future Works. The development of UDP-FL opens up several promising avenues for future research in DP-FL. For instance, adaptive privacy budget allocation strategies represent a significant area for improvement. Such approaches could dynamically adjust to changing data

distributions and client behaviors during the FL process. Another crucial direction is the integration of UDP-FL with other privacy-preserving techniques. Combining DP with secure multi-party computation or homomorphic encryption could provide multi-layered privacy protection [37]. This integration could address a broader range of privacy concerns and potentially mitigate some of the accuracy loss associated with DP alone.

3.7 Related Works

Differential Privacy and Rényi Differential Privacy. Differential privacy [9, 132] has been extensively studied and applied to various data analysis and machine learning tasks to ensure privacy protection. With the proposal of DP-SGD [10], more works focus on tightly tracking the privacy loss during a large number of training. With the formal definition of Rényi DP [11], tightly quantifying the privacy loss becomes more accessible. Works [82, 108, 133, 159, 160] discussing applying different mechanisms to preserve RDP.

In the context of FL, RDP has been employed to achieve more robust privacy protection. More works [110, 112] explored the use of RDP and shuffler [129, 130, 161] in FL systems to provide stronger privacy guarantees. This approach demonstrated improved privacy protection and utility tradeoffs compared to traditional DP-based methods. Geyer et. al. [115] propose an algorithm for client-sided differential privacy preserving federated optimization. Bhowmick et. al.[162] present practicable approaches to large-scale locally private model training that were previously impossible, showing theoretically and empirically that Bhowmick et. al.[162] can fit large-scale image classification and language models with little degradation in utility. Existing FL approaches either use secure multiparty computation (MPC) techniques, which are

vulnerable to inference, or differential privacy, which can lead to low accuracy given a large number of parties with relatively small amounts of data each. Truex et al. [163] present an alternative approach that utilizes both differential privacy and MPC to balance these tradeoffs. Li et al. [164] investigate the feasibility of applying differential privacy techniques to protect patient data in an FL setup. To avoid privacy leakage, Wu et al. [156] propose to add Rényi differential privacy (RDP) into FL-MAC. Girgis et al. [129, 130, 146] study privacy in a distributed learning framework, where clients collaboratively build a learning model through interactions with a server from whom needs privacy.

Differentially Private Federated Learning. DP-FL continues to evolve rapidly, with recent works addressing various challenges and proposing innovative solutions. Recent advancements in DP-FL have explored novel privacy mechanisms and frameworks. Building on this, Li et al. [165] developed FedMask, which preserves both data and model privacy through gradient masking and perturbation. The challenge of heterogeneity in FL has been addressed in recent DP-FL works. Wei et al. [126] proposed a personalized DP-FL framework that adapts to individual client characteristics while maintaining strong privacy guarantees. Similarly, Xu et al. [166] introduced FedCORP, a communication-efficient one-shot robust personalized FL framework that incorporates differential privacy. Zhu et al. [167] proposed an efficient DP-FL algorithm achieving optimal sample complexity with theoretical guarantees. The intersection of DP-FL with other privacy-preserving technologies has also been explored. Chen et al. [168] combined differential privacy with secure multi-party computation for enhanced privacy in collaborative learning. Fort et al. [169] investigated privacy amplification by iteration in FL, providing new insights into how iterative training processes can enhance privacy guarantees. Theoretical

advancements in DP-FL have also been made. Liu et al. [170] derived tighter privacy bounds for subsampled mechanisms in FL, improving upon existing privacy accounting methods and enabling more efficient use of privacy budgets. Ding et al.[171] proposed a novel approach that leverages LDP to achieve both communication efficiency and enhanced privacy in FL. Their method demonstrates improved model accuracy while maintaining strong privacy guarantees. Recent DP-FL research has focused on enhancing utility, efficiency, and privacy guarantees. Ding et al. [171] and Varun et al. [172] proposed local differential privacy approaches that improve communication efficiency, model accuracy, robustness against attacks, respectively. Addressing data heterogeneity, Luo et al. [173] combined differential privacy with multi-task learning for personalized models across diverse client data. Zhou et al. [174] developed tighter privacy composition bounds for FL, enabling more accurate privacy loss tracking over multiple training rounds. In the domain of large language models, Dagan et al. [175] proposed a DP-FL framework addressing the unique challenges of high-dimensional models. Sun et al. [176] introduced an adaptive differential privacy mechanism that dynamically adjusts privacy levels to optimize the privacy-utility trade-off during training. Zheng et al. [111] introduce federated f -DP, a new notion tailored explicitly to the federated setting, based on the framework of Gaussian differential privacy. Khanna et al. [177] outline a simple FL algorithm implementing DP to ensure privacy when training a machine learning model on data spread across different institutions.

Adversarial Attacks in Federated Learning. The evolution of DP-FL has been influenced by several key studies. Chen et al. [178] proposed generative models for private data generation, focusing on the effectiveness of DP in defending against model inversion and GAN-based

attacks. In the context of 5G networks, Liu et al. [179] developed a blockchain-based secure FL framework, enhancing privacy preservation for participants. Naseri et al. [180] conducted a comprehensive evaluation of Local and Central DP in FL, assessing their impact on privacy and robustness. Sun and Lyu [181] proposed FEDMD-NFDP, a federated model distillation framework incorporating a Noise-Free DP mechanism, effectively eliminating the risk of white-box inference attacks. Kerkouche et al. [182] introduced a new FL scheme, offering a balance between robustness, privacy, bandwidth efficiency, and model accuracy. Chen et al. [183] developed a decentralized, privacy-preserving global model training protocol for FL in P2P networks. Hossain et al. [184] demonstrated how DP could be exploited for stealthy and persistent model poisoning attacks in FL. Feng et al. [185] evaluated user-level DP in FL, specifically in the context of speech-emotion recognition systems. Lastly, Wang et al. [186] proposed a platform-free proof of FL consensus mechanism, focusing on sustainable blockchains and privacy protection in FL models. Salem et al. [123] proposed a game-based framework to unify definitions and analysis of privacy inference risks. They use reductions between games to relate notions like membership inference and attribute inference.

3.8 Proofs and Further Discussion

3.8.1 Proofs

Proof of Theorem 7

Proof. We first consider the define the distribution of the Staircase Mechanism.

$$f = \begin{cases} e^{-\rho\lambda}y & ||x||_1 \in [\rho\Delta, (\rho + \nu)\Delta] \\ e^{-(\rho+1)\lambda}y & ||x||_1 \in [(\rho + \nu)\Delta, (\rho + 1)\Delta] \end{cases}$$

From the Staircase definition, when $x < 0$, we will have $f(x - \Delta) = e^{-\lambda}f(x)$. By direct computation, we will get the following results.

$$\begin{aligned} \int_{-\infty}^0 g(x)dx &= \int_{-\infty}^0 f_v^\alpha(x) f_v^{(1-\alpha)}(x - \Delta)dx \\ &= \int_{-\infty}^0 f_v^\alpha(x) f_v^{(1-\alpha)}(x) e^{-(1-\alpha)\lambda}dx = e^{-(1-\alpha)\lambda} \int_{-\infty}^0 f(x)dx \\ &= \frac{1}{2} e^{-(1-\alpha)\lambda} \end{aligned}$$

when $x > \Delta$, we will have $f(x - \Delta) = e^{\lambda}f(x)$. Thus, $g(x) = f_v^\alpha(x) f_v^{(1-\alpha)}(x - \Delta)$.

$$\begin{aligned} \int_{\Delta}^{+\infty} g(x)dx &= \int_{\Delta}^{+\infty} f_v^\alpha(x) f_v^{(1-\alpha)}(x - \Delta)dx \\ &= \int_{\Delta}^{+\infty} f_v^\alpha(x) f_v^{(1-\alpha)}(x) e^{(1-\alpha)\lambda}dx = e^{(1-\alpha)\lambda} \int_{\Delta}^{+\infty} f_v(x)dx \\ &= e^{(1-\alpha)\lambda} \int_0^{+\infty} f_v(x)dx - e^{(1-\alpha)\lambda} \int_0^{\Delta} f_v(x)dx \\ &= \frac{1}{2} e^{-\lambda} e^{(1-\alpha)\lambda} \end{aligned}$$

The last case will be $0 < x < \Delta$. We will define the cases below.

$$\text{sgn}\left(\frac{1}{2} - \nu\right) = \begin{cases} 0 & \nu < \frac{1}{2} \\ 1 & \nu \geq \frac{1}{2} \end{cases} \quad (3.12)$$

$$\begin{aligned} f_\nu(x - \Delta) &= f_\nu(\Delta - x) \int_0^\Delta g(x) dx \\ &= \int_0^\Delta f_\nu^\alpha(x) f_\nu^{(1-\alpha)}(\Delta - x) dx \\ &= \{(e^{(\alpha-1)\lambda} + e^{-\alpha\lambda})(1 - \nu) + |2\nu - 1|e^{-\text{sgn}(\frac{1}{2}-\nu)\lambda}\} \Delta \beta(\nu) \\ &= \frac{1}{2}e^{(\alpha-1)\lambda} + \frac{1}{2}e^{-\alpha\lambda} + \{(e^{(\alpha-1)\lambda} + e^{-\alpha\lambda})(1 - \nu) \\ &\quad + |2\nu - 1|e^{-\text{sgn}(\frac{1}{2}-\nu)\lambda}\} \frac{1 - e^{-1}}{2(\nu + e^{-\lambda}(1 - \nu))} \end{aligned}$$

Thus, this completes the proof. $\square$

Proof of Theorem 13

Proof. Given loss function whose length is m after T training iteration, the true count and noisy count of parameter as w_i and w_i^* , respectively, we have:

$$\begin{aligned} T\mathbb{E}\left(\sum_{i=1}^m |w_i^* - w_i|\right) &= Tm\mathbb{E}|w_i + \mathcal{N}(0, \sigma^2) - w_i| \\ &= Tm \int_{-\infty}^{\infty} |x| f(x) dx = Tm \int_0^{\infty} 2x \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{x^2}{2\sigma^2}} dx \\ &= \sqrt{\frac{2m^2\sigma^2T^2}{\pi}} \end{aligned}$$

where $f(x)$ is the PDF of the Gaussian distribution, and σ is the noise multiplier for Gaussian mechanism computed by UDP-FL.

Thus, this completes the proof. $\square$

Proof of Theorem 14

Proof. Similarly to Proof in 3.8.1, we will calculate

$$\begin{aligned} T\mathbb{E}\left(\sum_{i=1}^m |w_i^* - w_i|\right) &= Tm \int_{-\infty}^{\infty} |x|f(x)dx \\ &= Tm \int_{-\infty}^{\infty} |x| - 2be^{-\frac{|x|}{b}} dx = Tmb \end{aligned}$$

where b is the noise multiplier for Laplace mechanism computed by UDP-FL. Thus, this completes the proof. $\square$

Proof of Theorem 6

Proof. Per Theorem 7, the upper bound of the RDP is $\frac{1}{2}e^{(\alpha-1)\lambda} + \frac{1}{2}e^{-\alpha\lambda} + \{(e^{(\alpha-1)\lambda} + e^{-\alpha\lambda})(1 - \nu) + |2\nu - 1|e^{-\text{sgn}(\frac{1}{2}-\nu)\lambda}\} \frac{1-e^{-1}}{2(\nu+e^{-\lambda}(1-\nu))}$. Also, we know that the optimal value of ν is $\frac{1}{1+e^{\lambda/2}}$.

Thus, we will set the upper bound to be less than the privacy budget γ .

$$\begin{aligned} L(x) &= \frac{1}{2}e^{(\alpha-1)\lambda} + \frac{1}{2}e^{-\alpha\lambda} + \{(e^{(\alpha-1)\lambda} + e^{-\alpha\lambda})(1 - \nu) \\ &\quad + |2\nu - 1|e^{-\text{sgn}(\frac{1}{2}-\nu)\lambda}\} \frac{1 - e^{-1}}{2(\nu + e^{-\lambda}(1 - \nu))} \end{aligned}$$

To tightly preserve the privacy, we generate the variable from $f(\lambda, \Delta, 1/2)$, we have

$$\begin{aligned}
 L(x) &= \frac{1}{2}e^{(\alpha-1)\lambda} + \frac{1}{2}e^{-\alpha\lambda} + \{(e^{(\alpha-1)\lambda} + e^{-\alpha\lambda})(1 - \frac{1}{2}) \\
 &\quad + |2\frac{1}{2} - 1|e^{-\text{sgn}(\frac{1}{2}-\frac{1}{2})\lambda}\} \frac{1 - e^{-1}}{2(\frac{1}{2} + e^{-\lambda}(1 - \frac{1}{2}))} \\
 &= \frac{1}{2}e^{(\alpha-1)\lambda} + \frac{1}{2}e^{-\alpha\lambda} + (\frac{1}{2}e^{(\alpha-1)\lambda} + \frac{1}{2}e^{-\alpha\lambda}) \frac{1 - \frac{1}{e}}{1 + \frac{1}{e^\lambda}} \\
 &< (\frac{1}{2}e^{(\alpha-1)\lambda} + \frac{1}{2}e^{-\alpha\lambda}) (\frac{2 - 1/e + e^{-\lambda}}{1 + e^{-\lambda}}) \\
 &< e^{(\alpha-1)\lambda} + e^{-\alpha\lambda} \\
 &< e^{(\alpha-1)\lambda} + 1
 \end{aligned}$$

Let $L(x) = \gamma$, we have $\lambda = \frac{\log(\gamma-1)}{\alpha-1}$.

Thus, this completes the proof. $\square$

Proof of Theorem 8

Proof. Similar to the proof in Appendix 3.8.1, we will calculate

$$\begin{aligned}
 \int_0^{+\infty} xf(x)dx &= \int_0^\Delta xf(x)dx + \int_\Delta^{+\infty} xf(x)dx \\
 &= \int_0^\Delta xf(x)dx + \int_0^{+\infty} (x + \Delta)f(x + \Delta)dx \\
 &= \int_0^\Delta xf(x)dx + e^{-\lambda} \int_0^{+\infty} xf(x)dx + \Delta e^{-\lambda} \int_0^{+\infty} f(x)dx
 \end{aligned}$$

Thus, we have

$$\begin{aligned}
mT\mathbb{E}[|x|] &= 2 \int_0^{+\infty} xf(x)dx \\
&= \frac{mT}{1-e^{-\lambda}}(v^2\Delta^2 + e^{-\lambda}\Delta^2 - e^{-\lambda}v^2\Delta^2 + \Delta e^{-\lambda})
\end{aligned}$$

where $f(x)$ is the PDF of the Staircase mechanism, and ρ, v is the noise multiplier for Staircase mechanism computed by UDP-FL.

Thus, this completes the proof. $\square$

Proof of Theorem 9

Proof. We will prove that if the original system converges in K rounds the expectation of the number of rounds $\bar{K}$ needed for $\tilde{\theta}_k$ to approach 1 ($w_n \rightarrow w$) increases by $O(\sigma^2)$ where for the staircase noise $\sigma = \Delta \cdot \frac{e^{\epsilon/2}}{e^{\epsilon}-1}$ [118]. We'll assume that the non-perturbed version converges in K rounds and find k in terms of K . Since UDP-FL on client machines exerts the mode connectivity algorithm when searching for optima, we need to investigate diffusion mechanism of the perturbed parameter $\tilde{\theta}$ and investigating how it affects the curve function $\phi_{\tilde{\theta}}(p)$.

Given the perturbed gradient update:

$$\tilde{\theta}_{t+1} = \tilde{\theta}_t - \eta(\nabla_{\theta}L(\tilde{\theta}_t) + \xi_t)$$

where ξ_t is the noise added by the staircase mechanism. The curve function with perturbed

parameter $\tilde{\theta}$ is:

$$\phi_{\tilde{\theta}}(p) = \begin{cases} 2(p\tilde{\theta} + (0.5 - p)w_1), & 0 \leq p \leq 0.5 \\ 2((p - 0.5)w_2 + (1 - p)\tilde{\theta}), & 0.5 \leq p \leq 1 \end{cases}$$

The loss function evaluated at the curve with perturbed parameter:

$$L(\tilde{\theta}) = \mathbb{E}_{p \sim U(0,1)}[\ell(\phi_{\tilde{\theta}}(p))]$$

Let's analyze the expected number of rounds k for the noisy parameter $\tilde{\theta}_k$ to approach the target value (similar to the convergence of θ in K rounds). Gradient Descent with Noise:

$$\tilde{\theta}_{t+1} = \tilde{\theta}_t - \eta(\nabla_{\theta} L(\tilde{\theta}_t) + \xi_t)$$

We know the non-perturbed algorithm converges in K rounds. For the noisy version, the convergence will be slower due to the added noise. The noise impact can be understood by examining how the noise affects the variance of the parameter updates over iterations. Let's denote the variance of the noise ξ_t as σ^2 . The variance of the parameter updates accumulates over iterations:

$$\text{Var}(\tilde{\theta}_k) = \sum_{t=1}^k \eta^2 \sigma^2$$

The noise introduces a diffusion term that slows down convergence, i.e., the number of iterations k needs to compensate for this diffusion. In the following we will show that extra term is bounded by $\frac{C\sigma^2}{\eta^2}$ where C is a constant that depends on the specific noise characteristics and the gradient landscape. The proof requires the following assumptions on the clients' loss functions

$\{F_k\}$ (standard assumption [187]).

Assumption 1. $\{F_k\}$'s are L -smooth: $F_k(v) \leq F_k(w) + (v - w)^T \nabla F_k(w) + \frac{L}{2} \|v - w\|_2^2, \forall v, w$.

Assumption 2. $\{F_k\}$'s are μ -strongly convex: $F_k(v) \geq F_k(w) + (v - w)^T \nabla F_k(w) + \frac{\mu}{2} \|v - w\|_2^2, \forall v, w$.

Assumption 3. Let ξ_t^k be sampled from the k -th device's data uniformly at random. The variance of stochastic gradients in each device is bounded: $E \|\nabla F_k(w_t^k, \xi_t^k) - \nabla F_k(w_t^k)\|^2 \leq \sigma_k^2, \forall k$.

Assumption 4. The expected squared norm of stochastic gradients is uniformly bounded, i.e., $E \|\nabla F_k(w_t^k, \xi_t^k)\|^2 \leq G^2, \forall k, t$.

Based on the findings from [187], the FedAvg algorithm's updates with only a subset of devices active can be explained as follows: Let w_t^k represent the model parameters on the k -th device at time step t , and let $\mathcal{I}_E$ denote the set of global synchronization steps, defined as $\mathcal{I}_E = \{nE \mid n = 1, 2, \dots\}$. If $t + 1$ is in $\mathcal{I}_E$, indicating a communication round, FedAvg engages all devices. For partial device participation, the updates proceed as:

$$v_k = w_k - \eta \nabla F(w_k, \xi_k), \quad \text{if } t + 1 \notin \mathcal{I}_E,$$

$$w_{t+1} = \sum_{k=1}^N p_k v_k, \quad \text{if } t + 1 \in \mathcal{I}_E.$$

Here, the variable v_k represents the immediate result of a single SGD step from w_k . The parameter w_{t+1}^k is viewed as the result after synchronization steps. In this context, we define two aggregated sequences: $v_t = \sum_{k=1}^N p_k v_t^k$ and $w_t = \sum_{k=1}^N p_k w_t^k$. The sequence v_{t+1} is derived from performing one SGD step on w_t . When $t + 1 \notin \mathcal{I}_E$, both sequences are unavailable.

If $t + 1 \in \mathcal{I}_E$, we can access w_{t+1} . For simplicity, we define $g_t = \sum_{k=1}^N p_k \nabla F_k(w_t^k)$ and $\bar{g}_t = \sum_{k=1}^N p_k \nabla F_k(w_t^k, \xi_t^k)$. Consequently, $v_{t+1} = w_t - \eta g_t$ and $\mathbb{E}[g_t] = \bar{g}_t$. Assuming that Assumptions 1 and 2 hold, and if $\eta_t \leq 1$, the following inequality is satisfied:

$$\begin{aligned} & \mathbb{E} [\|v_{t+1} - w^*\|^2] \\ & \leq (1 - \eta_t \mu) \mathbb{E} [\|w_t - w^*\|^2] + \eta_t^2 \mathbb{E} [\|\bar{g}_t - g_t\|^2] + 6L\eta_t^2 \Gamma + 2\mathbb{E} [w_t - w_t^k], \end{aligned}$$

where $\Gamma = F^* - \sum_{k=1}^N p_k F_k^* \geq 0$.

Assuming convergence of non-perturbed model when v_{t+1} follows the mode connectivity path (along θ_t), the path of UDP-FL will follow

$$\begin{aligned} & \mathbb{E} [\|\tilde{\theta}_{t+1} - w^*\|^2] \\ & \leq (1 - \eta_t \mu) \mathbb{E} [\|\tilde{\theta}_t - w^*\|^2] + \eta_t^2 \mathbb{E} [\|g_t - \bar{g}_t\|^2] + 6L\eta_t^2 \Gamma + 2\mathbb{E} [\tilde{\theta}_t - \bar{\theta}_t^k], \end{aligned}$$

where $\Gamma = F^* - \sum_{k=1}^N p_k \mathbb{E}_{p \sim U(0,1)} [\ell(\phi_{\bar{\theta}_t^k}(p))] \geq 0$. Only first term and last term will generate additional term when moving from non-perturbed to the perturbed path. Additional terms in both terms are variance of the noisy residue ζ_t which is $\Delta^2 \cdot \frac{e^\epsilon}{(e^\epsilon - 1)^2}$. $\square$

3.8.2 Rényi Privacy Accounting

The moment accountant technique, proposed by Abadi et al. [10], keeps track of the evaluations of the CGF at fixed locations. Leveraging Lemmas 5 and 6, it allows one to find the smallest ϵ given a desired δ or vice versa (Algorithm 7).

Algorithm 6: Mode Connectivity Curve Finding Procedure

Input: w_1, w_2 : Model parameters for two local minima

Output: Optimal curve parameters θ

- 1 Initialize θ as the midpoint: $\theta = 0.5(w_1 + w_2)$
- 2 Define loss function $\ell(w)$ (e.g., cross-entropy loss)
- 3 Define curve function $\phi_\theta(p)$:

$$\phi_\theta(p) = \begin{cases} 2(p\theta + (0.5 - p)w_1) & \text{for } 0 \leq p \leq 0.5, \\ 2((p - 0.5)w_2 + (1 - p)\theta) & \text{for } 0.5 < p \leq 1. \end{cases}$$

4 **while** θ not converges **do**

5 Sample $\hat{p}$ uniformly from $[0, 1]$

6 Compute loss: $L(\theta) = \ell(\phi_\theta(\hat{p}))$

7 Compute gradient: $\nabla_\theta L(\theta)$

8 Update θ using a gradient step: $\theta \leftarrow \theta - \eta \nabla_\theta L(\theta)$

9 **end**

10 **return** Optimal parameters θ

For instance, consider a mechanism $\mathcal{A}$ applying Gaussian noise with standard deviation σ to the output of each operation in DPSGD. The privacy of M is characterized by the parameter $\epsilon(\alpha, T)$, where α is the Rényi divergence order, and T is the number of operations applied to the dataset.

For Gaussian DPSGD, the privacy parameter $\epsilon(\alpha, T)$ is given by:

$$\epsilon(\alpha, T) = \frac{\alpha T}{2\sigma^2}$$

This equation shows how the privacy parameter ϵ is related to the Rényi divergence order α , the number of operations T , and the standard deviation of the Gaussian noise σ . By analyzing the privacy parameter $\epsilon(\alpha, T)$, one can determine the level of privacy provided by the mechanism M in Gaussian DPSGD. This analysis allows for the adjustment of parameters such as α , T , and σ to achieve the desired level of privacy, balancing privacy and utility in the

Algorithm 7: Rényi Accountant [108]

Input: Training iteration T , privacy budget ϵ , $\alpha \in [2, \infty)$
Output: Noise multiplier

```

1 Initialize noise multiplier
2 while noise multiplier is not optimal do
3   forall  $\alpha$  do
4     forall  $t \in T$  do
5       Compute Rényi divergence for current noise multiplier
6       Sum of all  $T$  rounds' Rényi divergence for current  $\alpha$ 
7       if the sum exceeds the privacy budget  $\epsilon$  then
8         Change a smaller noise multiplier
9       else if the sum is within the privacy budget then
10        Continue
11      Save the current noise multiplier and try a smaller noise multiplier until
        it is optimal
12    end
13  end
14 end
15 return the optimal noise multiplier

```

learning process.

3.8.3 On Convergence Assumptions

A summary of notations is listed in Table 3.2, which comprises constants typically employed in convergence analyses of gradient-based methods. These constants are considered as the key characteristics of the loss function F and are commonly assumed in various widely used cost functions [188], including the cross-entropy.

To establish the convergence of our non-Gaussian DP Federated Learning system (UDP-FL), we leverage a well-established result from the machine learning convergence literature. The provided Lemma [189] ensures that, under certain conditions for the neural network, the model converges to f_* . These conditions, namely *convexity* and *smoothness*, are satisfied for the commonly used cross-entropy loss function in backpropagation neural networks.

The convexity of cross-entropy is elucidated by examining its formulation in terms of Kullback-Leibler (KL) divergence: $H(p, q) = H(p) + D_{KL}(p||q)$. As $H(p)$ remains constant, attention is directed to the convex nature of the KL divergence component. Notably, KL divergence exhibits convexity for discrete pairs (p, q) , as demonstrated by (p_1, q_1) and (p_2, q_2) [190]. The proof of this convexity, utilizing the log-sum inequality [191], is provided here [192], offering a comprehensive understanding of KL divergence's convexity within the context of cross-entropy. Additionally, Berrada et al. [193] established the smoothness of cross entropy.




Fig. 3.6: Impact of Shuffler on UDP-FL with the privacy guarantee as Figure 3.4 (MNIST). We observe that: 1) the utility of UDP-FL can be marginally improved with the shuffler; 2) UDP-FL can converge slightly faster with the shuffler; and 3) UDP-FL using the Staircase mechanism performs the best.




Fig. 3.7: Impact of Shuffler on UDP-FL with the privacy guarantee as Figure 3.4 (Medical dataset, with similar observations to the MNIST dataset).

3.8.4 Privacy and Utility Analysis for Gaussian and Laplace Mechanisms

Table 3.3 summarizes the noise characteristics.

Gaussian Mechanism. The Gaussian mechanism is commonly used for achieving Rényi differential privacy [11, 108, 160]. To generate Gaussian noise based on Rényi differential privacy for federate learning, we can adjust the added noise multiplier according to the order α of the Rényi divergence. Specifically, we can use a scaled Gaussian noise distribution whose variance is proportional to the order of the Rényi divergence. Based on Theorem 11, the

Harmonizer in UDP-FL can transform the noise multiplier for the Gaussian mechanism to ensure tight privacy bound with the moments account for FL.

Theorem 11 (Proven in [160]). *Let $\epsilon > 0, \delta > 0, \alpha > 1$, for $\sigma^2 \geq \frac{\alpha\Delta^2}{2\epsilon}$, Gaussian mechanism $\mathcal{N}(\Delta, \sigma^2)$ satisfies (α, γ) -Rényi DP.*

Laplace Mechanism. Similarly, the Laplace mechanism can also be used in UDP-FL to support moments accountant for federated learning via Rényi differential privacy (via Theorem 12). Then, we first prove the RDP guarantee, and then state this new setting of moments accountant.

Theorem 12. *For any $\alpha > 1, \gamma > 0$, Laplace mechanism $\text{Lap}(\frac{\Delta}{b})$, satisfies (α, γ) -Rényi differential privacy, where*

$$\gamma = \frac{1}{b} \geq \frac{\Delta(\alpha - 1)}{\log\left(\frac{(2\alpha-1)\exp(\gamma(\alpha-1))+\alpha-1}{\alpha}\right)}.$$

Proof. Suppose we use the Laplace mechanism to get noisy query result. We firstly define Laplace distribution as $\Lambda(u, b) = \frac{1}{2b} \exp(\frac{-|x-u|}{b})$. Next, we will calculate the error bound of Rényi divergence between $\Lambda(0, b)$ and $\Lambda(\Delta, b)$.

$$\begin{aligned} \mathcal{D}_\alpha(\Lambda(0, b) || \Lambda(\Delta, b)) &= \frac{1}{\alpha - 1} \log \int_{-\infty}^{\infty} p(x)^\alpha q(x)^{1-\alpha} dx \\ &= \frac{1}{(\alpha - 1)(2\lambda)} \log \int_{-\infty}^0 \exp\left(\frac{\alpha x}{\lambda} + \frac{(1 - \alpha)(x - \Delta)}{\lambda}\right) dx \\ &\quad + \frac{1}{(\alpha - 1)(2\lambda)} \log \int_0^s \exp\left(\frac{-\alpha x}{\lambda} + \frac{(1 - \alpha)(x - \Delta)}{\lambda}\right) dx \\ &\quad + \frac{1}{(\alpha - 1)(2\lambda)} \log \int_s^{\infty} \exp\left(\frac{\alpha x}{\lambda} + \frac{(1 - \alpha)(x - \Delta)}{\lambda}\right) dx \\ &= \frac{1}{(\alpha - 1)} \log \left\{ \frac{\alpha}{2\alpha - 1} \exp\left(\frac{\Delta(\alpha - 1)}{\lambda}\right) + \frac{\alpha - 1}{2\alpha - 1} \exp\left(\frac{-\Delta\alpha}{\lambda}\right) \right\} \end{aligned} \tag{3.13}$$

Thus, this completes the proof. $\square$

This theorem will be used in Section 3.3.1 to check if the current noise multiplier will preserve the privacy budget or not.

Lemma 5. *The expectation of noise amplitude for UDP-FL with Gaussian mechanism is $\frac{2\sigma}{\sqrt{2\pi}}$ where $\sigma = \sqrt{\frac{\alpha\Delta_S^2}{2\epsilon_R}}$.*

Lemma 6. *The expectation of noise amplitude for UDP-FL with Laplace mechanism $f(\Delta, b)$ is Δ/b .*

Theorem 13. *The expectation of the ℓ_1 distance for the output model parameters preserved by UDP-FL with Gaussian mechanism after T training round is $\sqrt{\frac{2m^2\sigma^2}{\pi}}$ where m is the length of the loss function, σ is the noise multiplier computed by UDP-FL.*

Proof. See detailed proof in Appendix 3.8.1. $\square$

Theorem 14. *The expectation of the ℓ_1 distance for the output model parameters preserved by UDP-FL with Laplace mechanism after T training round is Tmb where m is the length of the loss function, b is the noise multiplier computed by UDP-FL.*

Proof. See detailed proof in Appendix 3.8.1. $\square$

3.8.5 Supplementary DRA Evaluations

We evaluate the data reconstruction attacks [124] on UDP-FL on CIFAR-10. These attacks aim to reconstruct training data points from a target model. The adversary first trains a separate reconstruction model on a dataset from a similar distribution as the target model's training data.

The reconstruction model learns to generate synthetic inputs that closely match real samples, using the predictions from the target model as feedback. By optimizing the synthetic inputs to minimally change the target model’s outputs, the reconstruction attack extracts information about the original training data. We assess the attack’s efficacy by measuring the mean squared error (MSE) and mean absolute error (MAE) between the reconstructed and original images.

Table 3.9: Evaluation for data reconstruction attacks.

ϵ	DP-SGD		UDP-FL (Laplace baseline)		UDP-FL (Staircase)		Non-Private	
	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
2	2.40	1.26	2.42	1.27	2.44	1.28	1.23	0.65
4	2.35	1.24	2.38	1.25	2.40	1.26	1.23	0.65
6	2.30	1.22	2.33	1.23	2.35	1.24	1.23	0.65
8	2.25	1.20	2.28	1.21	2.30	1.22	1.23	0.65

Table 3.9 emphasizes the effectiveness of the UDP-FL framework, particularly with its Staircase mechanism, in mitigating data reconstruction attacks. This configuration consistently exhibits slightly higher MSE and MAE compared to both DP-SGD and UDP-FL with Laplace baseline, suggesting a more robust defense against reconstruction attacks.

3.9 Conclusion

In this paper, we introduced UDP-FL, a novel framework for DP-FL that addresses the critical challenge of optimizing the tradeoff between privacy and accuracy. By harmonizing various DP mechanisms with the Rényi Accountant, UDP-FL achieves tighter privacy bounds and faster convergence compared to SOTA methods. Our experimental results demonstrate the superior performance of UDP-FL in terms of both privacy guarantees and model accuracy. Furthermore, we proposed a mode connectivity-based method for analyzing the convergence of DP-FL models,

providing valuable insights into the faster convergence. Through extensive evaluations, we also showed that UDP-FL exhibits substantial resilience against advanced privacy attacks, further validating the significant advancement in data protection in FL environments.

Chapter 4

DPED: Multi-Layer Noise Distillation for Privacy-Preserving Text Embeddings

Shuya Feng, Yuan Hong

4.1 Abstract

Training text embedding models under differential privacy constraints is challenging due to the high dimensionality of language data and the presence of rare, identifying linguistic features. We propose **DPED (Differentially Private Embedding Distillation)**, a framework that leverages teacher-student distillation with multi-layer noise injection to learn high-quality embeddings while providing differential privacy guarantees. DPED trains an ensemble of teacher models on disjoint subsets of sensitive text data, then transfers their knowledge to a student model through noisy aggregation at multiple layers. A rare-word-aware strategy adaptively handles infrequent words, improving privacy-utility trade-offs. Experiments on benchmark datasets demonstrate that DPED outperforms standard differentially private training methods, achieving substantially higher utility at the same privacy budget.

4.2 Introduction

Natural language data often contains sensitive information about individuals, posing privacy risks when used to train embedding models or other language representations [194]. Differential Privacy (DP) [132, 195] provides a formal framework to mitigate these risks by ensuring that the learned model does not inadvertently reveal details unique to any single training example. However, applying DP in language model training has proven difficult: standard DP training algorithms like DP-SGD [196] tend to substantially degrade the utility of learned representations [197], especially in settings with large vocabularies and uneven word frequencies [198]. The presence of rare words can lead to either excessive noise addition or inadvertent model memorization. This utility loss is pronounced for text embeddings, which need to capture fine-grained semantic relationships; the high dimensional embedding space and long-tailed word distribution exacerbate the challenges of private learning [199, 200, 201].

Recent research suggests that teacher-student distillation frameworks can achieve better privacy-utility trade-offs than direct gradient perturbation [202]. In particular, the PATE (Papernot et al. [203, 204]) introduced an ensemble of teacher models trained on disjoint data and a student model that learns from the teachers’ aggregated outputs. While PATE has been successfully applied to classification tasks with limited output domains, its application to high-dimensional embeddings for language data remains under-explored [205]. Naively extending teacher-student voting to vocabulary-sized outputs would incur high privacy costs, and the issues of rare word occurrences would persist. Our work differs from prior approaches in three key ways. First, we propose a multi-layer noise injection strategy that provides privacy guarantees while preserving more of the rich semantic information needed for high-quality embeddings.




Fig. 4.1: DPED Overview. Disjoint teachers add noise to their outputs, which are then aggregated; the student learns from these privatized signals.

Unlike previous approaches that add noise only to final outputs, we inject calibrated noise at both intermediate embedding layers and final prediction stages. Second, we introduce a rare-word-aware aggregation mechanism specifically designed to handle the privacy challenges posed by infrequent words. This component dynamically identifies queries involving rare tokens and applies adaptive privacy measures—either abstaining entirely or adding increased noise—to prevent privacy leakage while maintaining utility for common words. Third, we provide a comprehensive privacy analysis showing how our teacher partitioning and multi-layer approach leads to stronger privacy guarantees compared to equivalent single-model approaches, with tight bounds that quantify the privacy-utility benefits.

Our contributions are summarized as follows:

- We propose DPED, a novel framework for training differentially private text embeddings.

DPED combines teacher-student distillation with multi-layer noise injection, enabling the student model to learn from teachers' intermediate representations and outputs under strong privacy guarantees.

- We introduce a **rare-word-aware aggregation** strategy that improves privacy-utility trade-offs by treating low-frequency words and out-of-distribution queries with specially calibrated noise and thresholding.
- We provide a theoretical analysis proving that our method satisfies (ϵ, δ) -DP for the training dataset, with tight privacy bounds that demonstrate how partitioning data among teachers leads to stronger privacy guarantees compared to equivalent single-model approaches.

4.3 Related Work

Differential Privacy in Deep Learning. Differential Privacy [132, 195] provides a formal definition ensuring that an algorithm’s output does not reveal significant information about any individual training sample. DP-SGD [196] implements this by clipping per-sample gradients and adding Gaussian noise, enabling privacy-preserving neural networks including language models [206]. However, DP-SGD struggles with language tasks due to large output spaces and long-tailed distributions where rare token gradients are either heavily clipped or overwhelmed by noise [200]. Despite advances like the moments accountant for tracking privacy loss, the fundamental trade-off remains: strong privacy typically requires significant utility sacrifices for complex language tasks.

Teacher-student Distillation for Privacy. An alternative technology, Knowledge distillation framework [207] transfers information from models trained on sensitive data to a student model in a privacy-preserving manner. The PATE framework [203] trains teacher ensembles on disjoint private data subsets and uses their noisy aggregated votes to train a student model. Papernot et al.

[204] enhanced this approach with adaptive noise and Rényi differential privacy for handling larger class counts. While PATE succeeded with classification tasks, applying it to embedding learning introduces new challenges: the output space is significantly larger, making naive voting impractical, and the student needs to learn continuous representations rather than discrete labels. Our work extends this paradigm with richer distillation signals specifically designed for these embedding challenges.

Privacy-preserving Text Representations. Recent approaches have explored adding calibrated noise directly to word embeddings [199, 200], showing that privacy-utility trade-offs can be improved by adjusting noise according to word frequency [208]. However, these methods typically assume pre-trained embeddings that are subsequently sanitized, rather than training the embedding model itself with privacy guarantees. In contrast, our work focuses on learning the embedding model from scratch under DP constraints through a teacher-student process. Related efforts include federated learning with DP for language models [206] or DP-BERT [209, 210], though these approaches face similar challenges with vocabulary size and rare data that our method addresses through partitioning and distillation .

4.4 Methodology

Our goal is to train an embedding model on a sensitive text dataset D such that the model is differentially private with respect to D . We achieve this using a teacher-student framework with multiple points of noise injection.

4.4.1 Teacher Ensemble and Data Partitioning

We begin by splitting the private dataset $\mathcal{D}$ into N disjoint subsets $\mathcal{D}_1, \mathcal{D}_2, \dots, \mathcal{D}_N$ and training a teacher model T_i on each subset. Each teacher produces an intermediate representation $h_i(x)$ and output $y_i(x)$ for input x . This partitioning ensures each individual data record influences only one teacher model.

4.4.2 Multi-Layer Noisy Knowledge Aggregation

For each unlabeled query x , we collect outputs from all teachers and compute:

$$H(x) = \frac{1}{N} \sum_{i=1}^N h_i(x), \quad V(x) = \sum_{i=1}^N y_i(x) \quad (4.1)$$

where $H(x)$ is the average hidden representation and $V(x)$ is a vote count vector. To protect privacy, we inject noise into both:

$$\tilde{H}(x) = H(x) + n_h, \quad \tilde{V}(x) = V(x) + n_v \quad (4.2)$$

where $n_h \sim \mathcal{N}(0, \sigma_h^2 I_d)$ and $n_v \sim \mathcal{N}(0, \sigma_v^2 I_V)$. We then derive a privatized label $\hat{y}(x) = \arg \max_w \tilde{V}(x)[w]$. The student model is trained on tuples $\{x, \tilde{H}(x), \hat{y}(x)\}$ using a custom loss function:

$$\mathcal{L}_S = \frac{1}{M} \sum_{j=1}^M \left(\|h_S(x^j) - \tilde{H}(x^j)\|^2 + \mathcal{L}_{out}(y_S(x^j), \hat{y}(x^j)) \right) \quad (4.3)$$

This multi-layer guidance makes student training more data-efficient under noise constraints.

4.4.3 Rare-Word-Aware Aggregation

Low-frequency words are particularly risky: they can reveal unique user information if they appear in only a single teacher’s data. We introduce a threshold r on the teacher votes:

1. Compute $m = \max_w V(x)[w]$, the highest vote count for any token w .
2. If $m < r$, we either skip the query entirely (abstain) or add larger noise to $\tilde{V}(x)$.

In practice, we found abstaining on such rare patterns to be a simple and effective solution. This threshold mechanism amplifies privacy since outliers cannot sway the final student training if they do not meet a minimal teacher consensus.

The choice of r involves a trade-off. A higher r safeguards unique tokens by filtering out queries with insufficient teacher agreement, yet this can exclude many low-frequency words and reduce coverage. In contrast, a lower r preserves more queries but heightens privacy risk and may degrade performance if rare tokens are influenced by only one teacher.

4.5 Privacy Analysis

Our privacy analysis extends the PATE approach to multiple output components and uses composition across queries. We present the formal privacy guarantee of our method:

Theorem 15 (Privacy of DPED). *Let $\mathcal{A}$ be the DPED algorithm with N teachers trained on disjoint data subsets, noise parameters $\sigma_h \geq \frac{C_h \sqrt{2 \ln(1.25/\delta_q)}}{N \epsilon_q}$ for hidden representations and $\sigma_y \geq \frac{\sqrt{2 \ln(1.25/\delta_q)}}{\epsilon_q}$ for votes, rare-word threshold $r \geq 2$, and M unlabeled queries. Then algorithm $\mathcal{A}$ satisfies (ϵ, δ) -differential privacy, where ϵ and δ are derived from the per-query guarantees (ϵ_q, δ_q) using the moments accountant composition.*

Proof Sketch. The key insight is that a single data record affects at most one teacher’s outputs, and each teacher contributes to aggregation in limited ways. For each query x , the sensitivity of $H(x)$ is

bounded by $\frac{1}{N} \|h_k(x) - h'_k(x)\| \leq \frac{C_h}{N}$, where C_h is a constant (this can be enforced by clipping teacher representations). The sensitivity of $V(x)$ is at most 1 (one vote difference).

By calibrating noise according to these sensitivities, the Gaussian mechanism ensures that each query response satisfies (ϵ_q, δ_q) -DP. Using the moments accountant technique to compose across M queries, we achieve (ϵ, δ) -DP for the entire algorithm. The rare-word thresholding mechanism further strengthens privacy by filtering queries with insufficient teacher consensus, reducing the worst-case influence of any individual data point. The full detailed proof is provided in Appendix 4.7. $\square$

4.6 Experiments

4.6.1 Experimental Setup

Tasks and Datasets. We evaluate on: (1) word prediction using WikiText-2 (2M tokens), (2) sentiment classification with IMDb reviews, (3) WordSim-353 similarity, and (4) Google analogy completion.

Models. Each teacher model is a skip-gram architecture with 100-dimensional embeddings [211]. All systems use the same 4M-parameter skip-gram network (100-d embeddings). WikiText-2 is split into 10 shards (200k tokens each) to train the ten teachers; the student trains on 100k noisy queries. DP-SGD re-uses the architecture with norm-1 gradient clipping and Gaussian noise. PATE uses the same teacher ensemble but passes only noisy top-1 labels to the student (no hidden vectors).

4.6.2 Main Results: Utility vs Privacy

Table 1 shows performance across privacy levels. At $\epsilon \approx 1.0$, DPED achieves perplexity of 210.8 compared to 350.7 for DP-SGD and 290.4 for PATE, while showing significant improvements in WordSim similarity (0.57 vs. 0.45) and downstream sentiment accuracy (79.5% vs. 72.3%). As ϵ increases to 4.0, DPED nearly matches non-private performance (perplexity 130.7 vs. 120.5 non-private).

Notably, DPED at $\epsilon = 2$ outperforms DP-SGD at $\epsilon = 4$, demonstrating our method achieves




Fig. 4.2: Main experimental results on WikiText-2 (embedding training) and IMDb (downstream sentiment) under different privacy budgets. $\downarrow$ means lower is better, $\uparrow$ higher is better. Non-private is the upper bound with no noise. DP-SGD is training the embedding model with differentially private SGD. PATE is a baseline with teacher ensemble voting (single-layer). DPED is our full method with multi-layer distillation and rare-word-aware aggregator. Results are mean of 3 runs; standard deviations are shown for DP methods. Our approach consistently outperforms DP-SGD and PATE baselines in perplexity and downstream accuracy at equivalent (ϵ, δ) .

a given utility level at significantly lower privacy cost. The improvements in WordSim and analogies indicate that embeddings learned by DPED capture semantics better, likely due to the multi-layer guidance helping place words correctly in the embedding space.

4.6.3 Ablation Studies

Effect of Multi-Layer Distillation. We compare three variants: Full DPED (student gets both $\tilde{H}(x)$ and $\hat{y}(x)$), Output-only (student only gets $\hat{y}(x)$, like PATE), and Hidden-only (student gets $\tilde{H}(x)$ but not direct labels). At $\epsilon \approx 2$, combining hidden representation and output label guidance yields the best results (perplexity 150.3) compared to output-only (190.0) or hidden-only (250.4). We also observed faster convergence with multi-layer guidance, confirming that intermediate signals improve training efficiency under noise.

Table 4.1: Ablation on use of hidden representation in distillation ($\epsilon = 2$). Providing the student with the noisy hidden representation in addition to the noisy label significantly improves performance.

Student Training	PPL	WordSim	IMDb Acc
Output-only (like PATE)	190.0	0.56	80.2%
Hidden-only	250.4	0.49	75.5%
Full (output+hidden)	150.3	0.60	83.0%

Effect of Rare-Word Threshold. Disabling the rare-word threshold (setting $r = 1$) degrades performance substantially: perplexity increases from 210 to 240 at $\epsilon = 1$, with particularly pronounced negative effects on rare word representations. The thresholding mechanism improves both privacy and utility by preventing the student from learning spurious associations from infrequent tokens.

Varying Number of Teachers. We experimented with $N = 5$ and $N = 20$ teachers. With fewer teachers (5), performance degraded (PPL increased 10%) as noise per query had to be higher. With more teachers (20), we observed modest improvements (5-7% better PPL) as the privacy cost per query decreased, but computational costs increased with diminishing returns. $N = 10$ represents a good balance in our setting.

4.7 Detailed Privacy Analysis

In this section, we provide a more detailed analysis of the privacy guarantees of our DPED framework. We first formalize the notion of differential privacy used, then analyze how our multi-layer noise injection mechanism provides these guarantees.

4.7.1 Differential Privacy Background

A randomized algorithm $\mathcal{A}$ satisfies (ϵ, δ) -differential privacy if for any two neighboring datasets D and D' that differ in at most one record, and for all possible outputs $S \subseteq \text{Range}(\mathcal{A})$:

$$\Pr[\mathcal{A}(D) \in S] \leq e^\epsilon \cdot \Pr[\mathcal{A}(D') \in S] + \delta \quad (4.4)$$

The privacy parameter ϵ controls the strength of the privacy guarantee (smaller values indicate stronger privacy), while δ represents the probability of the guarantee failing.

4.7.2 Sensitivity Analysis

The sensitivity of our aggregation mechanisms directly affects the amount of noise needed to ensure privacy. Here we derive precise bounds:

Hidden Representation Sensitivity

For any neighboring datasets D and D' differing in one record, this record affects at most one teacher model, say T_k . The sensitivity of the average hidden representation $H(x)$ is:

$$\Delta H = \max_{D, D'} \|H_D(x) - H_{D'}(x)\| \quad (4.5)$$

$$= \max_{D, D'} \left\| \frac{1}{N} (h_k(x) - h'_k(x)) \right\| \leq \frac{C_h}{N} \quad (4.6)$$

where C_h is the maximum possible change in a single teacher's hidden representation due to one training example. In practice, we enforce this bound by clipping hidden representations to have norm at most C_h .

Vote Count Sensitivity

Similarly, for the vote count vector $V(x)$:

$$\Delta V = \max_{D, D'} \|V_D(x) - V_{D'}(x)\|_1 \leq 1 \quad (4.7)$$

This is because at most one teacher's vote can change, affecting at most one coordinate of $V(x)$ by at most 1.

4.7.3 Privacy of Multi-Layer Noise Mechanism

For a single query x , our mechanism releases both $\tilde{H}(x)$ and $\tilde{V}(x)$. By the Gaussian mechanism theorem and our sensitivity bounds, adding Gaussian noise with standard deviation $\sigma_h \geq \frac{C_h \sqrt{2 \ln(1.25/\delta_q)}}{N \epsilon_q}$ to $H(x)$ ensures that $\tilde{H}(x)$ is $(\epsilon_q/2, \delta_q/2)$ -DP. Similarly, adding noise with $\sigma_y \geq \frac{\sqrt{2 \ln(1.25/\delta_q)}}{\epsilon_q}$ to $V(x)$ ensures that $\tilde{V}(x)$ is $(\epsilon_q/2, \delta_q/2)$ -DP.

By basic composition, releasing both $\tilde{H}(x)$ and $\tilde{V}(x)$ provides (ϵ_q, δ_q) -DP for a single query.

4.7.4 Composition Across Multiple Queries

For a sequence of M queries, we use the moments accountant technique [196] to track the privacy loss more tightly than basic composition would allow. The moments accountant keeps track of the log moments of the privacy loss random variable, allowing for a more precise analysis of privacy cost accumulation.

For M queries, the total privacy cost using the moments accountant is approximately:

$$\epsilon \approx \epsilon_q \sqrt{2M \ln(1/\delta)} + M \epsilon_q (e^{\epsilon_q} - 1) \quad (4.8)$$

In practice, we calculate this precisely using the moments accountant implementation.

4.7.5 Privacy Amplification via Thresholding

Our rare-word-aware thresholding mechanism further improves privacy by abstaining on queries where consensus is low. This effectively reduces the worst-case influence of any single data point.

When using threshold $r > 1$, we ensure that results are only released when at least r teachers agree. This provides a form of privacy amplification: if a word appears in only one teacher’s training data, that word cannot affect the final student model through distillation because it will be filtered by the threshold mechanism.

4.7.6 Theorem: DPED Privacy Guarantee

Theorem 16. *The DPED algorithm with N teachers, noise scales σ_h and σ_y for hidden representations and output votes respectively, rare-word threshold $r \geq 2$, and processing M queries satisfies (ϵ, δ) -differential privacy with respect to the training dataset.*

The proof follows from the composition of privacy guarantees across multiple queries using the moments accountant, with each query’s response satisfying (ϵ_q, δ_q) -DP as shown above. The detailed proof is omitted for space constraints but follows the analysis structure outlined in this section.

4.8 Algorithm Pseudocode

For completeness, we provide detailed pseudocode for the DPED training procedure:

4.9 Conclusion

We presented DPED, a method for learning differentially private text embeddings via teacher-student distillation with multi-layer noise injection. Our approach achieves substantially better utility than prior methods at equivalent privacy levels through its dual-signal knowledge transfer and rare-word-aware

Algorithm 8: DPED: Differentially Private Embedding Distillation

Input: Private dataset $\mathcal{D}$, unlabeled query set $\mathcal{U}$, number of teachers N ,
noise scales σ_h, σ_y , rare-word threshold r

Output: Privacy-preserving student model S

/ Train teacher models on disjoint splits */*

- 1 Partition $\mathcal{D}$ into $\{\mathcal{D}_i\}_{i=1}^N$;
- 2 **for** $i \leftarrow 1$ **to** N **do**
- 3 $\lfloor$ Train teacher T_i on $\mathcal{D}_i$;
- /* Prepare student training data (noisy aggregation) */*
- 4 Initialize $\mathcal{T}_{\text{student}} \leftarrow \emptyset$;
- 5 **foreach** $x \in \mathcal{U}$ **do**
- /* Collect each teacher's output */*
- 6 **for** $i \leftarrow 1$ **to** N **do**
- 7 $\lfloor$ Collect $(h_i(x), y_i(x))$;
- 8 $H(x) \leftarrow \frac{1}{N} \sum_{i=1}^N h_i(x)$;
- 9 $V(x) \leftarrow \sum_{i=1}^N y_i(x)$;
- 10 $w^* \leftarrow \arg \max_w V(x)[w]$; *// top-voted word*
- 11 $m \leftarrow V(x)[w^*]$; *// vote count*
- 12 **if** $m < r$ **then**
- 13 $\lfloor$ Skip x ; *// low consensus*
- 14 **else**
- 15 Sample $n_h \sim \mathcal{N}(0, \sigma_h^2 I_d)$;
- 16 Sample $n_v \sim \mathcal{N}(0, \sigma_y^2 I_V)$;
- 17 $\tilde{H}(x) \leftarrow H(x) + n_h$;
- 18 $\tilde{V}(x) \leftarrow V(x) + n_v$;
- 19 $\hat{y}(x) \leftarrow \arg \max_w \tilde{V}(x)[w]$;
- 20 $\lfloor$ Add $(x, \tilde{H}(x), \hat{y}(x))$ to $\mathcal{T}_{\text{student}}$;
- /* Train student model */*
- 21 Initialize student model S ;
- 22 Train S on $\mathcal{T}_{\text{student}}$ to minimize:

$$\mathcal{L}_S = \frac{1}{|\mathcal{T}_{\text{student}}|} \sum_{(x, \tilde{H}, \hat{y}) \in \mathcal{T}_{\text{student}}} \|h_S(x) - \tilde{H}\|^2 + \mathcal{L}_{\text{out}}(y_S(x), \hat{y})$$

return S ;




Fig. 4.3: Cosine similarity of DPED and DP-SGD regarding different word frequency.

aggregation. Theoretical analysis provides formal privacy guarantees, while empirical results demonstrate DPED significantly narrows the performance gap between private and non-private NLP models. Future work could explore more efficient adaptations for larger Transformer architectures, building on DPED's foundation for privacy-preserving language representations.

Chapter 5

Usable Machine Learning Model Training with Differential Privacy

Shuya Feng, Seonhye Park, Doowon Kim, Yuan Hong

5.1 Abstract

Differential privacy (DP) has emerged as a fundamental framework for protecting sensitive data in machine learning applications, yet significant gaps remain between theoretical privacy guarantees and practical deployment. This paper presents a comprehensive multi-stakeholder analysis of Differentially Private Stochastic Gradient Descent (DP-SGD) adoption, examining perspectives across three key stakeholder groups: data contributors, machine learning engineers, and model recipients. We conduct a mixed-methods study employing online surveys, interactive educational interventions, and controlled experiments to investigate privacy-utility trade-offs across distinct data modalities (image vs. text classification) and implementation paradigms (white-box vs. black-box approaches). Our findings reveal stakeholder-specific concerns and requirements that significantly influence DP-SGD adoption decisions in production environments. We identify critical factors affecting deployment success, including utility threshold tolerances, privacy parameter interpretability, and implementation complexity preferences. These results provide actionable insights for improving DP-SGD usability and inform the design of stakeholder-centered privacy-preserving machine learning systems.

5.2 Introduction

Machine learning systems derive their predictive power from data, with performance typically improving as dataset size and richness increase. However, the collection and processing of personal and sensitive data has generated significant privacy concerns, amplified by high-profile data breaches, unauthorized data usage incidents, and the implementation of stringent data protection regulations such as the European Union’s General Data Protection Regulation (GDPR) and the California Consumer Privacy Act (CCPA). In response to these challenges, differential privacy (DP) has emerged as a mathematically rigorous framework for preserving individual privacy while enabling statistical analysis. By introducing carefully calibrated noise into computations, DP provides formal guarantees that the output of an analysis does not reveal the presence or absence of any individual record.

Despite its theoretical foundations and proven privacy guarantees, differential privacy remains underutilized in production machine learning systems. This adoption gap stems from multiple interconnected factors. First, practitioners find differentially private training methods, particularly Differentially Private Stochastic Gradient Descent (DP-SGD), technically challenging to implement and tune effectively. The method requires careful selection of hyperparameters—including noise multipliers, gradient clipping norms, and privacy budgets—that directly impact model utility. Second, organizational decision-makers often lack sufficient understanding of how DP translates into concrete compliance benefits or acceptable performance trade-offs. Third, individual data contributors may harbor skepticism about privacy-preserving technologies, questioning whether algorithmic noise truly protects their identities and whether such systems can be trusted.

To address these multifaceted challenges, we argue that a comprehensive, multi-stakeholder investigation of differential privacy adoption is essential. This work focuses specifically on DP-SGD applied to social media data, where individual users contribute personal content including text posts, images, and behavioral data. Social media platforms represent a particularly compelling use case: they aggregate vast quantities of user-generated content for applications including personalization, targeted advertising,




Fig. 5.1: Ecosystem of multi-stakeholder with DP-SGD.

and content analysis, creating substantial privacy risks through potential membership inference attacks, unintended data exposure, and misuse of personal information.

The deployment of DP-SGD in social media analytics involves three distinct stakeholder groups, each with unique motivations, constraints, and decision criteria:

1. **Data Contributors** are individual users who provide personal data in exchange for platform services such as personalized recommendations or automated content tagging. Their primary concern centers on trust: “Can I rely on the platform to protect my data from misuse or unauthorized disclosure?” While traditional privacy assurances (e.g., “we anonymize your data”) may be insufficient or unclear, formal privacy guarantees through DP-SGD could potentially increase data sharing willingness. Critical questions include how users perceive noise injection across different data modalities (images versus text) and their understanding of privacy parameter trade-offs.
2. **ML Engineers** are responsible for model design, training, and optimization. They face direct technical challenges from DP-SGD implementation, including accuracy degradation due to privacy noise, complex hyperparameter selection, and privacy budget management. Domain-specific considerations further complicate deployment: text-based tasks often involve smaller datasets and sparser feature representations that are more susceptible to noise, while image classification tasks may demonstrate greater robustness but still experience performance degradation. Engineers must

navigate these technical constraints while maintaining acceptable model utility.

3. **Model Recipients** include organizations, product managers, and technical teams that deploy trained models without directly contributing data or performing training. They require models that meet specific performance thresholds for business applications while satisfying regulatory compliance requirements and organizational privacy policies. DP-SGD can serve as evidence of privacy-by-design principles and regulatory compliance, but decision-makers must evaluate whether privacy-related utility losses justify deployment costs.

The interaction between these stakeholder groups is critical for successful DP-SGD adoption. Insufficient data contributor trust may limit dataset availability, inadequate engineering support may result in improper implementation or avoidance of DP methods, and recipient skepticism may prevent deployment of privacy-preserving models in favor of higher-accuracy alternatives. Conversely, solutions that simultaneously address trust, feasibility, and performance requirements across all stakeholder groups can enable broader DP-SGD deployment.

5.2.1 Research Questions

We structure our investigation around four research questions that systematically address stakeholder-specific concerns while collectively spanning the DP-SGD adoption challenge space:

- **RQ1 (Data Contributors - Trust):** How do data contributors perceive DP-SGD privacy guarantees, and to what extent does DP-SGD implementation increase their willingness to share data?

This question examines user-level privacy perceptions, investigating whether formal privacy guarantees translate into increased trust and data sharing behavior, and identifying factors that influence user confidence in privacy-preserving systems.

- **RQ2 (ML Engineers - Implementation):** What practical challenges do ML engineers encounter

when integrating DP-SGD into model training workflows, and what support mechanisms would mitigate implementation complexity?

This question focuses on technical implementation barriers, seeking to identify specific pain points in DP-SGD deployment (hyperparameter optimization, model quality management, computational overhead) and potential solutions.

- **RQ3 (Model Recipients - Performance):** How does DP-SGD impact model performance and deployment utility, and how do model recipients evaluate privacy-accuracy trade-offs?

This question addresses organizational deployment concerns, quantifying DP-SGD performance costs in realistic scenarios and understanding acceptable degradation levels for decision-makers across different application contexts.

- **RQ4 (Integration - Adoption Strategies):** What strategies can reconcile multi-stakeholder requirements to improve trust, simplify implementation, and optimize performance in DP-SGD deployment?

Building on insights from RQ1-RQ3, this question synthesizes findings into actionable recommendations, design principles, and best practices that facilitate practical DP-SGD adoption.

5.2.2 Contributions

This work provides several key contributions to privacy-preserving machine learning research and practice:

- **Comprehensive Multi-Stakeholder DP-SGD Analysis:** We present the first systematic investigation of DP-SGD adoption barriers from the perspectives of data contributors, ML engineers, and model recipients. Through empirical analysis of user attitudes, developer experiences, and organizational requirements, we identify specific factors that motivate or inhibit DP-SGD deployment across different contexts and stakeholder groups.

- **Evidence-Based DP-SGD Implementation Guidelines:** Based on controlled experiments and stakeholder feedback, we develop practical recommendations for DP-SGD configuration and deployment. These guidelines address privacy parameter selection (ϵ, δ) across application scenarios, hyperparameter optimization strategies (learning rate scheduling, gradient clipping) for minimizing utility loss [212], and architectural modifications that enhance noise robustness. Our recommendations provide practitioners with systematic approaches to effective DP-SGD implementation.
- **Trust and Usability Enhancement Framework:** We propose actionable strategies for improving DP-SGD trustworthiness and usability across the deployment ecosystem. For data contributors, we develop communication approaches that convey privacy guarantees through accessible explanations and visualization tools. For organizations, we provide integration guidance for incorporating DP-SGD into governance and compliance frameworks with transparent privacy reporting. For the research community, we identify critical usability improvements including automated privacy accounting and intuitive parameter selection interfaces that reduce practitioner barriers.

These contributions address both technical and human factors in DP-SGD deployment, providing privacy researchers, ML practitioners, and policymakers with evidence-based guidance for principled and practical differential privacy implementation.

5.3 Background and Related Works

5.3.1 Differential Privacy and DP-SGD

Differential Privacy (DP) provides a rigorous mathematical framework for quantifying privacy loss in data analysis. A randomized mechanism M with domain $\mathcal{D}$ satisfies (ϵ, δ) -differential privacy if for all pairs of adjacent datasets D, D' (differing in exactly one record) and for all measurable subsets O of outputs:

$$\Pr[M(D) \in O] \leq e^\epsilon \Pr[M(D') \in O] + \delta, \quad (5.1)$$

where $\epsilon \geq 0$ controls the privacy loss and $\delta \geq 0$ represents a failure probability that allows rare events in which privacy may be breached. When $\delta = 0$, the guarantee is termed “pure DP,” whereas allowing a small $\delta > 0$ provides “approximate DP,” permitting a small fraction of potential privacy breaches.

Intuitively, DP ensures that the presence or absence of any single record in the dataset has minimal impact on the mechanism’s outputs. This property implies that an adversary cannot reliably infer whether any individual’s data was included in the training set by inspecting the model’s behavior or outputs.

Differentially Private Stochastic Gradient Descent (DP-SGD) [10] adapts standard stochastic gradient descent to ensure that model parameters released after training satisfy differential privacy guarantees. The algorithm operates through three key mechanisms:

Gradient Clipping. In each training step, a random mini-batch of examples is selected through sub-sampling, and the gradient for each example g_i is computed. Each per-example gradient is then clipped to a maximum ℓ_2 norm C :

$$\tilde{g}_i = \frac{g_i}{\max(1, \|g_i\|_2/C)}. \quad (5.2)$$

This clipping operation ensures that no single sample’s gradient can exceed the threshold C , thereby limiting any individual’s influence on the parameter update.

Noise Injection. The clipped gradients are averaged across the mini-batch, and noise sampled from a Gaussian distribution (typically $\mathcal{N}(0, \sigma^2 C^2 I)$) is added before the parameter update step.

Privacy Accounting. Across multiple training iterations, each noise addition accumulates privacy cost. Advanced accounting methods, such as the moments accountant [10], provide tight bounds on the total

privacy expenditure to produce a final (ϵ, δ) guarantee for the entire training process.

By adjusting the noise scale σ and the number of training iterations, DP-SGD enables practitioners to navigate the fundamental trade-off between model utility and privacy protection. Stricter privacy requirements (smaller ϵ) generally necessitate increased noise or reduced training iterations, resulting in diminished model accuracy.

Several practical challenges arise when implementing DP-SGD:

- **Privacy-Utility Trade-off:** Increased noise injection (stronger privacy) directly degrades model performance. Empirical studies demonstrate that reducing ϵ from 8 to 2 in computer vision tasks can decrease accuracy by 6-8 percentage points [10].
- **Computational Overhead:** DP-SGD requires per-sample gradient computation and clipping, which can increase training time by a factor of 2-8 $\times$ compared to standard SGD, depending on implementation and hardware constraints.
- **Hyperparameter Sensitivity:** The choice of clipping threshold C , noise multiplier σ , and privacy parameters (ϵ, δ) significantly affects both privacy guarantees and model performance, requiring careful tuning.

To address these limitations, several algorithmic improvements have been developed:

Ghost Clipping reduces memory and computational overhead by avoiding explicit storage of per-example gradients. The technique performs gradient norm computation in a forward pass and applies clipping implicitly through loss scaling in the backward pass, effectively halving memory requirements [213].

Moments Accountant [10] provides tighter privacy loss estimation by analyzing higher-order moments of the privacy loss random variable, enabling reduced noise injection while maintaining identical privacy guarantees compared to basic composition theorems.

Adaptive Clipping dynamically adjusts the clipping threshold C based on gradient statistics

throughout training, reducing bias introduced by fixed clipping and improving convergence in high-noise regimes [214].

5.3.2 Related Work

Differential Privacy in Machine Learning

Differential privacy emerged as a gold standard for privacy-preserving data analysis [9, 215], initially focusing on aggregate statistical queries with Laplace or exponential noise mechanisms [216]. The extension to machine learning through DP-SGD [10] represented a significant advancement, enabling privacy-preserving neural network training while maintaining formal privacy guarantees.

Empirical evaluation of DP-SGD has revealed important limitations beyond the privacy-utility trade-off. Bagdasaryan and Shmatikov [217] demonstrated that privacy noise disproportionately affects model performance on minority or under-represented demographic groups, particularly in face recognition tasks. This finding highlights potential fairness implications of DP-SGD deployment, suggesting that privacy protection may inadvertently exacerbate algorithmic bias—a concern our study addresses from the model recipient perspective.

Privacy in Social Media and User-Generated Content

Social media platforms present unique privacy challenges due to the richness and sensitivity of user-generated content, including text posts, images, and behavioral patterns. Previous research has documented significant privacy risks in this domain through various attack vectors. Model inversion attacks [4, 97, 218] can reconstruct sensitive training images from deployed models, while membership inference attacks [1, 2, 93] can determine whether specific individuals’ data was used in training.

These vulnerabilities motivate the application of DP-SGD to social media data, where users maintain expectations of confidentiality despite data aggregation for AI services. However, the high-dimensional nature of social media content (images, free-form text) and the complexity of deep neural

networks present additional implementation challenges compared to traditional DP applications on simple aggregate statistics.

The distinction between local differential privacy (data privatization on user devices) and central differential privacy (server-side privacy guarantees) remains relevant for social media platforms. DP-SGD operates under the central DP model, providing protection even when adversaries have white-box access to model parameters and training procedures [10].

User Perceptions of Privacy-Preserving Technologies

A growing literature examines end-user attitudes toward privacy-preserving machine learning systems. Cummings et al. [219] conducted one of the first comprehensive surveys on user expectations of differential privacy, finding that users demonstrate concern about information leakage that DP is designed to prevent. Critically, their study revealed increased willingness to share personal data when users are informed about DP implementation, suggesting that privacy guarantees can positively influence data contribution decisions.

However, the effectiveness of privacy communication varies significantly based on presentation and framing. Xiong et al. [220] demonstrated that describing DP in terms of real-world implications and attack prevention, rather than technical definitions, substantially improves user comprehension and data sharing willingness. Similarly, Garrido et al. [221] and Kuhtreiber et al. [222] observed that clear, accessible descriptions of data protection mechanisms enhance user understanding and influence sharing decisions.

These findings underscore the importance of effective privacy communication: users are not inherently opposed to data sharing when privacy risks are adequately mitigated and transparently explained. Our work builds on this foundation by investigating data contributor attitudes specifically toward DP-SGD in social media contexts, examining how different implementation approaches (iterative versus final-model-only DP) affect user trust and sharing decisions.

Practitioner Adoption and Implementation Challenges

Despite theoretical advances, real-world DP adoption remains limited due to practical implementation barriers. Rogers et al. [223] note that successful deployment depends not only on theoretical soundness but also on usable tooling and effective communication across stakeholder roles.

Karegar et al. [224] explored interface design patterns to help data analysts reason about DP-protected outputs, while Ngong et al. [225] documented common implementation challenges developers encounter when using popular DP libraries. Their findings identify specific pain points including privacy parameter selection (choosing appropriate ϵ values), accuracy loss diagnosis, and hyperparameter optimization under privacy constraints.

These practitioner-focused studies highlight the gap between theoretical DP guarantees and practical deployment requirements, motivating our investigation of ML engineer perspectives on DP-SGD implementation across different data modalities and application contexts.

Model Deployment and Recipient Perspectives

The impact of DP-SGD on model recipients—organizations and end-users who deploy or consume trained models—remains under-explored in existing literature. While some studies examine user satisfaction with privacy-protected recommendation systems [221], comprehensive analysis of how privacy-utility trade-offs affect deployment decisions is limited.

Model recipients must balance multiple competing objectives: maintaining service quality, ensuring regulatory compliance, and managing computational costs. The fairness implications identified by Bagdasaryan and Shmatikov [217] may further complicate deployment decisions, as differential privacy can lead to uneven user experiences across demographic groups.

Our work addresses this gap by systematically investigating how model recipients evaluate DP-SGD trade-offs, their tolerance for utility degradation in exchange for privacy guarantees, and factors that influence their adoption decisions in practical deployment scenarios.




Fig. 5.2: Instructions of survey questions in Experiment 1, 2 and 3.

5.4 Overview of Experiment Design

We conduct a series of four experiments (with Experiment 2 comprising two modality-specific components, 2A and 2B) that systematically investigate multi-stakeholder perspectives on DP-SGD adoption. Our experimental design follows a sequential structure where insights from each study inform subsequent investigations, enabling iterative refinement of research questions and methodological approaches. Figure 5.2 illustrates the overall research workflow, while Table 5.1 provides a comprehensive summary of experimental parameters.

The sequential design ensures that findings from Experiment 1 (Data Contributors) directly

inform the design of Experiment 2 (ML Engineers), whose results subsequently shape Experiment 3 (Model Recipients). This cascading approach enables us to address stakeholder-specific concerns while maintaining coherence across the multi-stakeholder investigation.

5.4.1 Participant Recruitment and Demographics

Each experiment targets distinct populations representing key stakeholder groups in DP-SGD deployment scenarios:

Experiment 1 recruits active social media users who regularly share text content and images online, serving as proxies for data contributors in privacy-sensitive machine learning applications. Participants will be recruited through social media platforms, online communities, and university participant pools.

Experiments 2A and 2B engage machine learning engineers and data scientists with demonstrated experience in neural network training and deployment. Recruitment will target professional networks (LinkedIn, industry conferences), online ML communities (Reddit r/MachineLearning, Stack Overflow), and graduate programs in computer science and related fields.

Experiment 3 involves end-users of AI-driven social media features who consume model outputs without direct involvement in data contribution or model training. This population may overlap with Experiment 1 participants but focuses specifically on their role as model recipients.

Sample sizes are determined through power analysis balancing statistical requirements with practical constraints. Survey-based components (Experiments 1 and portions of Experiments 2A/2B) target 200-300 respondents to enable robust quantitative analysis of Likert-scale responses and between-group comparisons. Laboratory components (Experiment 2 implementation tasks, Experiment 3 user studies) involve 25-30 participants each, providing sufficient depth for qualitative observation and think-aloud protocol analysis while maintaining feasibility.

5.4.2 Data Modalities and Implementation Paradigms

Our investigation examines two primary data modalities—images and text—to assess whether modality characteristics influence DP-SGD adoption decisions and stakeholder perceptions. This dual-modality approach is motivated by fundamental differences in data representation, privacy sensitivity, and noise impact:

Image Data: High-dimensional, dense representations where DP noise may have different perceptual impacts and where visual privacy concerns (e.g., facial recognition, location inference) are prominent.

Text Data: Sparse, symbolic representations where DP noise may affect semantic meaning differently and where linguistic privacy concerns (e.g., authorship attribution, sentiment inference) predominate.

Experiment 2 explicitly separates these modalities into 2A (image classification) and 2B (text classification/NLP) to enable direct comparison of engineering challenges across data types. Experiment 1 incorporates both photo-sharing and text-posting scenarios to assess whether data contributors' privacy attitudes vary by modality. Experiment 3 evaluates both image-based services (DP-trained recommendation or tagging systems) and text-based services (DP-trained content analysis or sentiment classification) to understand recipient preferences across modalities.

Additionally, we investigate two DP implementation paradigms:

- **Iterative DP:** Privacy accounting across training iterations with intermediate model releases
- **Final-Model DP:** Privacy guarantees applied only to final model parameters

5.4.3 Experimental Protocol and Timeline

All studies adhere to institutional ethics approval requirements with informed consent procedures and explicit withdrawal rights, particularly crucial for data contributor studies given privacy sensitivity. Each experiment undergoes pilot testing to ensure clarity, feasibility, and methodological soundness.

The research timeline follows a structured progression:

Phase 1 (Months 1-2): Instrument development, pilot testing, and ethics approval for Experiment 1.

Phase 2 (Months 3-4): Experiment 1 execution and data analysis, with findings directly informing Experiment 2 design refinements.

Phase 3 (Months 5-7): Experiments 2A and 2B execution, conducted in parallel to maximize efficiency while enabling cross-modality comparison.

Phase 4 (Months 8-9): Experiment 3 design incorporating insights from prior phases, followed by execution and analysis.

Phase 5 (Months 10-12): Cross-experiment synthesis, validation studies, and comprehensive analysis.

Data collection combines quantitative metrics (survey ratings, performance measures, behavioral indicators) with qualitative insights (open-response survey items, think-aloud protocols, semi-structured interviews). This mixed-methods approach enables statistical analysis of stakeholder differences while capturing nuanced perspectives on DP-SGD adoption barriers and facilitators.

5.4.4 Methodological Innovations

Our experimental design incorporates several methodological innovations to enhance ecological validity and practical relevance:

Interactive Educational Tools: We develop web-based educational interfaces (Figure 5.3) that enable participants to explore DP-SGD concepts through guided interaction rather than passive consumption of technical descriptions. These tools facilitate more informed decision-making in subsequent experimental tasks.

Realistic Implementation Scenarios: Experiment 2 tasks utilize realistic datasets and implementation constraints that mirror production deployment challenges, rather than simplified academic


Table 5.1: Summary of Experimental Design

Exp.	Stakeholder Role	Modality	Methodology	Key Measures
1	Data Contributors	Text & Images	Online survey with hypothetical scenarios	Willingness to share (Likert scale), privacy concern ratings, DP comprehension, paradigm preferences
2A	ML Engineers	Images	Laboratory coding task + post-task survey	Implementation success (model accuracy), DP parameter selection, completion time, perceived difficulty, tooling preferences
2B	ML Engineers	Text	Laboratory coding task + post-task survey	As in 2A, with modality-specific challenges and optimization strategies
3	Model Recipients	Text & Images	Interactive prototype evaluation + comparative survey	User satisfaction ratings, perceived quality, trust metrics, privacy perceptions, deployment preferences

examples.

Comparative Evaluation Framework: Experiment 3 employs side-by-side comparison of DP-protected and standard models to isolate the impact of privacy mechanisms on user experience and acceptance.

This comprehensive experimental framework enables systematic investigation of DP-SGD adoption from multiple stakeholder perspectives while maintaining methodological rigor and practical relevance.



(a) Interactive DP-SGD Tutorial Tool

(b) Static DP-SGD Educational Material

Fig. 5.3: Web-based educational tools for DP-SGD concept delivery

5.5 Experiment 1: Data Contributors

5.5.1 Objectives and Rationale

Experiment 1 investigates how social media users perceive the use of their content in DP-SGD model training, addressing the fundamental question of whether differential privacy guarantees translate into increased data sharing willingness. Our investigation focuses on three primary research objectives: (1) quantifying the impact of DP-SGD implementation on user comfort with data sharing, (2) identifying user preferences between different DP paradigms and explanatory frameworks, and (3) determining what privacy assurances and constraints data contributors expect from privacy-preserving machine learning systems.

This experiment extends prior research demonstrating that users' data sharing willingness increases when privacy protections are clearly communicated [219, 220]. By examining social media contexts where data can be highly personal and sensitive, we aim to understand how DP-SGD explanations resonate with data contributors and what factors influence their trust in privacy-preserving systems.

A key innovation is our investigation of user preferences between iterative (white-box) and final-

model (black-box) DP paradigms. We examine whether explaining that "the model updates itself with your data many times but each update is protected" generates different trust levels compared to "your data is used once to produce a final private model." This distinction has practical implications for system design and user communication strategies.

5.5.2 Methodology

We conduct a between-subjects survey experiment with social media users recruited through established online platforms (Amazon Mechanical Turk, Prolific Academic) and targeted social media advertisements. Recruitment aims for demographic diversity across age, gender, technical background, and social media usage patterns. The survey is self-administered with an estimated completion time of 12-15 minutes.

Participants are presented with hypothetical scenarios describing their data usage in machine learning model training. For example: "Your photo posts will be used to improve an image recognition algorithm that curates content. The system employs differential privacy, adding carefully calibrated noise during training to ensure that no individual photo has a detectable impact on the final model." Similar scenarios describe text post usage for recommendation systems.

We employ a randomized experimental vignette design with three manipulated factors:

1. Privacy Mechanism Description:

- (a) **No Privacy (Control):** Standard model training without privacy protections
- (b) **Iterative DP (White-box):** DP-SGD described as protecting each training step through continuous noise injection
- (c) **Final-Model DP (Black-box):** DP described as sanitizing only the final model output

2. Privacy Level Communication: Qualitative framing of privacy-utility trade-offs (e.g., "high privacy protection with slightly reduced accuracy" versus "moderate privacy protection with maintained accuracy") rather than technical parameters, since numerical ϵ values may not be meaningful to non-experts.

3. **Data Modality:** Scenarios encompassing both image and text data to assess modality-specific privacy perceptions.

5.5.3 Measures and Instruments

Our measurement framework captures multiple dimensions of user attitudes and understanding:

Primary Outcome Measures:

- **Data Sharing Willingness:** Five-point Likert scale agreement with statements such as "I would be comfortable with my [posts/photos] being used in this way" (1=Strongly Disagree, 5=Strongly Agree)
- **Perceived Privacy Risk:** Concern ratings for statements like "How worried are you that your personal information could be inferred from the trained model?" (1=Not at all concerned, 5=Extremely concerned)

Secondary Measures:

- **Mechanism Understanding:** Comprehension assessment through multiple-choice questions (e.g., "According to the description, can the model memorize your exact post?" with options Yes/No/Uncertain)
- **Trust in Privacy Protection:** Likert scale rating of agreement with "I trust that this privacy technique will adequately protect my data"
- **Paradigm Preferences:** For participants exposed to multiple descriptions, explicit choice between approaches with free-text justification
- **Demographic and Background Variables:** Age, gender, education level, prior differential privacy familiarity, and existing privacy-protective behaviors

Qualitative Components: Open-ended questions capture nuanced concerns and suggestions, including "What additional information would help you decide whether to share your data?" and "Do you

have any remaining concerns about data usage in this scenario?"

5.5.4 Anticipated Challenges and Mitigation Strategies

Several methodological challenges require careful attention:

Accessible Explanation of DP-SGD: Communicating complex privacy mechanisms without technical jargon poses significant challenges. We address this through iterative pilot testing of multiple explanation formats, focusing on analogies and plain language descriptions. Pre-study pilots will assess comprehension and adjust explanations accordingly.

Social Desirability Bias: Participants may overstate comfort with privacy-protected data sharing. Mitigation strategies include: (1) emphasizing that honest opinions are valuable with no correct answers, (2) including attention and comprehension checks to identify low-quality responses, and (3) using validated scale items from prior privacy research.

Explanation Complexity: The iterative DP paradigm involves more complex concepts that may confuse some participants. We will systematically assess comprehension across explanation types and document confusion patterns to inform communication strategies.

Individual Difference Heterogeneity: Privacy attitudes vary substantially across individuals. Our design accounts for this through demographic stratification, privacy experience measures, and planned subgroup analyses.

5.5.5 Statistical Analysis Plan

We will employ both quantitative and qualitative analytical approaches:

Quantitative Analysis:

- **Primary Analysis:** Between-subjects ANOVA comparing data sharing willingness across privacy mechanism descriptions, with planned contrasts comparing each DP condition to the no-privacy control

- **Secondary Analyses:** T-tests examining modality differences, regression models incorporating demographic and background variables, and interaction analyses examining whether privacy level and mechanism type interact
- **Exploratory Analyses:** Clustering analysis to identify distinct user attitude profiles and decision tree analysis to identify key factors predicting sharing willingness

Qualitative Analysis: Thematic coding of open-ended responses using established qualitative analysis software, with inter-rater reliability assessment for code development and application.

5.5.6 Expected Outcomes and Implications

Based on prior research, we hypothesize that any DP protection will increase data sharing willingness compared to no protection [220]. Additionally, we expect the iterative DP description may generate slightly higher trust due to its emphasis on continuous protection, though its complexity might reduce comprehension for some participants.

We anticipate identifying distinct user segments with varying privacy preferences and risk tolerances. Some users may demand strong privacy guarantees regardless of utility impact, while others may prioritize service quality. Understanding these preference distributions will inform privacy parameter selection in subsequent experiments.

The proportion of users willing to share under different privacy levels will directly influence Experiment 2 parameter selection. If users generally accept $\epsilon \approx 1$ (strong privacy) with moderate utility loss, Experiment 2 will focus on achieving good performance in that regime. Conversely, if users prefer $\epsilon \approx 5$ (moderate privacy) with maintained utility, engineering efforts will target that parameter range.

5.5.7 Integration with Subsequent Experiments

Findings from Experiment 1 will directly inform the design of Experiments 2A and 2B in several ways:

Parameter Selection: User-acceptable privacy levels will guide the (ϵ, δ) ranges that ML engineers attempt to achieve in implementation tasks.

Communication Strategies: Effective explanation approaches identified in Experiment 1 will be incorporated into engineer training materials and recipient evaluation frameworks in Experiment 3.

Paradigm Focus: Strong user preferences for iterative versus final-model DP will influence the implementation approaches emphasized in Experiment 2.

Concern Identification: Qualitative concerns raised by data contributors will inform the guidelines and best practices developed from the multi-stakeholder investigation.

This integration ensures that our multi-stakeholder approach remains coherent while addressing real user needs and preferences throughout the DP-SGD adoption process.

5.6 Experiments 2A & 2B: ML Engineers

Experiments 2A and 2B investigate the practical implementation challenges and decision-making processes that ML engineers encounter when deploying DP-SGD in realistic training scenarios. We separate this investigation into two complementary studies based on data modality: Experiment 2A focuses on image classification tasks representative of social media photo processing, while Experiment 2B examines text analytics tasks reflecting social media post analysis. This modality-specific approach enables systematic comparison of DP-SGD implementation challenges across different data types, capturing potential differences in noise sensitivity, computational requirements, and optimization strategies.

Our primary objectives are to: (1) observe and document practical implementation barriers engineers face when applying DP-SGD, (2) analyze decision-making processes for privacy parameter selection and hyperparameter optimization, (3) assess engineer attitudes toward different DP paradigms from implementation and performance perspectives, and (4) quantify achievable privacy-utility trade-offs across data modalities.

5.6.1 Participant Recruitment and Selection

We recruit ML engineers and data scientists with demonstrated experience in neural network development and deployment. Inclusion criteria require: (1) at least intermediate machine learning knowledge with hands-on deep learning experience, (2) familiarity with standard ML frameworks (PyTorch, TensorFlow), and (3) experience with model training and hyperparameter optimization. We aim for participants with varying levels of differential privacy familiarity, from complete novices to those with theoretical knowledge but limited implementation experience.

Recruitment strategies include professional networking platforms (LinkedIn), emails, university computer science departments, and industry partnerships. Target sample size is 25-30 participants per experiment (2A and 2B), providing sufficient depth for qualitative analysis while enabling quantitative comparison of implementation outcomes.

5.6.2 Experiment 2A: Image Classification Implementation

Task Design and Materials

Participants complete a controlled implementation task in either laboratory or remote screen-sharing environments. The task involves training a convolutional neural network for image classification using a curated dataset simulating social media images (e.g., facial recognition, object classification, or scene understanding).

We provide participants with a comprehensive implementation environment including a pre-configured CNN architecture in PyTorch with standard training pipeline, access to established DP libraries (Opacus, TensorFlow Privacy) with documentation and tutorials, standardized hardware configuration ensuring consistent performance measurement, and automated accuracy assessment and privacy accounting tools.

Participants receive the following directive:

"Your task is to modify this image classification model to satisfy differential privacy

requirements. The trained model must provide formal privacy guarantees ensuring that no individual training image can be identified or reconstructed from the model parameters. Your objective is to achieve the highest possible classification accuracy while meeting the specified privacy constraint of $\epsilon = X$ with $\delta = 10^{-5}$."

The target privacy parameter ϵ will be informed by Experiment 1 findings. If data contributors demand strict privacy guarantees, we will set $\epsilon \leq 2$; if moderate privacy levels are acceptable, we will use $\epsilon \approx 5 - 8$.

Data Collection and Observation Protocol

Sessions last approximately 90-120 minutes with concurrent think-aloud protocols where participants verbalize their decision-making processes. We collect multi-modal data through technical implementation metrics, process documentation, and qualitative insights.

Technical implementation metrics capture DP-SGD integration steps including optimizer replacement, noise multiplier selection, and gradient clipping configuration. We also document hyperparameter optimization trajectories such as learning rate schedules, batch size adjustments, and epoch selection. Performance outcomes include final model accuracy, convergence behavior, and training duration, while privacy accounting measures track achieved (ϵ, δ) values and privacy budget allocation strategies.

Process documentation involves screen recording capturing implementation decisions and debugging processes, audio recording of think-aloud commentary (with participant consent), time-stamped interaction logs with the development environment, and systematic documentation of error patterns and debugging strategies employed. Think-aloud protocols capture decision rationale, confusion points, tool usability feedback, and mental model development regarding DP-SGD mechanics.

5.6.3 Experiment 2B: Text Analytics Implementation

Experiment 2B follows an identical methodological framework but focuses on text analytics tasks representative of social media content analysis. Participants train models for sentiment analysis, topic classification, or next-word prediction using social media post datasets.

Task-Specific Considerations

Text analytics presents distinct challenges for DP-SGD implementation across several dimensions. Pre-processing complexity involves tokenization, embedding strategies, and vocabulary management under privacy constraints. Gradient characteristics differ substantially, as sparse gradient patterns in text models may respond differently to clipping and noise injection compared to dense image representations. Sequential dependencies in recurrent or transformer architectures may present unique privacy accounting challenges, while evaluation requires text-specific performance measures such as F1-score, perplexity, and BLEU scores under privacy constraints.

We systematically compare text and image implementation challenges, documenting modality-specific optimization strategies and performance characteristics. Cross-modal analysis examines relative difficulty perceptions across data modalities, transfer of DP-SGD knowledge and strategies between tasks, modality-specific performance degradation patterns, and preferred implementation approaches for different data types.

5.6.4 Expected Findings and Hypotheses

Based on existing literature and preliminary observations, we hypothesize several key patterns. Regarding performance, we expect model accuracy will decrease monotonically with stricter privacy constraints (smaller ϵ), image classification tasks will exhibit greater sensitivity to DP noise than text analytics due to dense, high-dimensional representations, and text models may require different optimization strategies due to sparse gradient characteristics.

For implementation patterns, we anticipate common implementation challenges will include gradient clipping threshold selection, noise multiplier tuning, and privacy budget allocation. Engineers with prior DP exposure will likely achieve better privacy-utility trade-offs through more effective hyperparameter selection, while tool usability and documentation quality will significantly influence implementation success and satisfaction.

Regarding paradigm preferences, engineers may express preference for simpler privacy paradigms (final-model perturbation) despite library support for DP-SGD, particularly if implementation complexity is high. Additionally, modality-specific preferences may emerge based on relative implementation difficulty and performance outcomes.

5.6.5 Integration with Multi-Stakeholder Framework

Findings from Experiments 2A and 2B directly inform Experiment 3 design and the broader multi-stakeholder analysis through several integration mechanisms. Performance benchmarking provides realistic performance expectations for Experiment 3 evaluation—if engineers consistently achieve 80% accuracy at $\epsilon = 1$ for image tasks, Experiment 3 can examine whether model recipients accept this performance level compared to 90

Implementation insights inform communication strategies for both data contributors (Experiment 1 follow-up) and model recipients (Experiment 3). For example, if engineers consistently express difficulty explaining privacy parameters to stakeholders, Experiment 3 can test different explanation approaches. Paradigm evaluation provides technical grounding for user preference investigations, with discrepancies between engineer and user preferences (e.g., engineers preferring final-model DP while users prefer iterative DP) highlighting important design trade-offs.

Technical challenges and successful strategies identified in implementation studies directly contribute to practical DP-SGD deployment guidelines, ensuring recommendations are grounded in real implementation experience rather than theoretical considerations alone. This systematic integration ensures

that our multi-stakeholder investigation captures both technical feasibility constraints and stakeholder preference requirements, enabling realistic and actionable recommendations for DP-SGD adoption.

5.7 Experiment 3: Model Recipients

5.7.1 Objectives and Rationale

Experiment 3 investigates the perspectives of model recipients—end-users who consume the outputs of machine learning models trained on social media data. This stakeholder group represents the final link in the DP-SGD deployment chain, where theoretical privacy guarantees and implementation efforts ultimately translate into user experience outcomes. Our investigation addresses critical questions about the practical viability of privacy-preserving machine learning systems from the recipient perspective.

Specifically, we examine whether users perceive quality differences in AI-driven services when models employ DP-SGD training, how privacy awareness influences user satisfaction and system preference, and whether explicit communication about privacy protections can offset potential utility degradation in user perception. This investigation is essential because successful DP-SGD deployment requires not only data contributor acceptance and engineering feasibility, but also recipient tolerance for any resulting changes in service quality.

The experiment addresses three primary research objectives: (1) quantifying user-perceptible impacts of DP-SGD on model output quality and relevance, (2) assessing how privacy awareness influences user satisfaction and system preference, and (3) determining user willingness to accept utility trade-offs in exchange for enhanced privacy protections.

5.7.2 Experimental Design and Methodology

We employ a user experience evaluation focused on assessing recipient attitudes toward DP-SGD trained models and their tolerance for privacy-utility trade-offs. The experimental context utilizes realistic model outputs and performance scenarios derived from Experiments 2A and 2B to understand how recipients

evaluate privacy-preserving machine learning systems in practical deployment contexts.

System Implementation

We develop evaluation materials that present realistic model performance scenarios rather than live system comparisons. Using performance data and model outputs from Experiments 2A and 2B, we create representative examples of DP-SGD system behavior including accuracy levels, recommendation quality, and content relevance that participants can evaluate. These materials reflect actual privacy-utility trade-offs observed in the engineering experiments, ensuring ecological validity while focusing on recipient utility assessment rather than implementation details.

Experimental Protocol

Participants evaluate model performance scenarios and outputs without requiring live system interaction or technical implementation knowledge. We present participants with realistic performance summaries (e.g., "System A achieves 85

The experimental protocol focuses on utility assessment and privacy trade-off preferences rather than live system interaction. Participants review performance scenarios and sample outputs that demonstrate realistic privacy-utility trade-offs observed in the engineering studies. The evaluation involves structured assessment of different performance levels corresponding to various privacy parameters, allowing participants to express preferences for different privacy-utility combinations without requiring technical understanding of DP-SGD implementation.

5.7.3 Measurement Framework

Our assessment framework captures multiple dimensions of user experience and privacy perception through validated instruments and novel measures developed specifically for DP-SGD evaluation contexts.

Primary outcome measures include acceptability ratings of different privacy-utility trade-off scenarios, assessed through Likert scale evaluations of various performance levels (e.g., "How acceptable

would 78% accuracy be if it provided strong privacy protection?"). We measure deployment preferences through choice scenarios presenting different combinations of accuracy and privacy protection. Utility tolerance measures examine the maximum performance degradation participants would accept in exchange for privacy guarantees.

Secondary measures examine privacy valuation and organizational decision-making factors relevant to model deployment. Trust and privacy perception questions assess the value participants place on formal privacy guarantees versus service performance. Deployment scenario evaluation examines how different organizational contexts (e.g., healthcare recommendations vs. entertainment content) affect privacy-utility preferences. Willingness to pay measures explore whether participants would accept subscription costs or reduced functionality in exchange for privacy protection.

Fairness and diversity perception measures examine whether users detect differences in content representation, demographic bias, or topic coverage between system variants. Open-ended questions capture unanticipated observations and nuanced preferences through prompts such as "Describe any differences you noticed between the two systems" and "How important is privacy protection in recommendation systems to you?"

5.7.4 Anticipated Challenges and Mitigation Strategies

Several methodological challenges require careful attention to ensure valid and interpretable results.

Effect size calibration presents a significant challenge—privacy-utility trade-offs must be perceptible to participants while remaining realistic. If differences are too subtle, participants may not form meaningful preferences; if too pronounced, results may not reflect realistic deployment scenarios. We address this through iterative pilot testing with privacy parameter adjustment informed by Experiment 2 findings to achieve moderate but detectable performance gaps.

Explanation bias represents another concern, as privacy protection framing might inherently bias participants toward the DP-SGD system regardless of actual performance. Our mitigation strategy

involves neutral explanation language, counterbalanced disclosure timing, and explicit assessment of participant understanding to ensure informed but unbiased evaluation.

Prototype limitations may affect ecological validity since our experimental system necessarily simplifies real-world recommendation complexity. We address this through realistic content diversity, authentic interaction patterns, and focus on relative rather than absolute performance assessment.

Individual difference variation in privacy preferences and recommendation system experience may introduce significant response heterogeneity. Our analysis plan accounts for this through demographic stratification, privacy attitude measurement, and planned subgroup analyses examining different user segments.

5.7.5 Statistical Analysis Plan

Our analytical approach combines quantitative comparison of system performance with qualitative examination of user reasoning and preferences.

Primary quantitative analyses employ paired t-tests comparing mean quality ratings between DP-SGD and non-private model conditions, with effect size estimation to assess practical significance. Secondary analyses examine privacy information impact through repeated measures ANOVA with disclosure timing as a between-subjects factor. Behavioral measures undergo similar paired comparison analysis with appropriate transformations for non-normal distributions.

Preference analysis utilizes chi-square tests for categorical preference choices with logistic regression to identify predictors of DP-SGD system preference. Trade-off analysis examines relationships between quality perception, privacy valuation, and ultimate system preference through correlation and regression analysis.

Qualitative analysis employs systematic thematic coding of open-ended responses to identify recurring themes in user reasoning, unexpected concerns, and suggestions for improvement. Inter-rater reliability assessment ensures coding consistency and validity.

Subgroup analyses examine whether privacy attitudes, technical background, or demographic characteristics moderate responses to DP-SGD systems, providing insights into user segments most and least receptive to privacy-preserving approaches.

5.7.6 Expected Outcomes and Integration

We anticipate several key findings that will complete our multi-stakeholder analysis of DP-SGD adoption. Regarding performance perception, we expect users will detect quality differences between systems when privacy protection significantly impacts recommendation accuracy, but that explicit privacy information may partially offset negative quality perceptions. For privacy trade-off preferences, we anticipate substantial individual variation, with some users strongly preferring privacy protection despite quality costs while others prioritize performance regardless of privacy implications.

Integration with prior experiments will reveal critical adoption barriers and facilitators. If data contributors demand strong privacy guarantees (Experiment 1) but recipients cannot tolerate the resulting quality degradation (Experiment 3), while engineers struggle with implementation complexity (Experiments 2A/2B), this tension will inform our deployment guidelines and highlight areas requiring technical advancement or communication improvement.

Conversely, if recipients appreciate privacy protections and tolerate moderate quality reductions, while data contributors show increased sharing willingness and engineers successfully implement DP-SGD, this convergence suggests favorable conditions for practical deployment. Such findings would support development of comprehensive adoption strategies addressing all stakeholder needs.

The complete multi-stakeholder perspective enables realistic assessment of DP-SGD viability and identification of specific interventions needed to facilitate successful privacy-preserving machine learning deployment in social media contexts.

5.8 Discussion

In this section, we will synthesize findings from all experiments (once completed) and discuss their implications for theory and practice. Although we have not yet run the studies, we can outline the expected integration of results and the kinds of interpretation we will perform. The general discussion will likely cover the following themes:

1. **Contrasting Stakeholder Perspectives:** We expect to compare what data contributors want, what ML engineers can do, and what model recipients prefer. For example, a possible scenario is:
 - Data contributors (Exp1) strongly favor rigorous privacy (perhaps indicating willingness to share data only if ϵ is very low or if certain assurances are given).
 - ML engineers (Exp2) find that such rigorous privacy dramatically lowers model accuracy or is hard to implement, and they might prefer a middle ground or need better tools.
 - Model recipients (Exp3) may notice the lowered accuracy and not entirely welcome it, unless they highly value the ethical aspect.

This could create a privacy–utility tension among stakeholders. Our discussion will explore how to reconcile these differences. Perhaps guidelines might suggest transparency with end-users: if users understand the reason for a slight drop in personalization (i.e., privacy), they might accept it more readily, bridging the contributor and recipient positions. We will also consider if there are misconceptions to clear up: maybe contributors in Exp1 overestimated what DP guarantees (some might think it protects against all evils, whereas engineers know it doesn’t stop all forms of leakage or misuse outside the model). These misconceptions, if found, will be discussed in light of how to better communicate DP’s true guarantees and limits to non-experts.

2. **White-Box vs Black-Box Paradigms:** A specific point of discussion will be the comparison of the two DP paradigms across studies. We will piece together:

- **Contributors' trust:** Did they trust one paradigm more? Suppose Exp1 revealed that when told “iterative DP,” participants felt slightly more secure than “final-only DP.” We would discuss why that might be—perhaps iterative DP sounds like continuous oversight, whereas final-only might sound like they only care at the end. On the other hand, it’s possible users didn’t meaningfully distinguish them, focusing only on whether privacy was mentioned at all.
- **Engineers' preference:** If our engineers in Exp2 indicated a clear technical preference (maybe they found white-box DP-SGD straightforward with libraries, or maybe some said they wished they could just noise the output to save accuracy), we will discuss the practicality and risks of those approaches. The white-box DP (DP-SGD) is known to provide strong protection even if adversaries have full knowledge, so we might argue that while engineers might desire simpler black-box methods, those could give weaker guarantees or be harder to audit. We will consider whether any promising hybrid approaches emerged from participants’ suggestions (some might mention federated learning or other frameworks).
- **Recipients:** Though model recipients might not directly know the difference, if we had any follow-ups about it, we’ll note that they mostly care about outcomes, not the internal process. This implies that the distinction is more critical for contributors and engineers. Therefore, guidelines could state that from a user perspective, what matters is the net effect on privacy and quality, not how it’s achieved—giving flexibility to engineers as long as privacy claims hold.

3. **Practical Guidelines for Adoption of DP-SGD:** The culmination of the discussion will be a set of recommended practices informed by the entire study. Some expected guidelines might include:

- *Communication Guidelines:* How to effectively communicate privacy measures to users. For example, if Exp1 and Exp3 show that a straightforward, implication-focused description worked best, we will advocate for adopting such language in user-facing privacy policies

(e.g., “We use a technique that guarantees no single one of your posts can be identified from our AI models” rather than “we use DP with $\epsilon = 1$ ”).

- *Privacy Budget Choices:* Recommendations on setting ϵ and other DP parameters in practice, balancing user expectations and model utility. For instance, suppose our findings indicate that users are satisfied with moderate privacy if utility stays high (or perhaps they demand strong privacy but then are unhappy with the results as recipients). We might suggest an ϵ range that appears acceptable. We will also stress the importance of making this choice transparent and perhaps even giving users some control if feasible.
- *Engineer Tooling and Training:* Based on Exp2, we expect to highlight where engineers need support. If many struggled with the math or interpretation of privacy loss over training, we might recommend better tooling (e.g., libraries that automatically adjust noise to meet a target accuracy, or visualization tools to show the privacy-utility frontier).
- *Iterative Design and Pilot Deployments:* We might recommend that organizations pilot DP-SGD on less critical features first, gather feedback (like our study did in microcosm), then scale up. This mirrors our multi-phase study: incorporate user feedback early (Exp1) and engineer feedback (Exp2) before exposing a broad user base (Exp3).
- *Policy and Ethical Considerations:* We will discuss how our findings relate to broader ethical AI principles. For example, if DP-SGD was found to exacerbate some biases and users noticed unfairness, we’d recommend pairing privacy efforts with bias mitigation techniques to ensure one ethical goal doesn’t undermine another. We’ll emphasize that differential privacy should not be considered in isolation: it intersects with transparency, fairness, and user control.
- *When Not to Use DP-SGD:* Perhaps a counterintuitive guideline but important: if our results show that in some scenarios DP-SGD’s cost is too high and users don’t value it enough (maybe for certain low-sensitivity data, the platform might decide not to use DP if it’s not

needed), focus DP where it matters most.

4. **Limitations and Future Work:** We will candidly address limitations of our study to temper our guidelines. For example, our participant samples (maybe mostly US-based or tech-aware users) might not represent all populations—privacy perceptions can vary by culture. Also, our experiments were simulations; real-world behavior might differ (users might say they care about privacy in a survey but act differently in practice). We will suggest that future work do field trials or test our guidelines in a live system. Additionally, any surprising findings (if, say, engineers were far more capable with DP than assumed, or users overwhelmingly didn't care about privacy) will be discussed with possible explanations and further research questions. We will connect back to related work: do our findings support or contradict prior studies? For instance, if we find that explaining DP boosted sharing willingness (Exp1), that reinforces Cummings et al.'s [219] results; if we find recipients didn't care about privacy, that might add a new insight that just because contributors want privacy doesn't mean recipients will reward it.

5. **The Path Forward for DP-SGD in Social Media:** We will conclude on an optimistic but realistic note: differential privacy offers a way to protect users at scale, but it must be implemented with attention to human factors. Our multi-stakeholder study will have illuminated a path where platforms can align the interests of users, engineers, and recipients. We might propose a conceptual model or framework that shows how to balance privacy budget, utility, and stakeholder satisfaction. For example, a diagram or conceptual table might map each stakeholder's needs to certain DP-SGD configurations or processes. In summary, the general discussion will merge all the findings into a coherent narrative and provide actionable recommendations. It will emphasize the novelty of our approach (covering all three roles) and argue that you cannot successfully deploy privacy-preserving ML without considering the ecosystem of people around it. This aligns with recent calls in the literature for socio-technical perspectives on privacy technology.

Chapter 6

Conclusions

This dissertation has advanced the field of differentially private machine learning through comprehensive theoretical, algorithmic, and practical contributions. The research has addressed fundamental challenges in privacy-preserving machine learning across multiple domains, from infinite data streams to federated learning, and from theoretical frameworks to real-world deployment considerations. We introduced DPED (Differentially Private Embedding Distillation), the first framework for training differentially private text embeddings through teacher-student distillation with multi-layer noise injection. This work established novel theoretical foundations for privacy-preserving representation learning, proving convergence guarantees and providing tight privacy bounds that demonstrate how data partitioning among teachers leads to stronger privacy guarantees compared to single-model approaches. Through DPI (Differential Privacy for Infinite Data Streams), we solved the long-standing problem of unbounded privacy leakage in continuous data processing. Our framework provides the first user-level differential privacy solution for infinite streaming applications, introducing sensitivity compression techniques and novel boosting mechanisms that maintain bounded privacy loss while maximizing accuracy over time. UDP-FL (Universally Harmonizing Differential Privacy Mechanisms for Federated Learning) demonstrated how to harmonize diverse differential privacy mechanisms beyond the commonly used Gaussian approach. By leveraging Rényi Differential Privacy as a mediator, we showed how optimal mechanisms like the Staircase mechanism can significantly improve the privacy-utility tradeoff in federated learning scenarios, achieving faster convergence and higher accuracy compared to state-of-the-art methods. Across all contributions, we pro-




Fig. 6.1: Research roadmap: building secure and privacy-centric data solutions.

vided extensive experimental validation on diverse datasets and applications, demonstrating substantial improvements over existing methods. Our frameworks have shown resilience against advanced privacy attacks including membership inference, data reconstruction, and attribute inference attacks, validating their practical effectiveness in protecting sensitive information. While the core of this dissertation focuses on advancing differential privacy, my broader research has also addressed other critical areas of data science, security, and machine learning. These works include enhancing the security of distributed ledger systems through the analysis of block withholding attacks in blockchain [226], developing novel workflows for large language models in multi-role collaborative reranking [227], creating efficient methods for generative active learning [228], and improving body composition estimation from 3D body scans [229].

Building upon these contributions and the insights from our multi-stakeholder analysis framework outlined in the thesis proposal, we hope to have a complete research circle that motivate us to prove more privacy-preserving solutions.

6.0.1 Human-Centered Privacy-Preserving AI

The most pressing need is to bridge the gap between technical privacy guarantees and human understanding and acceptance. Our proposed multi-stakeholder study represents the first step toward this goal, but extensive work remains:

- **User-Centric Privacy Design.** Future research must develop privacy mechanisms that align

with user mental models and expectations. This includes creating adaptive privacy systems that can adjust protection levels based on user preferences, data sensitivity, and contextual factors. Research should focus on developing interpretable privacy metrics that users can understand and trust.

- **Stakeholder-Aware Privacy Frameworks.** Building on our multi-stakeholder approach, future work should develop comprehensive frameworks that simultaneously consider data contributors, ML engineers, and model recipients. This includes creating tools for transparent privacy communication, developing privacy-preserving systems that maintain usability for practitioners, and ensuring that end-users receive both privacy protection and satisfactory service quality.
- **Empirical Privacy Studies.** Large-scale empirical studies are needed to understand how privacy-preserving technologies perform in real deployment scenarios across different populations, cultures, and application domains. This research should inform the development of privacy policies and technical standards that reflect actual user needs and concerns.

6.0.2 Scalable Privacy-Preserving AI for Complex Modern Applications

As AI systems become more sophisticated and data-intensive, privacy-preserving techniques must evolve to handle emerging challenges:

- **Large Language Models and Foundation Models.** Current privacy-preserving techniques face significant challenges when applied to large language models and foundation models. Future research should develop specialized privacy mechanisms that can handle the scale and complexity of these models while maintaining their emergent capabilities. This includes investigating privacy-preserving fine-tuning, prompt engineering under privacy constraints, and federated learning for foundation models.
- **Multi-Modal and Cross-Domain Privacy.** Modern AI applications increasingly involve multi-modal data (text, images, audio, video) and cross-domain learning. Privacy-preserving techniques

must evolve to handle the unique challenges posed by correlations across modalities and the complex privacy implications of cross-domain information leakage.

- **Real-Time and Edge Computing Scenarios.** The proliferation of edge computing and real-time AI applications requires privacy-preserving techniques that can operate under strict latency and computational constraints. Future research should focus on lightweight privacy mechanisms, hardware-accelerated privacy-preserving computation, and privacy-preserving edge-cloud collaboration.

6.0.3 Advanced Threat Models and Next-Generation Privacy Guarantees

As adversarial capabilities evolve, privacy-preserving AI must anticipate and defend against increasingly sophisticated attacks:

- **Beyond Differential Privacy.** While differential privacy provides strong theoretical guarantees, future research should explore complementary and alternative privacy notions that address its limitations. This includes developing privacy frameworks that better handle correlated data, temporal dependencies, and complex auxiliary information scenarios.
- **Adversarial Privacy Research.** Future work should systematically study the robustness of privacy-preserving AI systems against adaptive adversaries who can manipulate training data, exploit model updates over time, or leverage external information sources. This includes developing privacy-preserving techniques that remain robust under realistic threat models.
- **Privacy-Preserving AI Governance.** As privacy-preserving AI becomes more prevalent, research is needed on governance frameworks that can verify, audit, and regulate these systems. This includes developing standardized privacy testing methodologies, creating compliance frameworks for privacy-preserving AI deployment, and establishing best practices for privacy-preserving AI system lifecycle management.

Bibliography

- [1] Reza Shokri, Marco Stronati, Congzheng Song, and Vitaly Shmatikov. Membership inference attacks against machine learning models. In *IEEE S&P*, 2017.
- [2] Nicholas Carlini, Steve Chien, Milad Nasr, Shuang Song, Andreas Terzis, and Florian Tramer. Membership inference attacks from first principles. In *S&P*, pages 1897–1914. IEEE, 2022.
- [3] Yuxin Wen, Arpit Bansal, Hamid Kazemi, Eitan Borgnia, Micah Goldblum, Jonas Geiping, and Tom Goldstein. Canary in a coalmine: Better membership inference with ensembled adversarial queries. *arXiv preprint arXiv:2210.10750*, 2022.
- [4] Jonas Geiping, Hartmut Bauermeister, Hannah Dröge, and Michael Moeller. Inverting gradients-how easy is it to break privacy in federated learning? *NeurIPS*, 33:16937–16947, 2020.
- [5] Yugeng Liu, Rui Wen, Xinlei He, Ahmed Salem, Zhikun Zhang, Michael Backes, Emiliano De Cristofaro, Mario Fritz, and Yang Zhang. ML-Doctor: Holistic risk assessment of inference attacks against machine learning models. In *USENIX Security 22*. USENIX Association, August 2022.
- [6] Helen Nissenbaum. Privacy in context: Technology, policy, and the integrity of social life. In *Privacy in context*. Stanford University Press, 2009.
- [7] Alessandro Acquisti, Laura Brandimarte, and George Loewenstein. Privacy and human behavior in the age of information. *Science*, 347(6221):509–514, 2015.
- [8] Cynthia Dwork, Frank McSherry, Kobbi Nissim, and Adam D. Smith. Calibrating noise to sensitivity in private data analysis. In *Theory of Cryptography Conference*, 2006.
- [9] Cynthia Dwork, Frank McSherry, Kobbi Nissim, and Adam Smith. Calibrating noise to sensitivity in private data analysis. In *TCC 2006*. Springer, 2006.
- [10] Martin Abadi, Andy Chu, Ian Goodfellow, H Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. Deep learning with differential privacy. In *CCS*, 2016.
- [11] Ilya Mironov. Rényi differential privacy. In *2017 IEEE 30th computer security foundations symposium (CSF)*. IEEE, 2017.
- [12] Úlfar Erlingsson, Vasyl Pihur, and Aleksandra Korolova. Rappor: Randomized aggregatable privacy-preserving ordinal response. In *Proceedings of the 21st ACM Conference on Computer and Communications Security, CCS '14*. ACM, 2014. Introduces RAPPOR, deployed in Chrome for privacy-preserving telemetry via randomized response.

- [13] Apple and Google. Exposure notification privacy-preserving analytics (enpa). White paper, Apple Inc., April 2021. Describes ENPA, used alongside Exposure Notification System for differential-privacy based metrics for PHAs.
- [14] David Smith, OpenDP, and Microsoft Azure Team. Whitenoise: An open source library to protect data with differential privacy. online article / toolkit, May 2020. Introduces Microsoft's WhiteNoise differential-privacy toolkit developed in collaboration with Harvard OpenDP.
- [15] Shuya Feng, Meisam Mohammady, Han Wang, Xiaochen Li, Zhan Qin, and Yuan Hong. Dpi: Ensuring strict differential privacy for infinite data streaming. In *IEEE S&P*, 2024.
- [16] Shuya Feng, Meisam Mohammady, Hanbin Hong, Shenao Yan, Ashish Kundu, Binghui Wang, and Yuan Hong. Universally harmonizing differential privacy mechanisms for federated learning: Boosting accuracy and convergence. *arXiv preprint arXiv:2407.14710*, 2024.
- [17] Shuya Feng and Yuan Hong. Dped: Multi-layer noise distillation for privacy-preserving text embeddings. *Under submission*, 2025.
- [18] Shuya etc.. Feng. Dped: Multi-layer noise distillation for privacy-preserving text embeddings. *On going work*, 2025.
- [19] Pablo Ameigeiras, Juan J Ramos-Munoz, Jorge Navarro-Ortiz, and Juan M Lopez-Soler. Analysis and modelling of youtube traffic. *Transactions on Emerging Telecommunications Technologies*, 23(4):360–377, 2012.
- [20] Marina Thottan, Guanglei Liu, and Chuanyi Ji. Anomaly detection approaches for communication networks. In *Algorithms for next generation networks*, pages 239–261. Springer, 2010.
- [21] P Ranjitha. Streaming analytics over real-time big data. *Global Journal of Computer Science and Technology*, 2015.
- [22] Cynthia Dwork, Moni Naor, Toniann Pitassi, and Guy N Rothblum. Differential privacy under continual observation. In *Proceedings of the forty-second ACM Symposium on Theory of Computing*, pages 715–724, 2010.
- [23] Jean Bolot, Nadia Fawaz, Shanmugavelayutham Muthukrishnan, Aleksandar Nikolov, and Nina Taft. Private decayed predicate sums on streams. In *Proceedings of the 16th International Conference on Database Theory*, pages 284–295, 2013.
- [24] T.-H. Hubert Chan, Elaine Shi, and Dawn Song. Private and continual release of statistics. *ACM Trans. Inf. Syst. Secur.*, 14(3), nov 2011.
- [25] Yan Chen, Ashwin Machanavajjhala, Michael Hay, and Gerome Miklau. Pegasus: Data-adaptive differentially private stream processing. In *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*, page 1375–1388, 2017.
- [26] Cynthia Dwork. Differential privacy in new settings. In *Proceedings of the twenty-first annual ACM-SIAM symposium on Discrete Algorithms*, pages 174–183. SIAM, 2010.
- [27] Victor Perrier, Hassan Jameel Asghar, and Dali Kaafar. Private continual release of real-valued data streams. *arXiv preprint arXiv:1811.03197*, 2018.

- [28] Tianhao Wang, Joann Qiongna Chen, Zhikun Zhang, Dong Su, Yueqiang Cheng, Zhou Li, Ninghui Li, and Somesh Jha. Continuous release of data streams under both centralized and local differential privacy. In *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security*, page 1237–1253, 2021.
- [29] Georgios Kellaris, Stavros Papadopoulos, Xiaokui Xiao, and Dimitris Papadias. Differentially private event sequences over infinite streams. *Proc. VLDB Endow.*, 7(12):1155–1166, aug 2014.
- [30] Qian Wang, Yan Zhang, Xiao Lu, Zhibo Wang, Zhan Qin, and Kui Ren. Rescuedp: Real-time spatio-temporal crowd-sourced data publishing with differential privacy. In *The 35th Annual IEEE International Conference on Computer Communications*, pages 1–9. IEEE, 2016.
- [31] Liyue Fan and Li Xiong. An adaptive approach to real-time aggregate monitoring with differential privacy. *IEEE Transactions on knowledge and data engineering*, 26(9):2094–2106, 2013.
- [32] Liyue Fan, Li Xiong, and Vaidy Sunderam. Differentially private multi-dimensional time series release for traffic monitoring. In *IFIP Annual Conference on Data and Applications Security and Privacy*, pages 33–48. Springer, 2013.
- [33] Han Wang, Shangyu Xie, and Yuan Hong. Videodp: A flexible platform for video analytics with differential privacy. *Proc. Priv. Enhancing Technol.*, 2020.
- [34] Han Wang, Yuan Hong, Yu Kong, and Jaideep Vaidya. Publishing video data with indistinguishable objects. In *Proceedings of the 23rd International Conference on Extending Database Technology*, pages 323–334, 2020.
- [35] Bingyu Liu, Shangyu Xie, Han Wang, Yuan Hong, Xuegang Ban, and Meisam Mohammady. VTDP: privately sanitizing fine-grained vehicle trajectory data with boosted utility. *IEEE Trans. Dependable Secur. Comput.*, 18(6):2643–2657, 2021.
- [36] Cynthia Dwork. Differential privacy: A survey of results. In *Theory and Applications of Models of Computation*, volume 4978, pages 1–19, 2008.
- [37] Yuan Hong, Jaideep Vaidya, Haibing Lu, Panagiotis Karras, and Sanjay Goel. Collaborative search log sanitization: Toward differential privacy and boosted utility. *TDSC*, 2015.
- [38] Úlfar Erlingsson, Vasyl Pihur, and Aleksandra Korolova. RAPPOR: randomized aggregatable privacy-preserving ordinal response. In *Proceedings of the 2014 ACM SIGSAC Conference on Computer and Communications Security*, pages 1054–1067, 2014.
- [39] Cynthia Dwork and Aaron Roth. The algorithmic foundations of differential privacy. *Foundations and Trends® in Theoretical Computer Science*, 9(3–4):211–407, 2014.
- [40] Cynthia Dwork, Guy N Rothblum, and Salil Vadhan. Boosting and differential privacy. In *FOCS*, pages 51–60, 2010.
- [41] Moritz Hardt and Guy N. Rothblum. A multiplicative weights mechanism for privacy-preserving data analysis. In *2010 IEEE 51st Annual Symposium on Foundations of Computer Science*, pages 61–70, 2010.
- [42] Cynthia Dwork, Moni Naor, Omer Reingold, Guy N Rothblum, and Salil Vadhan. On the complexity of differentially private data release: efficient algorithms and hardness results. In *Proceedings of the 41st annual ACM Symposium on Theory of Computing*, pages 381–390, 2009.

- [43] Yijie Wang, Xiaoyong Li, Xiaoling Li, and Yuan Wang. A survey of queries over uncertain data. *Knowledge and information systems*, 37(3):485–530, 2013.
- [44] Mohamed A Soliman, Ihab F Ilyas, and Kevin Chen-Chuan Chang. Probabilistic top-k and ranking-aggregate queries. *ACM Transactions on Database Systems (TODS)*, 33(3):1–54, 2008.
- [45] Ihab F Ilyas, Walid G Aref, and Ahmed K Elmagarmid. Supporting top-k join queries in relational databases. *The VLDB journal*, 13:207–221, 2004.
- [46] Xin Li, Young Jin Kim, Ramesh Govindan, and Wei Hong. Multi-dimensional range queries in sensor networks. In *Proceedings of the 1st international conference on Embedded networked sensor systems*, pages 63–75, 2003.
- [47] DUTTA GAURAV. Real-time covid 19 data, 2023.
- [48] Iman Sharafaldin, Arash Habibi Lashkari, Saqib Hakak, and Ali A. Ghorbani. Developing realistic distributed denial of service (ddos) attack dataset and taxonomy. In *2019 International Carnahan Conference on Security Technology (ICCST)*, pages 1–8, 2019.
- [49] Foreign Agricultural Service, Department of Agriculture. Production, supply, and distribution database. <http://apps.fas.usda.gov/psdonline>, 2020.
- [50] Angelo Coluccia, Alessandro D’Alconzo, and Fabio Ricciato. Distribution-based anomaly detection via generalized likelihood ratio test: A general maximum entropy approach. *Computer Networks*, 57(17):3446–3462, 2013.
- [51] Jinwon An and Sungzoon Cho. Variational autoencoder based anomaly detection using reconstruction probability. *Special lecture on IE*, 2(1):1–18, 2015.
- [52] Bo Zong, Qi Song, Martin Renqiang Min, Wei Cheng, Cristian Lumezanu, Daeki Cho, and Haifeng Chen. Deep autoencoding gaussian mixture model for unsupervised anomaly detection. In *International conference on learning representations*, 2018.
- [53] Thomas Defard, Aleksandr Setkov, Angelique Loesch, and Romaric Audigier. Padim: a patch distribution modeling framework for anomaly detection and localization. In *ICPR International Workshops and Challenges*, pages 475–489. Springer, 2021.
- [54] Hafiz Salman Asif, Periklis A. Papakonstantinou, and Jaideep Vaidya. How to accurately and privately identify anomalies. In *Proceedings of the ACM SIGSAC Conference on Computer and Communications Security*, pages 719–736. ACM, 2019.
- [55] Jairo Alonso Giraldo, Alvaro A. Cárdenas, Murat Kantarcioglu, and Jonathan Katz. Adversarial classification under differential privacy. In *NDSS*, 2020.
- [56] Meisam Mohammady, Han Wang, Lingyu Wang, Mengyuan Zhang, Yosr Jarraya, Suryadipta Majumdar, Makan Pourzandi, Mourad Debbabi, and Yuan Hong. Dpoad: Differentially private outsourcing of anomaly detection through iterative sensitivity learning. *arXiv preprint arXiv:2206.13046*, 2022.
- [57] Min Du, R. Jia, and Dawn Xiaodong Song. Robust anomaly detection and backdoor attack detection via differential privacy. *ArXiv*, abs/1911.07116, 2020.
- [58] Markus Goldstein and Andreas Dengel. Histogram-based outlier score (hbos): A fast unsupervised anomaly detection algorithm. *KI-2012: poster and demo track*, 1:59–63, 2012.

- [59] Badrul Sarwar, George Karypis, Joseph Konstan, and John T Riedl. Application of dimensionality reduction in recommender system-a case study. 2000.
- [60] Xun Zhou, Jing He, Guangyan Huang, and Yanchun Zhang. Svd-based incremental approaches for recommender systems. *Journal of Computer and System Sciences*, 81(4):717–733, 2015.
- [61] Fei Tony Liu, Kai Ming Ting, and Zhi-Hua Zhou. Isolation-based anomaly detection. *ACM Transactions on Knowledge Discovery from Data (TKDD)*, 6(1):3, 2012.
- [62] Yaling Zhang, Na Liu, and Shangping Wang. A differential privacy protecting k-means clustering algorithm based on contour coefficients. *PloS one*, 13(11):e0206832, 2018.
- [63] Chunyong Yin, Lingfeng Shi, Ruxia Sun, and Jin Wang. Improved collaborative filtering recommendation algorithm based on differential privacy protection. *The Journal of Supercomputing*, 76:5161–5174, 2020.
- [64] Jessica Lin, Eamonn Keogh, Stefano Lonardi, and Bill Chiu. A symbolic representation of time series, with implications for streaming algorithms. In *Proceedings of the 8th ACM SIGMOD Workshop*, 2003.
- [65] Sirish Chandrasekaran and Michael J Franklin. Streaming queries over streaming data. In *VLDB'02: Proceedings of the 28th International Conference on Very Large Databases*, pages 203–214. Elsevier, 2002.
- [66] Xiaochen Li, Zhan Qin, Kui Ren, Chen Gong, Shuya Feng, Yuan Hong, and Tianhao Wang. Delay-allowed differentially private data stream release. In *32nd Annual Network and Distributed System Security Symposium, NDSS 2025, San Diego, California, USA, February 24-28, 2025*. The Internet Society, 2025.
- [67] Hamid Ebadi, David Sands, and Gerardo Schneider. Differential privacy: Now it's getting personal. *POPL '15*, page 69–81, New York, NY, USA, 2015. Association for Computing Machinery.
- [68] Ziqi Liu, Yu-Xiang Wang, and Alexander Smola. Fast differentially private matrix factorization. In *Proceedings of the 9th ACM Conference on Recommender Systems*, page 171–178, 2015.
- [69] Darakhshan Mir, Shan Muthukrishnan, Aleksandar Nikolov, and Rebecca N Wright. Pan-private algorithms via statistics on sketches. In *Proceedings of the thirtieth ACM SIGMOD-SIGACT-SIGART symposium on Principles of database systems*, pages 37–48, 2011.
- [70] T-H Hubert Chan, Mingfei Li, Elaine Shi, and Wenchang Xu. Differentially private continual monitoring of heavy hitters from distributed streams. In *International Symposium on Privacy Enhancing Technologies Symposium*, pages 140–159. Springer, 2012.
- [71] Robert E Schapire and Yoram Singer. Improved boosting algorithms using confidence-rated predictions. In *Proceedings of the 11th annual conference on Computational learning theory*, pages 80–91, 1998.
- [72] Avrim Blum, Cynthia Dwork, Frank McSherry, and Kobbi Nissim. Practical privacy: The sulq framework. pages 128–138, 06 2005.
- [73] Han Wang, Hanbin Hong, Li Xiong, Zhan Qin, and Yuan Hong. L-SRR: local differential privacy for location-based services with staircase randomized response. In *Proceedings of the ACM Conference on Computer and Communications Security*, pages 2809–2823, 2022.

- [74] Matthew Joseph, Aaron Roth, Jonathan Ullman, and Bo Waggoner. Local differential privacy for evolving data. *NeurIPS*, 2018.
- [75] Xiaochen Li, Weiran Liu, Jian Lou, Yuan Hong, Lei Zhang, Zhan Qin, and Kui Ren. Local differentially private heavy hitter detection in data streams with bounded memory. In *Proceedings of the 2024 ACM SIGMOD International Conference on Management of Data*, 2024.
- [76] Ergute Bao, Yin Yang, Xiaokui Xiao, and Bolin Ding. Cgm: an enhanced mechanism for streaming data collection with local differential privacy. *Proceedings of the VLDB Endowment*, 14(11):2258–2270, 2021.
- [77] Xuebin Ren, Liang Shi, Weiren Yu, Shusen Yang, Cong Zhao, and Zongben Xu. Ldp-ids: Local differential privacy for infinite data streams. In *Proceedings of the 2022 International Conference on Management of Data*, pages 1064–1077, 2022.
- [78] Nicolas Papernot. Machine learning at scale with differential privacy in TensorFlow. In *2019 USENIX Conference on Privacy Engineering Practice and Respect*, 2019.
- [79] Martin Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security*, page 308–318, 2016.
- [80] Jaideep Vaidya, Basit Shafiq, Anirban Basu, and Yuan Hong. Differentially private naive bayes classification. In *2013 IEEE/WIC/ACM International Conferences on Web Intelligence*, pages 571–576, 2013.
- [81] Nan Wu, Farhad Farokhi, David Smith, and Mohamed Ali Kaafar. The value of collaboration in convex machine learning with differential privacy. In *2020 IEEE Symposium on Security and Privacy (SP)*, pages 304–317. IEEE, 2020.
- [82] Meisam Mohammady, Shangyu Xie, Yuan Hong, Mengyuan Zhang, Lingyu Wang, Makan Pourzandi, and Mourad Debbabi. R2dp: A universal and automated approach to optimizing the randomization mechanisms of differential privacy for utility metrics with no known optimal distributions. In *CCS'20*, 2020.
- [83] Kang Wei, Jun Li, Ming Ding, Chuan Ma, Howard H Yang, Farhad Farokhi, Shi Jin, Tony QS Quek, and H Vincent Poor. Federated learning with differential privacy: Algorithms and performance analysis. *IEEE Transactions on Information Forensics and Security*, 15:3454–3469, 2020.
- [84] Arik Friedman and Assaf Schuster. Data mining with differential privacy. In *Proceedings of the 16th ACM SIGKDD international conference on Knowledge discovery and data mining*, pages 493–502, 2010.
- [85] Han Wang, Jayashree Sharma, Shuya Feng, Kai Shu, and Yuan Hong. A model-agnostic approach to differentially private topic mining. In *The 28th ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, pages 1835–1845, 2022.
- [86] Frank McSherry and Ilya Mironov. Differentially private recommender systems: Building privacy into the netflix prize contenders. In *Proceedings of the 15th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, page 627–636, 2009.
- [87] Chaochao Chen, Huiwen Wu, Jiajie Su, Lingjuan Lyu, Xiaolin Zheng, and Li Wang. Differential private knowledge transfer for privacy-preserving cross-domain recommendation. In *Proceedings of the ACM Web Conference*, page 1455–1465, 2022.

- [88] Tingting Feng, Yuchun Guo, and Yishuai Chen. A differential private collaborative filtering framework based on privacy-relevance of topics. In *2016 IEEE Symposium on Computers and Communication (ISCC)*, pages 946–951, 2016.
- [89] Rina Okada, Kazuto Fukuchi, and Jun Sakuma. Differentially private analysis of outliers. In *Joint European Conference on Machine Learning and Knowledge Discovery in Databases*, 2015.
- [90] Spencer Bloch. *Higher regulators, algebraic K-theory, and zeta functions of elliptic curves*. Number 11. American Mathematical Soc., 2000.
- [91] P. Mateev. On the entropy of the multinomial distribution. *Theory of Probability & Its Applications*, 23(1):188–190, 1978.
- [92] H. Brendan McMahan, Daniel Ramage, Kunal Talwar, and Li Zhang. Learning differentially private recurrent language models. In *International Conference on Learning Representations*, 2018.
- [93] Jamie Hayes, Luca Melis, George Danezis, and Emiliano De Cristofaro. Logan: Membership inference attacks against generative models. In *Proceedings on Privacy Enhancing Technologies*, volume 2019, pages 133–152, 2019.
- [94] Jingwen Zhang, Jiale Zhang, Junjun Chen, and Shui Yu. Gan enhanced membership inference: A passive local attack in federated learning. In *ICC. IEEE*, 2020.
- [95] Milad Nasr, Reza Shokri, and Amir Houmansadr. Comprehensive privacy analysis of deep learning: Passive and active white-box inference attacks against centralized and federated learning. In *IEEE S&P. IEEE*, 2019.
- [96] Xiao Jin, Pin-Yu Chen, Chia-Yi Hsu, Chia-Mu Yu, and Tianyi Chen. Cafe: Catastrophic data leakage in vertical federated learning. *NeurIPS*, 2021.
- [97] Haimei Gong, Liangjun Jiang, Xiaoyang Liu, Yuanqi Wang, Omary Gastro, Lei Wang, Ke Zhang, and Zhen Guo. Gradient leakage attacks in federated learning. *Artificial Intelligence Review*, 56 (Suppl 1):1337–1374, 2023.
- [98] Zhuohang Li, Jiaxin Zhang, Luyang Liu, and Jian Liu. Auditing privacy defenses in federated learning via generative gradient leakage. In *CVPR*, 2022.
- [99] Karan Ganju, Qi Wang, Wei Yang, Carl A Gunter, and Nikita Borisov. Property inference attacks on fully connected neural networks using permutation invariant representations. In *CCS*, 2018.
- [100] Chen Chen, Lingjuan Lyu, Han Yu, and Gang Chen. Practical attribute reconstruction attack against federated learning. *IEEE Transactions on Big Data.*, 2022.
- [101] Lingjuan Lyu and Chen Chen. A novel attribute reconstruction attack in federated learning. *arXiv preprint arXiv:2108.06910*, 2021.
- [102] Caridad Arroyo Arevalo, Sayedeh Leila Noorbakhsh, Yun Dong, Yuan Hong, and Binghui Wang. Task-agnostic privacy-preserving representation learning for federated learning against attribute inference attacks. In *AAAI*, 2024.
- [103] Mohammed Adnan, Shivam Kalra, Jesse C Cresswell, Graham W Taylor, and Hamid R Tizhoosh. Federated learning and differential privacy for medical image analysis. *Scientific reports*, 2022.

- [104] Qinbin Li, Zhaomin Wu, Zeyi Wen, and Bingsheng He. Privacy-preserving gradient boosting decision trees. In *AAAI*, 2020.
- [105] Xiaokui Xiao, Guozhang Wang, and Johannes Gehrke. Differential privacy via wavelet transforms. *IEEE Transactions on knowledge and data engineering*, 23(8):1200–1214, 2010.
- [106] Yuqing Zhu, Jinshuo Dong, and Yu-Xiang Wang. Optimal accounting of differential privacy via characteristic function. In *AISTATS*, 2022.
- [107] Hua Wang, Sheng Gao, Huanyu Zhang, Milan Shen, and Weijie J Su. Analytical composition of differential privacy via the edgeworth accountant. *arXiv preprint arXiv:2206.04236*, 2022.
- [108] Yu-Xiang Wang, Borja Balle, and Shiva Prasad Kasiviswanathan. Subsampled rényi differential privacy and analytical moments accountant. In *AISTATS*, 2019.
- [109] Tim van Erven and Peter Harremos. Rényi divergence and kullback-leibler divergence. *IEEE Transactions on Information Theory*, 2014.
- [110] Naman Agarwal, Peter Kairouz, and Ziyu Liu. The skellam mechanism for differentially private federated learning. *NIPS*, 2021.
- [111] Qinqing Zheng, Shuxiao Chen, Qi Long, and Weijie Su. Federated f-differential privacy. In *International Conference on Artificial Intelligence and Statistics*, 2021.
- [112] Aleksei Triastcyn and Boi Faltings. Federated learning with bayesian differential privacy. In *2019 IEEE International Conference on Big Data*. IEEE, 2019.
- [113] Lu Shi, Jiangang Shu, Weizhe Zhang, and Yang Liu. Hfl-dp: Hierarchical federated learning with differential privacy. In *GLOBECOM*, 2021.
- [114] Kang Wei, Jun Li, Ming Ding, Chuan Ma, Howard H. Yang, Farhad Farokhi, Shi Jin, Tony Q. S. Quek, and H. Vincent Poor. Federated learning with differential privacy: Algorithms and performance analysis. *TIFS*, 2020.
- [115] Robin C Geyer, Tassilo Klein, and Moin Nabi. Differentially private federated learning: A client level perspective. *arXiv preprint arXiv:1712.07557*, 2017.
- [116] Rui Hu, Yuanxiong Guo, Hongning Li, Qingqi Pei, and Yanmin Gong. Personalized federated learning with differential privacy. *IEEE Internet of Things Journal*, 2020.
- [117] Jie Fu, Yuan Hong, Xinpeng Ling, Leixia Wang, Xun Ran, Zhiyu Sun, Wendy Hui Wang, Zhili Chen, and Yang Cao. Differentially private federated learning: A systematic review. *CoRR*, abs/2405.08299, 2024.
- [118] Quan Geng and Pramod Viswanath. The optimal mechanism in differential privacy. In *IEEE ISIT*, pages 2371–2375. IEEE, 2014.
- [119] Timur Garipov, Pavel Izmailov, Dmitrii Podoprikin, Dmitry P. Vetrov, and Andrew Gordon Wilson. Loss surfaces, mode connectivity, and fast ensembling of dnns. In *NeurIPS*, 2018.
- [120] Pu Zhao, Pin-Yu Chen, Payel Das, Karthikeyan Natesan Ramamurthy, and Xue Lin. Bridging mode connectivity in loss landscapes and adversarial robustness. In *ICLR*, 2020.

- [121] Akhilesh Gotmare, Nitish Shirish Keskar, Caiming Xiong, and Richard Socher. Using mode connectivity for loss landscape analysis. *CoRR*, 2018.
- [122] Brendan McMahan, Eider Moore, Daniel Ramage, Seth Hampson, and Blaise Aguera y Arcas. Communication-efficient learning of deep networks from decentralized data. In *Artificial intelligence and statistics*, pages 1273–1282, 2017.
- [123] Ahmed Salem, Giovanni Cherubin, David Evans, Boris Köpf, Andrew Paverd, Anshuman Suri, Shruti Tople, and Santiago Zanella-Béguelin. Sok: Let the privacy games begin! a unified treatment of data inference privacy in machine learning. In *IEEE S&P*. IEEE, 2023.
- [124] Matt Fredrikson, Somesh Jha, and Thomas Ristenpart. Model inversion attacks that exploit confidence information and basic countermeasures. In *CCS*, pages 1322–1333, 2015.
- [125] Ligeng Zhu, Zhijian Liu, and Song Han. Deep leakage from gradients. *NeurIPS*, 32, 2019.
- [126] Kang Wei, Jun Li, Chuan Ma, Ming Ding, Wen Chen, Jun Wu, Meixia Tao, and H. Vincent Poor. Personalized federated learning with differential privacy and convergence guarantee. *IEEE TIFS*, 2023.
- [127] Xinpeng Ling, Jie Fu, and Zhili Chen. Adaptive local steps federated learning with differential privacy driven by convergence analysis. *arXiv preprint arXiv:2308.10457*, 2023.
- [128] Lichao Sun, Jianwei Qian, and Xun Chen. Ldp-fl: Practical private aggregation in federated learning with local differential privacy. *arXiv preprint arXiv:2007.15789*, 2020.
- [129] Antonious M Girgis, Deepesh Data, Suhas Diggavi, Peter Kairouz, and Ananda Theertha Suresh. Shuffled model of federated learning: Privacy, accuracy and communication trade-offs. *IEEE journal on selected areas in information theory*, 2(1):464–478, 2021.
- [130] Antonious M Girgis, Deepesh Data, Suhas Diggavi, Ananda Theertha Suresh, and Peter Kairouz. On the rényi differential privacy of the shuffle model. In *ACM SIGSAC Conference on Computer and Communications Security*, 2021.
- [131] Zhuotao Lian, Weizheng Wang, and Chunhua Su. Cofel: Communication-efficient and optimized federated learning with local differential privacy. In *ICC 2021-IEEE International Conference on Communications*. IEEE, 2021.
- [132] Cynthia Dwork. Differential privacy. In *Automata, Languages and Programming: 33rd International Colloquium, ICALP 2006, Venice, Italy, July 10-14, 2006, Proceedings, Part II 33*. Springer, 2006.
- [133] Ilya Mironov, Kunal Talwar, and Li Zhang. R\'enyi differential privacy of the sampled gaussian mechanism. *arXiv preprint arXiv:1908.10530*, 2019.
- [134] Christopher M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer-Verlag, Berlin, Heidelberg, 2006.
- [135] Peter Auer et al. Exponentially many local minima for single neurons. In *NIPS*, volume 8, 1995.
- [136] Anna Choromanska et al. The loss surfaces of multilayer networks. In *Proceedings of the Eighteenth International Conference on Artificial Intelligence and Statistics*, 2015.

- [137] Jeong Joon Park, Peter Florence, Julian Straub, Richard Newcombe, and Steven Lovegrove. Deep sdf: Learning continuous signed distance functions for shape representation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, 2019.
- [138] Ricky TQ Chen, Yulia Rubanova, Jesse Bettencourt, and David K Duvenaud. Neural ordinary differential equations. *NeurIPS*, 31, 2018.
- [139] Thee Chanyaswad, Alex Dytso, H Vincent Poor, and Prateek Mittal. Mvg mechanism: Differential privacy under matrix-valued query. In *CCS*, 2018.
- [140] Pranav Shankar Subramani, Nicholas Vadivelu, and Gautam Kamath. Enabling fast differentially private SGD via just-in-time compilation and vectorization. In A. Beygelzimer, Y. Dauphin, P. Liang, and J. Wortman Vaughan, editors, *NeurIPS*, 2021.
- [141] Sebastian U. Stich. Local sgd converges fast and communicates little. In *International Conference on Learning Representations*, 2019.
- [142] Xiang Li, Kaixuan Huang, Wenhao Yang, Shusen Wang, and Zhihua Zhang. On the convergence of fedavg on non-iid data. In *International Conference on Learning Representations*, 2020.
- [143] Hiroki Masuda, Kentaro Kita, Yuki Koizumi, Junji Takemasa, and Toru Hasegawa. Model fragmentation, shuffle and aggregation to mitigate model inversion in federated learning. In *2021 IEEE International Symposium on Local and Metropolitan Area Networks (LANMAN)*, pages 1–6, 2021.
- [144] Ruixuan Liu, Yang Cao, Hong Chen, Ruoyang Guo, and Masatoshi Yoshikawa. Flame: Differentially private federated learning in the shuffle model. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 8688–8696, 2021.
- [145] Úlfar Erlingsson, Vitaly Feldman, Ilya Mironov, Ananth Raghunathan, Kunal Talwar, and Abhradeep Thakurta. Amplification by shuffling: From local to central differential privacy via anonymity. In *SODA*, 2019.
- [146] Antonious Girgis, Deepesh Data, and Suhas Diggavi. Renyi differential privacy of the subsampled shuffle model in distributed learning. *NeurIPS*, 2021.
- [147] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 1998.
- [148] Sammy Kus. Medical mnist. *Mendeley Data*, 2022.
- [149] Alex Krizhevsky, Vinod Nair, and Geoffrey Hinton. Cifar-10 (canadian institute for advanced research). Technical report, University of Toronto, 2009.
- [150] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *CVPR*, pages 770–778, 2016.
- [151] Daniel J Beutel, Taner Topal, Akhil Mathur, Xinchu Qiu, Javier Fernandez-Marques, Yan Gao, et al. Flower: A friendly federated learning research framework. *arXiv preprint arXiv:2007.14390*, 2020.
- [152] Samuel Yeom, Irene Giacomelli, Matt Fredrikson, and Somesh Jha. Privacy risk in machine learning: Analyzing the connection to overfitting. In *2018 IEEE 31st Computer Security Foundations Symposium (CSF)*, pages 268–282. IEEE, 2018.

- [153] Vincent Bindschaedler, Reza Shokri, and Carl A. Gunter. Plausible deniability for privacy-preserving data synthesis. *Proc. VLDB Endow.*, 2017.
- [154] Zhifei Zhang, Yang Song, and Hairong Qi. Age progression/regression by conditional adversarial autoencoder. In *CVPR*, pages 5810–5818, 2017.
- [155] Sai Praneeth Karimireddy, Satyen Kale, Mehryar Mohri, Sashank Reddi, Sebastian Stich, and Ananda Theertha Suresh. Scaffold: Stochastic controlled averaging for federated learning. In *ICML*. PMLR, 2020.
- [156] Xing Wu, Zhaowang Liang, and Jianjia Wang. Fedmed: A federated learning framework for language modeling. *Sensors*, 2020.
- [157] Xiaotong Yuan and Ping Li. On convergence of fedprox: Local dissimilarity invariant bounds, non-smoothness and beyond. *NeurIPS*, 2022.
- [158] Tao Hong, Pierre Pinson, Shu Fan, Hamidreza Zareipour, Alberto Troccoli, and Rob Hyndman. Probabilistic energy forecasting: Global energy forecasting competition 2014 and beyond. *International Journal of Forecasting*, 2016.
- [159] Joseph Geumlek, Shuang Song, and Kamalika Chaudhuri. Renyi differential privacy mechanisms for posterior sampling. *NeurIPS*, 30, 2017.
- [160] Han Wang, Jayashree Sharma, Shuya Feng, Kai Shu, and Yuan Hong. A model-agnostic approach to differentially private topic mining. In *KDD’22*, 2022.
- [161] Vitaly Feldman, Audra McMillan, and Kunal Talwar. Stronger privacy amplification by shuffling for rényi and approximate differential privacy. In *SODA*. SIAM, 2023.
- [162] Abhishek Bhowmick, John Duchi, Julien Freudiger, Gaurav Kapoor, and Ryan Rogers. Protection against reconstruction and its applications in private federated learning. *arXiv preprint arXiv:1812.00984*, 2018.
- [163] Stacey Truex, Nathalie Baracaldo, Ali Anwar, Thomas Steinke, Heiko Ludwig, Rui Zhang, and Yi Zhou. A hybrid approach to privacy-preserving federated learning. In *Proceedings of the 12th ACM workshop on artificial intelligence and security*, pages 1–11, 2019.
- [164] Qinbin Li, Zeyi Wen, Zhaomin Wu, Sixu Hu, Naibo Wang, Yuan Li, Xu Liu, and Bingsheng He. A survey on federated learning systems: vision, hype and reality for data privacy and protection. *IEEE Transactions on Knowledge and Data Engineering*, 2021.
- [165] Jing Li, Daoxin Lin, Huaxin Shu, Yuan Wang, and Xindong Liu. Federated learning with data and model privacy preservation. *arXiv preprint arXiv:2306.08005*, 2023.
- [166] Mingzhe Xu, Gen Li, Jiaojiao Zheng, Qian Xiao, and Jian Liu. Federated learning with privacy-preserving and model protection. *Inf. Sci.*, 2022.
- [167] Haonan Zhu and Ruoxi Ding. Efficient differentially private federated learning for heterogeneous data. *arXiv preprint arXiv:2306.08750*, 2023.
- [168] Xiaochen Chen, Yongqiang Li, Yuncheng Wang, Jiacheng Liu, and Kim-Kwang Raymond Choo. Federated learning with differential privacy and secure multiparty computation. *IEEE TIFS*, 2023.

- [169] Stanislav Fort, Stefanie Günther, Zoltán Szabó, Andrew McMillan, and Vitaly Feldman. Privacy amplification by iteration in federated learning. *arXiv preprint arXiv:2305.15046*, 2023.
- [170] Xinwei Liu, Penghui Zhao, Chao Li, and Jiawen Zhang. Tighter privacy bounds for subsampled mechanisms in federated learning. *arXiv preprint arXiv:2304.02140*, 2023.
- [171] Hao Ding, Xiangyu Gao, Ming Li, and Qian Tang. Differentially private federated learning with local randomization and adaptive optimization. *IEEE Transactions on Information Forensics and Security*, 2023.
- [172] Matta Varun, Shuya Feng, Han Wang, Shamik Sural, and Yuan Hong. Towards accurate and stronger local differential privacy for federated learning with staircase randomized response. In *CODASPY*, pages 307–318, 2024.
- [173] Yitong Luo, Jiaheng Zhang, Qiaoyu Tan, Yu-Xiang Wang, and Qiang Yang. Federated multi-task learning with differential privacy. *arXiv preprint arXiv:2307.03449*, 2023.
- [174] Wei Zhou, Jiacheng Mao, Ankai Xu, Shutian Wan, Qianqian Gong, Xin Li, Yi Wen, and Jian Li. Tight privacy composition analysis for federated learning. In *ICML*, 2023.
- [175] Idan Dagan, Tomer Gafni, Oleksii Romanenko, and Idit Keidar. Federated learning of large language models with differential privacy. *arXiv preprint arXiv:2306.10635*, 2023.
- [176] Lichao Sun, Jianwei Qian, and Xiang Chen. Federated learning with adaptive differential privacy. *IEEE TDSC*, 2023.
- [177] Amol Khanna, Vincent Schaffer, Gamze Gürsoy, and Mark Gerstein. Privacy-preserving model training for disease prediction using federated learning with differential privacy. In *EMBC*, 2022.
- [178] Qingrong Chen, Chong Xiang, Minhui Xue, Bo Li, Nikita Borisov, Dali Kaarfar, and Haojin Zhu. Differentially private data generative models. *arXiv preprint arXiv:1812.02274*, 2018.
- [179] Yi Liu, Jialiang Peng, Jiawen Kang, Abdullah M. Iliyasu, Dusit Niyato, and Ahmed A. Abd El-Latif. A secure federated learning framework for 5g networks. *arXiv preprint arXiv:2001.05637*, 2020.
- [180] Mohammad Naseri, Jamie Hayes, and Emiliano De Cristofaro. Local and central differential privacy for robustness and privacy in federated learning. *arXiv preprint arXiv:2009.03561*, 2020.
- [181] Lichao Sun and Lingjuan Lyu. Federated model distillation with noise-free differential privacy. *arXiv preprint arXiv:2012.04187*, 2020.
- [182] Raouf Kerkouche, Gergely Ács, and Claude Castelluccia. Federated learning in adversarial settings. *arXiv preprint arXiv:2001.05641*, 2020.
- [183] Qian Chen, Zilong Wang, Wenjing Zhang, and Xiaodong Lin. Ppt: A privacy-preserving global model training protocol for federated learning in p2p networks. *arXiv preprint arXiv:2101.02281*, 2021.
- [184] Md Tamjid Hossain, Shafkat Islam, Shahriar Badsha, and Haoting Shen. Desmp: Differential privacy-exploited stealthy model poisoning attacks in federated learning. *arXiv preprint arXiv:2102.03070*, 2021.

- [185] Tiantian Feng, Raghuveer Peri, and Shrikanth Narayanan. User-level differential privacy against attribute inference attack of speech emotion recognition in federated learning. *arXiv preprint arXiv:2202.01684*, 2022.
- [186] Yuntao Wang, Haixia Peng, Zhou Su, Tom H Luan, Abderrahim Benslimane, and Yuan Wu. A platform-free proof of federated learning consensus mechanism for sustainable blockchains. *arXiv preprint arXiv:2202.01884*, 2022.
- [187] Xiang Li, Kaixuan Huang, Wenhao Yang, Shusen Wang, and Zhihua Zhang. On the convergence of fedavg on non-iid data. *arXiv preprint arXiv:1907.02189*, 2019.
- [188] Léon Bottou, Frank E Curtis, and Jorge Nocedal. Optimization methods for large-scale machine learning. *Siam Review*, 2018.
- [189] Benjamin Recht et al. Hogwild: A lock-free approach to parallelizing stochastic gradient descent. 2011.
- [190] Joram Soch, The Book of Statistical Proofs, Maja, Pietro Monticone, Thomas J. Faulkenberry, Alex Kipnis, Kenneth Petrykowski, Carsten Allefeld, Heiner Atze, Adam Knapp, Ciarán D. McInerney, Lo4ding00, and amvosk. StatProofBook/StatProofBook.github.io: StatProofBook 2023, 2024.
- [191] Yuhong Yang. Elements of information theory, 2008.
- [192] Anub Rao. A comprehensive study on cross-entropy convexity, 2010.
- [193] Leonard Berrada, Andrew Zisserman, and M. Pawan Kumar. Smooth loss functions for deep top-k classification. *ArXiv*, 2018.
- [194] Xinyu Yang, Zichen Wen, Wenjie Qu, Zhaorun Chen, Zhiying Xiang, Beidi Chen, and Huaxiu Yao. Memorization and privacy risks in domain-specific large language models. In *ICLR 2024 Workshop on Reliable and Responsible Foundation Models*, 2013.
- [195] Cynthia Dwork, Frank McSherry, Kobbi Nissim, and Adam Smith. Calibrating noise to sensitivity in private data analysis. In *Proceedings of the Third Theory of Cryptography Conference (TCC)*, TCC '06, pages 265–284. Springer, 2006.
- [196] Martín Abadi, Andy Chu, Ian Goodfellow, H. Brendan McMahan, Ilya Mironov, Kunal Talwar, and Li Zhang. Deep learning with differential privacy. In *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security, CCS '16*, pages 308–318. ACM, 2016.
- [197] Yu Zheng, Wenchao Zhang, Yonggang Zhang, Wei Song, Kai Zhou, and Bo Han. Rethinking improved privacy-utility trade-off with pre-existing knowledge for dp training. *arXiv preprint arXiv:2409.03344*, 2024.
- [198] Lijie Hu, Ivan Habernal, Lei Shen, and Di Wang. Differentially private natural language models: Recent advances and future directions. In Yvette Graham and Matthew Purver, editors, *Findings of the Association for Computational Linguistics: EACL 2024*, pages 478–499, St. Julian's, Malta, March 2024. Association for Computational Linguistics.
- [199] Nicolas Fernandes, Mark Dras, and Annabelle McIver. Generalised differential privacy for text document processing. In *Proceedings of the 8th International Conference on Principles of Security and Trust*. Springer, 2019.

- [200] Oluwaseyi Feyisetan, Borja Balle, Thomas Drake, and Timour K. Diethe. Privacy- and utility-preserving textual analysis via calibrated multivariate perturbations. In *Proceedings of the 13th ACM International Conference on Web Search and Data Mining*, WSDM '20, pages 178–186, 2020.
- [201] Stefan Arnold. Inspecting the representation manifold of differentially-private text. In Ivan Habernal, Sepideh Ghanavati, Vijayanta Jain, Timour Igamberdiev, and Shomir Wilson, editors, *Proceedings of the Sixth Workshop on Privacy in Natural Language Processing*, pages 53–59, Albuquerque, New Mexico, April 2025. Association for Computational Linguistics.
- [202] Bochao Liu, Jianghu Lu, Pengju Wang, Junjie Zhang, Dan Zeng, Zhenxing Qian, and Shiming Ge. Privacy-preserving student learning with differentially private data-free distillation. In *2022 IEEE 24th International Workshop on Multimedia Signal Processing (MMSP)*, pages 01–06, 2022.
- [203] Nicolas Papernot, Martín Abadi, Úlfar Erlingsson, Ian Goodfellow, and Kunal Talwar. Semi-supervised knowledge transfer for deep learning from private training data. In *Proceedings of the 5th International Conference on Learning Representations*, 2017.
- [204] Nicolas Papernot, Shuang Song, Ilya Mironov, Ananth Raghunathan, Kunal Talwar, and Úlfar Erlingsson. Scalable private learning with PATE. In *Proceedings of the 6th International Conference on Learning Representations*, 2018.
- [205] Dominik Fay, Jens Sjölund, and Tobias J. Oechtering. Private learning via knowledge transfer with high-dimensional targets. In *ICASSP 2022 - 2022 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, pages 3873–3877, 2022.
- [206] H. Brendan McMahan, Daniel Ramage, Kunal Talwar, and Li Zhang. Learning differentially private recurrent language models. In *Proceedings of the 6th International Conference on Learning Representations*, 2018.
- [207] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. arXiv preprint arXiv:1503.02531, 2015.
- [208] Yujun Wang, Xiangyang Luo, and Konstantina Palla. Neighbourhood-aware differential privacy for word embeddings. In *Findings of ACL (IJCNLP)*, 2023.
- [209] Jeremy Beutel and et al. Private fine-tuning of large language models with differential privacy. In *MLSys*, 2022.
- [210] Rohan Anil, Shuang Song, and et al. Large-scale differentially private BERT. In *ArXiv*, 2021.
- [211] Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. Efficient estimation of word representations in vector space. In *ICLR (Workshop)*, 2013.
- [212] Matthieu Moreau and Tarek Benkhelif. Dpsgd strategies for cross-silo federated learning. In *2021 International Conference on Communications, Computing, Cybersecurity, and Informatics (CCCI)*, pages 1–5, 2021.
- [213] Joshua B Lee and Dan Suci. Scaling up differentially private deep learning with fast per-example gradient clipping. In *Advances in Neural Information Processing Systems*, volume 35, pages 10168–10181, 2022.

- [214] Galen Andrew, Om Thakkar, Brendan McMahan, and Swaroop Ramaswamy. Differentially private learning with adaptive clipping. *NeurIPS*, 2021.
- [215] Cynthia Dwork, Aaron Roth, et al. The algorithmic foundations of differential privacy. *Foundations and Trends® in Theoretical Computer Science*, 2014.
- [216] Frank McSherry and Kunal Talwar. Mechanism design via differential privacy. In *48th IEEE Symposium on Foundations of Computer Science (FOCS 2007)*, pages 94–103. IEEE, 2007.
- [217] Eugene Bagdasaryan and Vitaly Shmatikov. Differential privacy has disparate impact on model accuracy. In *NeurIPS Workshop on Privacy Preserving Machine Learning*, 2019. arXiv:1905.12101.
- [218] Matthew Fredrikson, Somesh Jha, and Thomas Ristenpart. Model inversion attacks that exploit confidence information and basic countermeasures. In *ACM CCS 2015*, pages 1322–1333. ACM, 2015.
- [219] Rachel Cummings, Brennan McMahan, and David Lie. A large-scale survey of public perception, knowledge, and concerns about differential privacy. In *CHI 2021*, pages Paper 123, 15 pages. ACM, 2021.
- [220] Weijian Xiong, Shikun Han, and Arvind Narayanan. Explaining differential privacy: Focusing on implications improves end-user understanding. In *31st USENIX Security Symposium*, pages 1795–1812, 2022.
- [221] Pedro Garrido, Raquel Hill, and Oriana Caivano. Mind the gap: How description styles affect understanding of differential privacy. In *Privacy Enhancing Technologies Symposium (PETS 2022)*, pages 196–215, 2022.
- [222] Jan Kührtreiber, Alexandra Brill, and Melanie Volkamer. Framing privacy: Evaluating explanations of differential privacy for novice users. In *CSCW 2022*, pages Paper 121, 25 pages. ACM, 2022.
- [223] Ryan Rogers, Anne Brown, Ilya Mironov, and Peter Kairouz. Sok: Differential privacy for machine learning — a practical roadmap. In *44th IEEE Symposium on Security and Privacy (S&P 2023)*, pages 1567–1585. IEEE, 2023.
- [224] Fatemeh Karegar, Silvia Lindtner, and Florian N. Kriegel. Helping data analysts understand differential privacy: Interface metaphors and mental models. In *CHI 2022*, pages Paper 252, 17 pages. ACM, 2022.
- [225] Ivoline C. Ngong, Brad Stenger, Joseph P. Near, and Yuanyuan Feng. Evaluating the usability of differential privacy tools with data practitioners. In *20th Symposium on Usable Privacy and Security (SOUPS 2024)*, 2024.
- [226] Shuya Feng, Jia He, and Maggie X. Cheng. Security analysis of block withholding attacks in blockchain. In *ICC 2021 - IEEE International Conference on Communications*, pages 1–6, 2021.
- [227] Can Jin, Hongwu Peng, Anxiang Zhang, Nuo Chen, Jiahui Zhao, Xi Xie, Kuangzheng Li, Shuya Feng, Kai Zhong, Caiwen Ding, et al. Rankflow: A multi-role collaborative reranking workflow utilizing large language models. In *Companion Proceedings of the ACM on Web Conference 2025*, pages 2484–2493, 2025.
- [228] Hanbin Hong, Shenao Yan, Shuya Feng, Yan Yan, and Yuan Hong. Galot: Generative active learning via optimizable zero-shot text-to-image generation. *arXiv preprint arXiv:2412.16227*, 2024.

- [229] Boyuan Feng, Yijiang Zheng, Ruting Cheng, Shuya Feng, Khashayar Vaziri, and James Hahn. Enhanced body composition estimation from 3d body scans.